

# Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀  
playing 🤖 with  
Python 🐍 and AI/ML!  


## Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	<b>Lab:</b> Python Startup: Basics	<b>Lab:</b> Feed-forward and Graph NNs	<b>Lab:</b> Graph Learning with PyTorch	<b>Lab:</b> Containers and Reproducibility	<b>Lab:</b> Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	<b>Lab:</b> Work environments, Python everywhere	<b>Lab:</b> Simple Transformer and LLM Use	<b>Lab:</b> Agent Frameworks	<b>Lab:</b> CI/CD Pipelines	<b>Lab:</b> Graph-based Assimilation
12:45–13:30	<b>Lunch Break</b>				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	<b>Lab:</b> GRIB, NetCDF and Obs Visualisation	<b>Lab:</b> RAG Pipeline	<b>Lab:</b> DAWID Exploration	<b>Lab:</b> Anemoi Training Pipeline	<b>Lab:</b> Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	<b>The AI Transformation</b>	Learning from Observations Only
16:15–17:00	<b>Lab:</b> Torch Tensors and First Neural Net	<b>Lab:</b> Radar, SAT and Multimodal Data	<b>Lab:</b> MLflow Hands-on	<b>Lab:</b> How work style could change	<b>Lab:</b> ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

## Lecture 2: Jupyter Notebooks, APIs and Servers

### Goal of this Lecture

- ▶ Work productively with Jupyter Notebooks
- ▶ Understand Notebooks as part of a scientific workflow
- ▶ Prepare the ground for reproducible ML experiments

### Focus

- ▶ Not UI details, but how things fit together
- ▶ From exploration to engineering

### Topics Overview

- ▶ Jupyter Notebooks: Kernel, Server, Browser
- ▶ Environments and package management
- ▶ Markdown, magic commands, shell integration
- ▶ Data visualization as quality control
- ▶ APIs as a structuring principle
- ▶ Local, library and web APIs
- ▶ Native code integration (Fortran / C++)

# Jupyter Notebook

- ▶ Cells with comments
- ▶ Cells with code
- ▶ Cells with output

jupyter Jupyter\_Demo1 Last Checkpoint: 20 seconds ago

File Edit View Run Kernel Settings Help Trusted JupyterLab ⌂ (ropy) C

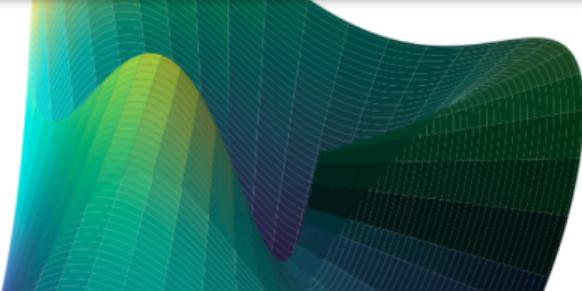
[1]:

```
import numpy as np, matplotlib.pyplot as plt
from matplotlib.colors import LightSource

u = np.linspace(0, 2*np.pi, 320); v = np.linspace(-0.6, 0.6, 80); U, V = np.meshgrid(u, v)

# Möbius-like twisted surface with extra complexity
X = (1 + V*np.cos(U/2)) * np.cos(U); Y = (1 + V*np.cos(U/2)) * np.sin(U)
Z = V*np.sin(U/2) + 0.3*np.sin(3*U)

ls = LightSource(azdeg=22, altdeg=65); rgb = ls.shade(Z, cmap=plt.cm.viridis, vert_exag=0.9)
plt.figure(figsize=(12,12)); ax = plt.axes(projection="3d")
ax.plot_surface(X, Y, Z, facecolors=rgb, linewidth=0); ax.set_axis_off(); plt.show()
```



## Why Jupyter Notebooks in Science and ML?

### Why Notebooks Matter

- ▶ Rapid exploration of data and ideas
- ▶ Immediate feedback via plots and diagnostics
- ▶ Combine code, results and explanation

### In Research Contexts

- ▶ Hypothesis testing and prototyping
- ▶ Understanding data and model behavior
- ▶ Bridging theory and implementation

### Why They Scale Beyond Prototyping

- ▶ Documentation of decisions and assumptions
- ▶ Reproducible experiments (when done right)
- ▶ Natural interface to libraries and APIs
- ▶ Gateway to larger workflows and services

### Key Message

- ▶ Notebooks are tools, not products
- ▶ Value comes from disciplined usage

## Jupyter Architecture: Browser – Server – Kernel

### Three Core Components

- ▶ **Browser** User interface: notebooks, plots, interaction
- ▶ **Jupyter Server** Manages files, sessions, security
- ▶ **Kernel** Executes Python code, holds state

### Key Idea

- ▶ UI and computation are **decoupled**

### Why This Matters

- ▶ Code runs in the **kernel**, not in the browser
- ▶ Kernel state persists across cells
- ▶ Server and kernel may run **remotely**
- ▶ Multiple notebooks can share one kernel

### Strength and Risk

- ▶ Powerful interactive workflow
- ▶ Risk of hidden state and irreproducibility

## Reproducibility in Jupyter Notebooks

### Why Reproducibility Matters

- ▶ Results must be repeatable
- ▶ Experiments must be explainable
- ▶ Others (and you later) must trust them

### Typical Problems

- ▶ Hidden kernel state
- ▶ Cells executed out of order
- ▶ Undocumented dependencies

### Four Simple Rules

- ▶ Restart kernel and **Run All**
- ▶ Clear, explicit import and parameter cells
- ▶ Save outputs (plots, files, artefacts)
- ▶ Document dependencies and environment

### Key Message

- ▶ A notebook is an **executable document**

## Installing Packages in Jupyter Notebooks

### Recommended Workflow

- ▶ Install packages either outside in the right virtual environment or directly in the notebook
- ▶ Use pip install or !pip install
- ▶ Packages are installed into the running kernel

### Why This Works

- ▶ Jupyter uses the kernel's Python environment
- ▶ pip is typically bound to this Python

### Good Practice: Check Once

Check active environment

```
1 import sys
2 print(f"Python executable: {sys.executable}")
```

### Key Message

- ▶ Use pip where you run your code
- ▶ Verify the environment if something looks wrong

## Markdown and Narrative Computing

### Why Markdown Matters

- ▶ Makes notebooks **readable**
- ▶ **Explains** intent, not just results
- ▶ Turns experiments into documents

### Narrative Computing

- ▶ Code, text and results in one place
- ▶ **Reasoning** becomes explicit
- ▶ Supports review and reuse

### Minimum You Should Use

- ▶ Headings (#, ##)
- ▶ Bullet lists (-)
- ▶ Inline formulas ( $x^2$ )
- ▶ Short explanations

### Markdown Cells

- ▶ Change cell type: Esc → M

### Key Message

- ▶ A notebook should read like a **lab notebook**

## Magic Commands and Shell Access

### Why Magic Commands

- ▶ Speed up interactive work
- ▶ Reduce boilerplate code
- ▶ Support exploration and debugging

### Line vs Cell Magics

- ▶ % for single-line commands
- ▶ %% for whole cells

### The Few You Really Need

- ▶ %timeit (runtime)
- ▶ %%writefile (save code)
- ▶ %%bash (run shell)
- ▶ !pip (install packages)
- ▶ !ls (files)

### Key Message

- ▶ Use magic and shell commands  
**sparingly**

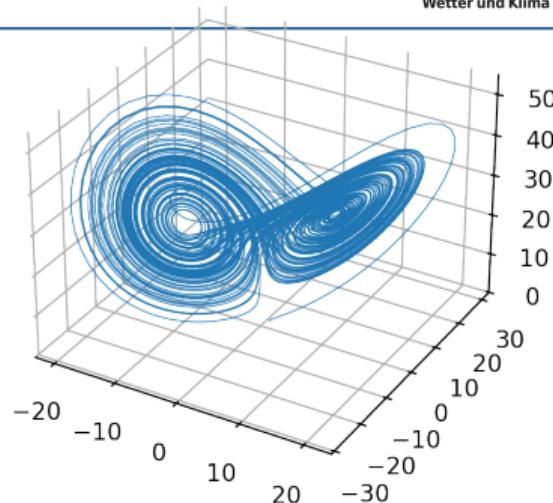
## Visualization as Quality Control

### Why Visualization Matters

- ▶ See patterns and anomalies
- ▶ Detect bugs early
- ▶ Build intuition about data and models

### Typical Questions

- ▶ Does this look reasonable?
- ▶ Are scales and units correct?
- ▶ Are outliers expected?



### Static vs Interactive

- ▶ Static: documentation, papers
- ▶ Interactive: exploration, debugging

### Key Message

- ▶ Plot early, plot often!

## Matplotlib

- ▶ Low-level, explicit control
- ▶ Publication-ready figures
- ▶ Default for scientific Python

### Matplotlib (static)

```
1 plt.plot(x, y)
2 plt.xlabel("x");
3 plt.ylabel("y")
4 plt.savefig("plot.png")
```

## When to use

- ▶ Final figures
- ▶ Full control over layout

## Seaborn

- ▶ Built on Matplotlib
- ▶ Statistical defaults
- ▶ Fast insight into distributions

### Seaborn (statistical)

```
1 sns.kdeplot(x=data)
```

## Plotly

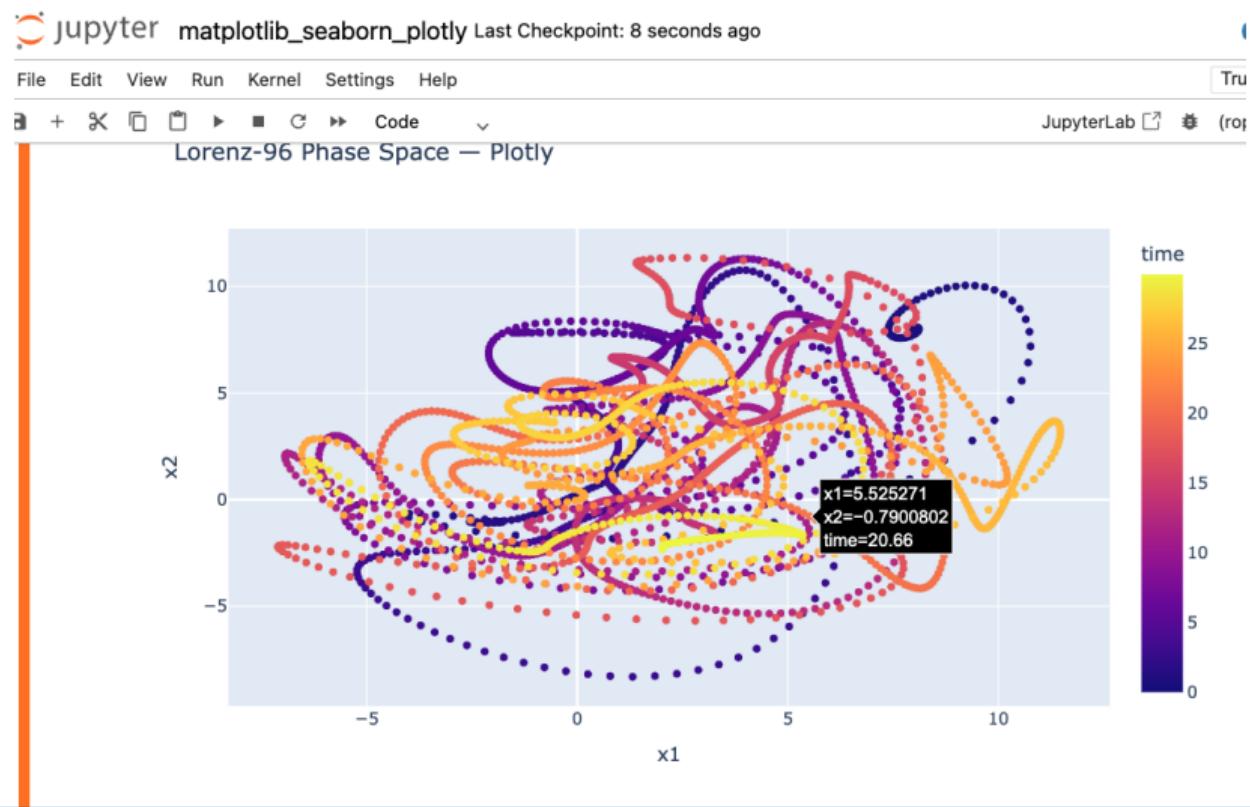
- ▶ Interactive plots
- ▶ Zoom, hover, selection

### Plotly (interactive)

```
1 fig = px.scatter(df, x="x",
2                   y="y"); fig.show()
```

## Interactive Plot via Plotly

- ▶ Explore your data
- ▶ Learn, check!
- ▶



## What Is an API — and What Is Not

### What an API Is

- ▶ A **defined interface** to call functionality
- ▶ Clear inputs, outputs and behavior
- ▶ Stable contract between components

### Typical API Examples

- ▶ Python functions and classes
- ▶ Library interfaces (e.g. NumPy)
- ▶ Web endpoints (HTTP/REST)

### What an API Is *Not*

- ▶ Just installable code
- ▶ A script without defined entry points
- ▶ A collection of loosely coupled functions

### Important Distinction

- ▶ Installation is **distribution**
- ▶ An API is **how you call the code**

### Key Message

- ▶ Packages may **expose** APIs
- ▶ Installation alone does not define one

## APIs as a Scaling Principle

### Core Idea

- ▶ APIs defines a contract on how functionality is accessed
- ▶ Clear separation of interface and implementation
- ▶ Same principle from notebooks to services

### Why This Matters

- ▶ Logic becomes reusable and testable
- ▶ Notebooks stay thin and readable
- ▶ Systems can grow without rewrites

### Example 1: Local API (project code)

#### Function API usage

```
1 from model import forecast  
2 y = forecast(x, params)
```

### Example 2: Web API (service)

#### HTTP request

```
1 response = requests.get(  
2     "/weather",  
3     params={"city": "Berlin", "  
        date": "2025-01-20"})  
4 data = response.json()
```

### Key Insight

- ▶ APIs shift complexity behind a stable boundary

## REST Essentials — Tasks as Resources / Agent Framework Elements

### Resources

- ▶ A task is a **resource**
- ▶ Identified by a URL
- ▶ Has a state and metadata
- ▶ `/tasks`
- ▶ `/tasks/<id>`

### HTTP Methods

- ▶ POST — create a task
- ▶ GET — inspect task or list
- ▶ PUT — change task state
- ▶ DELETE — remove a task

### Status Codes (Minimum)

- ▶ 200 OK
- ▶ 201 Created
- ▶ 400 Bad Request
- ▶ 404 Not Found

## Task Server Setup — Minimal Flask Example

### Goal

- ▶ Tasks as REST resources
- ▶ Explicit task state
- ▶ Minimal server logic

### States

- ▶ created
- ▶ checked
- ▶ executing
- ▶ completed / failed

task\_server.py

```
1 from flask import Flask, request, jsonify
2 app = Flask(__name__)
3 tasks = {}; i = 1
4
5 @app.post("/tasks")
6 def create():
7     global i
8     if not request.json: abort(400)
9     t = {"id": i, "state": "created", "data":
10         ": request.json}
11     tasks[i] = t; i += 1
12     return jsonify(t), 201
13 app.run()
```

## Testing the Task REST API

### What We Test

- ▶ Send JSON to the server (POST /tasks)
- ▶ Server creates a new task
- ▶ Server assigns **ID** and initial **state**

### Expected Result

- ▶ HTTP status 201 Created
- ▶ JSON response with:
  - ▶ task id
  - ▶ state = created
  - ▶ echoed input data

#### Test via curl

```
1 curl -X POST http://127.0.0.1:5000/
  tasks -H "Content-Type:
  application/json" -d '{"type":'
  "demo", "params": {"x": 1}}'
```

#### Minimal Python Task Creation

```
1 import requests
2 r = requests.post(
3     "http://127.0.0.1:5000/tasks",
4     json={"type": "demo", "params": {"x":
5         1}})
6 print(r.status_code); print(r.json())
```

```
{'id': 1, 'state': 'created',
'data': {'type': 'demo',
'params': {'x': 1}}}
```

## Task API — Server Endpoints

### Purpose

- ▶ Expose tasks via REST
- ▶ Read-only inspection
- ▶ Server owns all state

### Server Resource Endpoints

- ▶ `/tasks` — task collection
- ▶ `/tasks/<id>` — single task

### Semantics

- ▶ Stateless client
- ▶ Explicit URLs
- ▶ JSON responses

```
task_endpoints.py
```

```
1 from flask import Flask, jsonify,
2     abort
3 app = Flask(__name__)
4 tasks = {}
5
6 @app.get("/tasks")
7 def list_tasks():
8     return jsonify(list(tasks.
9         values()))
10
11 @app.get("/tasks/<int:i>")
12 def get_task(i):
13     return jsonify(tasks[i]) if i
14         in tasks else abort(404)
15
16 app.run()
```

## Querying Tasks — Inspecting Server State

### Goal

- ▶ Inspect existing tasks
- ▶ Read state and metadata
- ▶ No client-side state

### REST Principle

- ▶ Tasks are resources
- ▶ Identified by URL
- ▶ Read via GET

### Endpoints

- ▶ GET /tasks
- ▶ GET /tasks/<id>

#### List all Tasks

```
1 import requests
2 r = requests.get("http
                  ://127.0.0.1:5000/tasks")
3 print(r.status_code)
4 print(r.json())
```

#### Describe one task

```
1 import sys, requests
2 tid = int(sys.argv[1])
3 r = requests.get(f"http
                  ://127.0.0.1:5000/tasks/{tid}")
4 print(r.status_code)
5 print(r.json())
```

## Native Code Integration: Fortran and C++ in Python & ML

### Why Native Code Matters

- ▶ Decades of validated Fortran in NWP
- ▶ High-performance kernels in C++
- ▶ Tight control over memory and execution
- ▶ Reuse of trusted implementations

### Typical Use Cases

- ▶ Physical parameterizations
- ▶ Linear operators, solvers, kernels
- ▶ Observation operators
- ▶ Legacy model components

### Python as the Orchestration Layer

- ▶ Python controls the workflow
- ▶ Native code provides compute kernels
- ▶ Clean separation via APIs

### Integration Options (Overview)

- ▶ ctypes – explicit C-compatible interfaces
- ▶ f2py – automatic Fortran bindings
- ▶ pybind11 – modern C++ bindings
- ▶ Shared libraries: .so / .dylib

## Fortran with C Bindings: A Stable Interface

### Why C Bindings?

- ▶ Fortran and Python do **not** talk directly
- ▶ The common denominator is the **C ABI**
- ▶ Stable, explicit, language-independent

### Key Concept

- ▶ Fortran exposes functions as **C-compatible symbols**
- ▶ No name mangling
- ▶ Well-defined data types

#### Minimal Fortran Example

```
1 function f_sin_cos(x) result(f) bind(C)
2   use iso_c_binding
3   real(c_double), intent(in) :: x
4   real(c_double) :: f
5   f = sin(x) * cos(x)
6 end function
```

#### What This Ensures

- ▶ Symbol name is predictable
- ▶ Argument layout follows C rules
- ▶ Callable from Python, C, C++

## Calling Fortran from Python with ctypes

### Role of Python

- ▶ Python loads the shared library
- ▶ Defines the function signature
- ▶ Manages data exchange

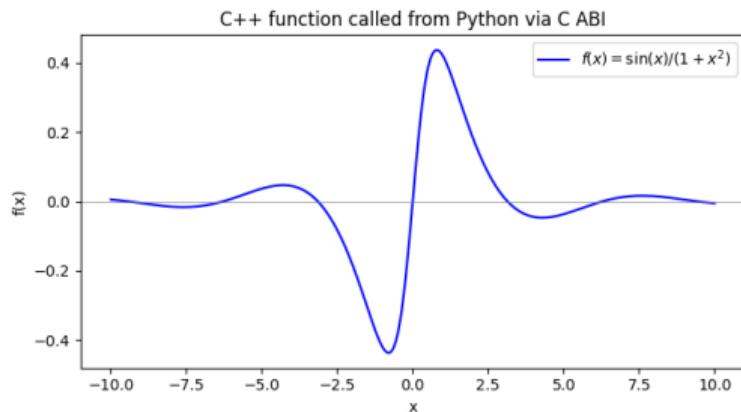
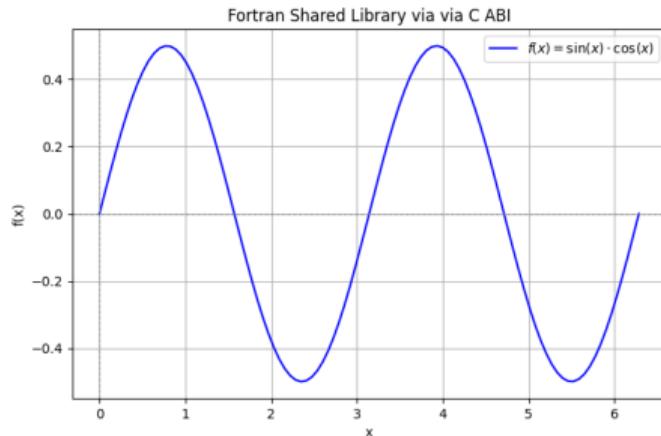
### Important Detail

- ▶ Fortran arguments are passed by reference
- ▶ Python must pass pointers
- ▶ Native Fortran computation
- ▶ Controlled Python interface
- ▶ No performance-critical Python loop

#### Minimal Python Interface

```
1 import ctypes
2
3 lib = ctypes.CDLL("./fortran_interface.so")
4 lib.f_sin_cos.argtypes = [
5     ctypes.POINTER(ctypes.c_double) ]
6 lib.f_sin_cos.restype = ctypes.c_double
7
8 x = ctypes.c_double(1.0)
9 y = lib.f_sin_cos(ctypes.byref(x))
```

## Native Code Integration — Fortran vs. C++ from Python



Fortran via ISO\_C\_BINDING + ctypes  
Numerics compiled, called from Python

C++ via C-compatible ABI + ctypes  
Explicit interface, manual symbol control

- ▶ Python orchestrates, native code computes
- ▶ Identical workflow: compile → load → call → plot

## Exposing a C++ Function via the C ABI

### Role of C++

- ▶ Implements numerical logic
- ▶ Compiled into a shared library
- ▶ Exposes a stable C ABI

### Key Requirement

- ▶ Use `extern "C"`
- ▶ Avoid C++ name mangling

#### Minimal C++ Interface

```
1 #include <cmath>
2
3 extern "C" double f_cpp(double x)
4 {
5     return std::sin(x) / (1.0 + x*x);
6 }
```

- ▶ Plain C-compatible symbol
- ▶ No templates, no classes
- ▶ ABI-safe function signature

## Calling C++ from Python with ctypes

### Role of Python

- ▶ Python loads the shared library
- ▶ Defines the **binary interface**
- ▶ Controls execution and visualization

### Important Detail

- ▶ C++ function uses **C ABI**
- ▶ Arguments passed **by value**
- ▶ Native C++ computation
- ▶ Explicit ABI contract
- ▶ Minimal Python overhead

#### Minimal Python Interface

```
1 import ctypes
2
3 lib = ctypes.CDLL("./cpp_interface.so")
4 lib.f_cpp.argtypes = [ctypes.c_double]
5 lib.f_cpp.restype = ctypes.c_double
6
7 x = 1.0
8 y = lib.f_cpp(x)
```

## Remote Jupyter — Principle

### Core Idea

- ▶ Compute: **remote** (server, HPC, cloud)
- ▶ Visualization: **local** in the browser
- ▶ Browser setup remains unchanged

### Typical Scenarios

- ▶ HPC login or compute nodes
- ▶ Cloud virtual machines
- ▶ Office workstation / bastion host

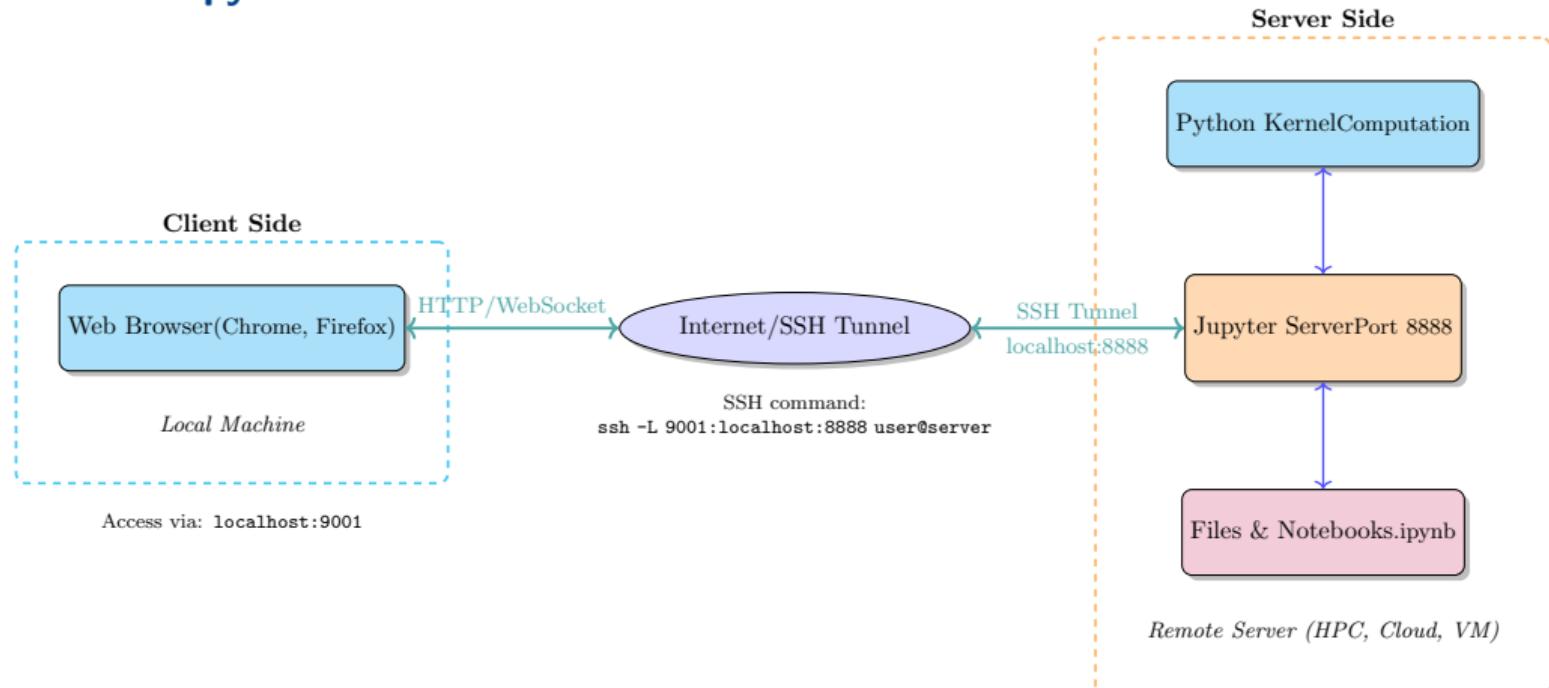
### Security Principle

- ▶ **No** open notebook ports to the network
- ▶ Never expose :8888 on a public IP
- ▶ Access exclusively via **SSH tunnels**

### Key Takeaway

- ▶ Jupyter only listens **locally**
- ▶ SSH provides secure transport

## Remote Jupyter — Network Architecture



## Remote Jupyter — Port Forwarding Recipe

### Step 1: Start Jupyter Remotely

#### Remote (Linux)

```
1 jupyter notebook \
2   --no-browser \
3   --port=8888
```

### Step 2: Create SSH Tunnel

#### Local (Bash / PowerShell)

```
1 ssh -N -L 9001:localhost:8888 \
2       user@remote-host
```

### Step 3: Open in Browser

- ▶ Access token appears in the terminal
- ▶ Port is **local only** on the remote server

- ▶ <http://localhost:9001>
- ▶ Use token from the remote log
- ▶ **Adjust port** if needed

To access the server, open this file in a browser:

<file:///hpc/uhome/rpotthas/.local/share/jupyter/runtime/jpserver-1662241-open.html>

Or copy and paste one of these URLs:

<http://localhost:8888/tree?token=369b469c83e8eaa369ee02ea443d994865ca93e4185bb385>

<https://127.0.0.1:8888/tree?token=369b469c83e8eaa369ee02ea443d994865ca93e4185bb385>

## Remote Jupyter — Firewalls & Troubleshooting

### Typical Problems

- ▶ Local port already in use
- ▶ Remote port 8888 blocked by firewall
- ▶ Unstable SSH connection
- ▶ Multi-hop access (Jump / Bastion host)

### Important Note

- ▶ A blocked 8888 is **not an error**
- ▶ SSH tunnels do not require open inbound ports

### Best Practices

- ▶ Remote:
  - ▶ -ip=127.0.0.1
  - ▶ no public binding
- ▶ Local:
  - ▶ choose a free port (9001, 9002, ...)
- ▶ Infrastructure:
  - ▶ Bastion / jump host via -J

### Recommendation

- ▶ **Never** expose Jupyter notebook ports publicly