

Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀
playing 🤖 with
Python 🐍 and AI/ML!
 

Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	Lab: Python Startup: Basics	Lab: Feed-forward and Graph NNs	Lab: Graph Learning with PyTorch	Lab: Containers and Reproducibility	Lab: Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	Lab: Work environments, Python everywhere	Lab: Simple Transformer and LLM Use	Lab: Agent Frameworks	Lab: CI/CD Pipelines	Lab: Graph-based Assimilation
12:45–13:30	Lunch Break				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	Lab: GRIB, NetCDF and Obs Visualisation	Lab: RAG Pipeline	Lab: DAWID Exploration	Lab: Anemoi Training Pipeline	Lab: Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	The AI Transformation	Learning from Observations Only
16:15–17:00	Lab: Torch Tensors and First Neural Net	Lab: Radar, SAT and Multimodal Data	Lab: MLflow Hands-on	Lab: How work style could change	Lab: ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

Working with Meteorological Data for AI and ML

Context

- ▶ Python-based workflows
- ▶ Meteorological fields for AI and machine learning
- ▶ Real operational datasets

Why this lecture?

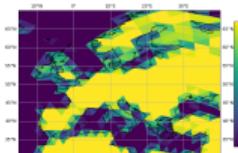
- ▶ ML needs structured input data
- ▶ Models operate on fields, not files
- ▶ Visualization builds physical intuition

Core skills we build

- ▶ Access meteorological fields
- ▶ Understand grids and metadata
- ▶ Convert and preprocess data
- ▶ Visualize fields and observations

Outcome

- ▶ Data ready for ML pipelines
- ▶ Reproducible Python workflows



Installing and Exploring ecCodes

What is ecCodes?

- ▶ ECMWF library for GRIB handling
- ▶ Standard tool in operational NWP
- ▶ Python interface available

Why we need it

- ▶ Forecast data comes as GRIB
- ▶ We need access to fields and metadata
- ▶ Independent of model (IFS, ICON, ...)

Installation and test

```
1 # Linux / WSL
2 sudo apt update
3 sudo apt install eccodes libeccodes-
   tools
4
5 # macOS (Homebrew)
6 brew install eccodes
7
8 # Python bindings
9 pip install eccodes
10
11 # Test installation
12 grib_ls -V
13 python -c "import eccodes; print(
   eccodes.codes_get_api_version())"
```

Using ecCodes

- ▶ See what is inside a GRIB file
- ▶ Identify fields
- ▶ Get size and structure

Key idea

- ▶ A GRIB file contains multiple messages
- ▶ Each message holds one field
- ▶ ecCodes iterates message by message

Inspecting a GRIB file

```
1 import eccodes
2
3 f = open("icon_eu_t2m_latest.grib2", "rb")
4 while True:
5     gid = eccodes.codes_grib_new_from_file(f)
6     if gid is None:
7         break
8
9     short = eccodes.codes_get(gid, "shortName")
10    level = eccodes.codes_get(gid, "level")
11    size = eccodes.codes_get_size(gid, "values")
12    print(short, level, size)
13    eccodes.codes_release(gid)
14 f.close()
```

IFS Open Data — Access via Client Library

- ▶ IFS data is not exposed as simple directories
- ▶ Access is service-based and load-controlled
- ▶ Small subsets can be retrieved efficiently

Key idea

- ▶ Use ECMWF's official open-data client
- ▶ Request fields by parameter, step, level
- ▶ Result is a standard GRIB2 file

Downloading IFS fields with ecmwf-opendata

```
1 from ecmwf.opendata import Client
2
3 client = Client(
4     source="ecmwf",
5     model="ifs",
6 )
7
8 client.retrieve(
9     time=0,
10    type="fc",
11    step=24,
12    param=["2t", "msl"],
13    target="ifs_2t.grib2")
```

Plotting an IFS Field (Regular Grid)

- ▶ IFS open data uses a regular lat–lon grid
- ▶ Field values form a 2-D NumPy array
- ▶ Plotting is straightforward

Important

- ▶ Grid size is fixed by the model
- ▶ figsize controls readability only

Plotting IFS 2 m temperature

```
1 import eccodes
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 f = open("ifs_2t.grib2", "rb")
6 gid = eccodes.codes_grib_new_from_file(f)
7
8 nx = eccodes.codes_get(gid, "Ni")
9 ny = eccodes.codes_get(gid, "Nj")
10 values = eccodes.codes_get_array(gid, "values")
11 field = values.reshape(ny, nx)
12
13 plt.figure(figsize=(7, 3.5))
14 plt.imshow(field, origin="lower")
```

Cartopy — Map Projections for Meteorological Fields

- ▶ Cartopy provides projection-aware plotting
- ▶ Separates data coordinates from map projection
- ▶ Built on top of matplotlib

In practice

- ▶ Define projection, e.g. PlateCarree
- ▶ Coastlines, borders, land masks
- ▶ Define geographic extent

Cartopy-based map plot

```
1 import cartopy.crs as ccrs
2 import cartopy.feature as cfeature
3
4 fig, ax = plt.subplots( figsize=(10,5),
5                         subplot_kw={"projection": ccrs.
6                           PlateCarree()})
7
8 ax.coastlines()
9 ax.add_feature(cfeature.BORDERS)
10
11 ax.pcolormesh(lon, lat, field,
12                 transform=ccrs.PlateCarree(),
13                 cmap="jet" )
```

ICON Open Data — File-Based Access

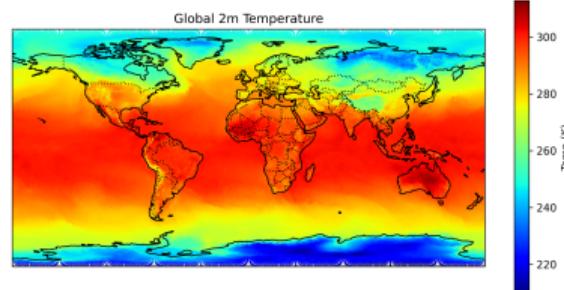
- ▶ ICON data is published as static GRIB files
- ▶ Organised by model, run, and parameter
- ▶ Distributed via DWD Open Data server

Typical workflow

- ▶ Query directory listing
- ▶ Identify latest available timestamp
- ▶ Download and decompress .grib2.bz2

Material provided

- ▶ Python script: icon_download_t2m.py
- ▶ Jupyter notebook: ICON download & inspection



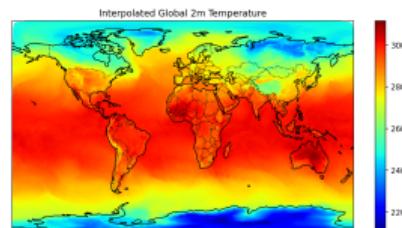
See code and notebooks for a robust, reusable implementation.

ICON Fields — Interpolation and Dateline Handling

- ▶ ICON uses an unstructured triangular grid
 - ▶ Direct plotting requires triangulation
 - ▶ Interpolation needed for:
 - ▶ smooth visualization
 - ▶ map projections
-
- ### Implementation
- ▶ See Python script:
`fix_dateline_triangles.py`
(adapted from ICON tutorial)

Key challenge

- ▶ Triangles crossing the date line
- ▶ Incorrect rendering in global projections
- ▶ Requires explicit geometric handling



Observational Data — Purpose and Characteristics

- ▶ Observations anchor models to reality
- ▶ Essential for:
 - ▶ verification
 - ▶ data assimilation
 - ▶ ML training and validation

Typical properties

- ▶ Irregular spatial distribution
- ▶ Sparse and heterogeneous
- ▶ Strong metadata dependence

Common observation types

- ▶ SYNOP (surface stations)
- ▶ TEMP / radiosondes
- ▶ AIREP / AMDAR (aircraft)
- ▶ Satellite observations

Key difference to models

- ▶ No grid
- ▶ Point-based measurements
- ▶ Representativeness matters

SYNOP Data — Reading Raw Observations

- ▶ SYNOP stored in NetCDF
- ▶ One record per station
- ▶ Variables include:
 - ▶ latitude, longitude
 - ▶ observed temperature

Minimal requirement

- ▶ Read coordinates
- ▶ Read physical quantity
- ▶ Preserve missing values

Python tools

- ▶ netCDF4.Dataset
- ▶ numpy

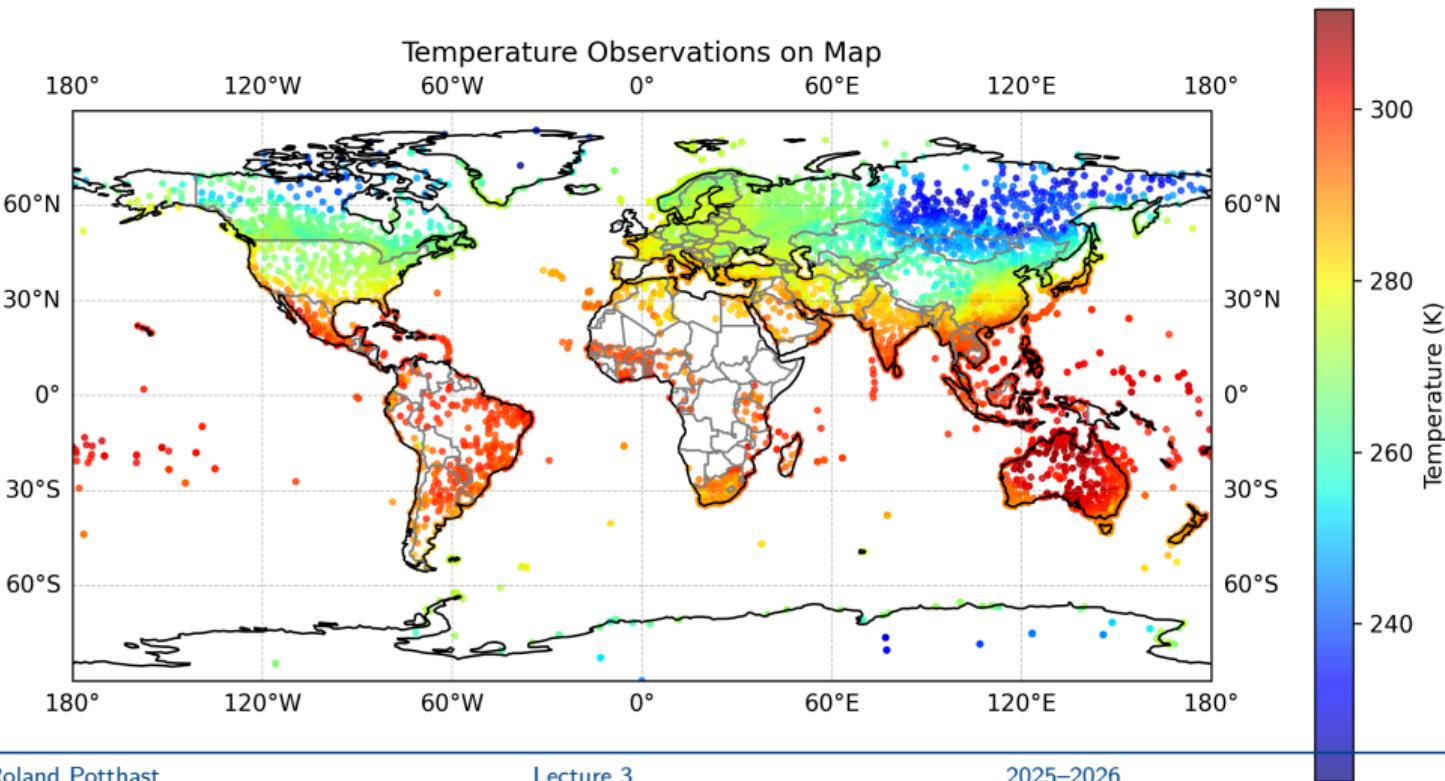
Variables used

- ▶ MLAH — latitude
- ▶ MLOH — longitude
- ▶ MTDBT — temperature

Implementation

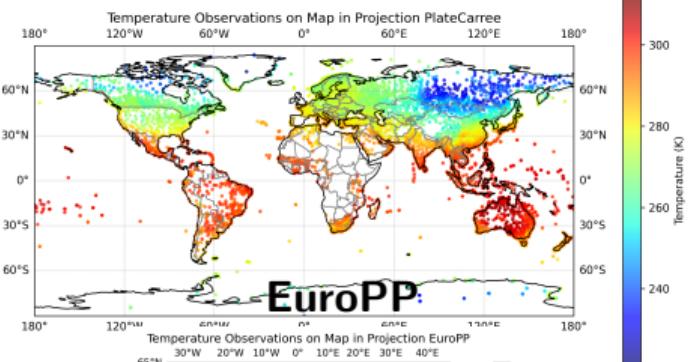
- ▶ See function: `read_synop_data()`

SYNOP Observations — Point-Based Visualization

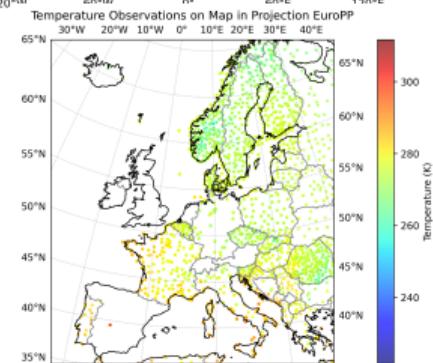


SYNOP Observations — Map Projections

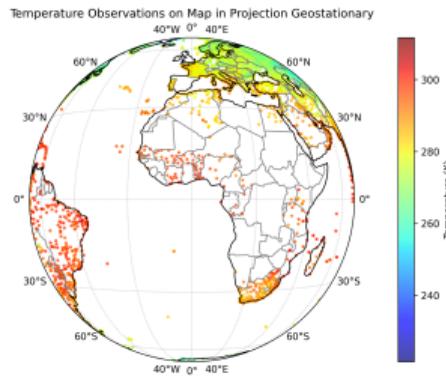
Plate Carree



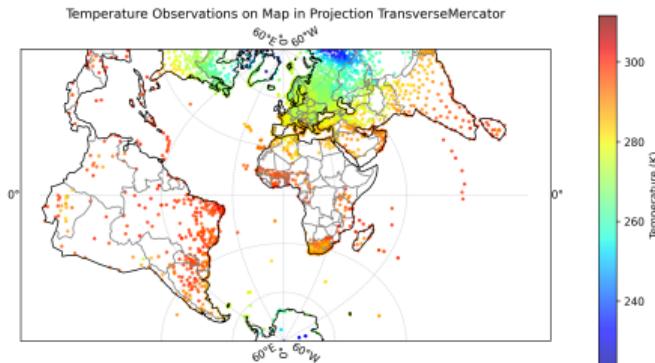
EuroPP



Geostationary



Transverse Mercator



Feedback Files — Linking Models and Observations

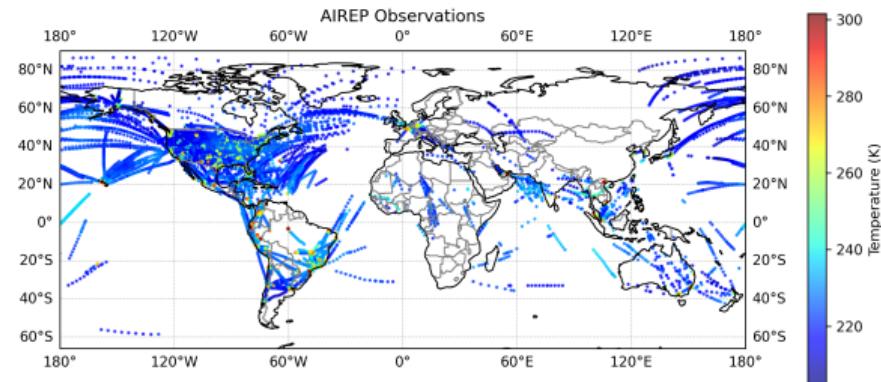
- ▶ Feedback files store model–observation pairs
- ▶ Generated during data assimilation
- ▶ Central diagnostic product

Why they matter

- ▶ Quantify model bias and error
- ▶ Basis for monitoring systems
- ▶ Ideal input for ML analysis

Contained information

- ▶ Observation value
- ▶ Model equivalent ($H(x)$)
- ▶ Quality control flags
- ▶ Metadata (time, location, type)



Feedback Files — Structure and Indexing

- ▶ Stored as structured NetCDF files
- ▶ Separate layers for:
 - ▶ report metadata
 - ▶ individual observations

Key idea

- ▶ One report may contain many observations
- ▶ Explicit indexing links both layers

Core variables (excerpt)

i_body	(nreport)	start index of report
l_body	(nreport)	number in report
obs	(nobs)	observation
veri_data	(nmodel,nobs)	model equivalent $H(x)$

Interpretation

- ▶ `i_body:l_body` maps reports → `obs`
- ▶ `obs` and `veri_data` are aligned
- ▶ Enables direct O-B, O-A diagnostics

AIREP Feedback — Statistical Diagnostics

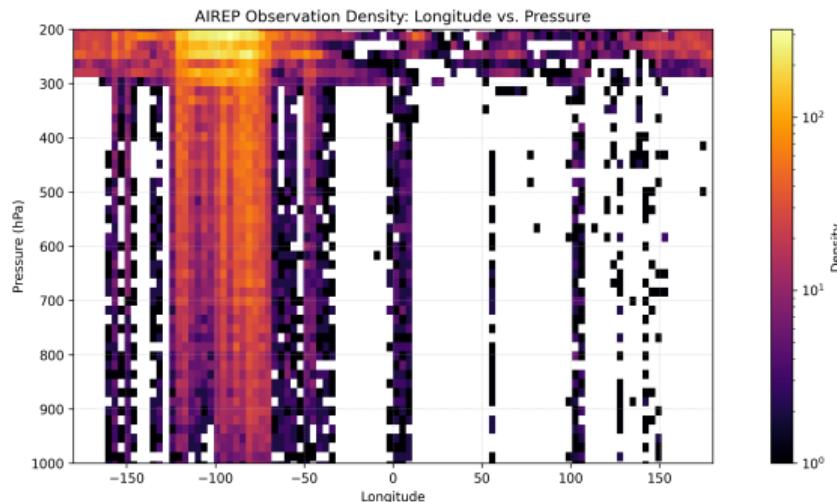
- ▶ AIREP / AMDAR provide aircraft observations
- ▶ Strong vertical and regional sampling biases
- ▶ Feedback files enable systematic evaluation

Typical diagnostics

- ▶ Observation density vs height
- ▶ Mean and spread of innovations (O–B)
- ▶ Model-dependent biases

Why this matters

- ▶ Identifies representativeness errors
- ▶ Reveals flow- and height-dependent biases

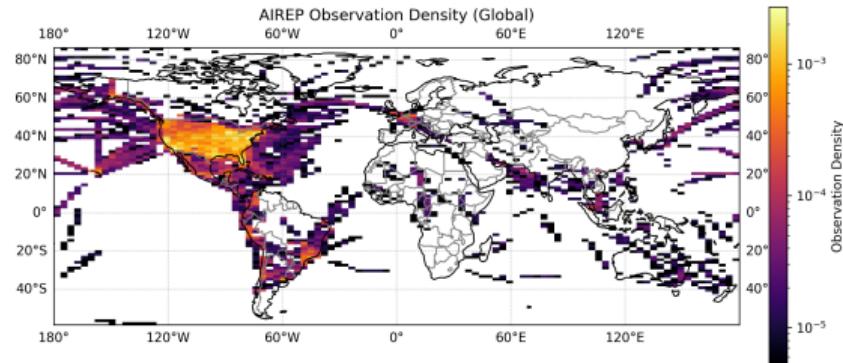


AIREP Feedback — Horizontal Sampling

- ▶ AIREP sampling follows air traffic routes
- ▶ Highly inhomogeneous horizontal coverage
- ▶ Strong land–sea and regional contrasts

Key implications

- ▶ Data density reflects logistics, not physics
- ▶ Large unsampled regions persist
- ▶ Bias risk if learning ignores sampling



Global horizontal density of AIREP observations

Operational relevance

- ▶ Affects verification statistics
- ▶ Affects DA weighting
- ▶ Critical for ML generalisation

Why GRIB? — Operational Requirements

- ▶ GRIB = GRIdded Binary
- ▶ Designed for operational NWP
- ▶ Optimised for:
 - ▶ large data volumes
 - ▶ fast I/O
 - ▶ network distribution

Core design goals

- ▶ Compact encoding
- ▶ Self-describing metadata
- ▶ Bit-level compression

Why it is still used

- ▶ Global models produce terabytes per day
- ▶ Multiple grids and vertical coordinates
- ▶ Strict real-time constraints

Operational reality

- ▶ WMO standard
- ▶ Tooling is mature and stable
- ▶ Backward compatibility matters

Alternatives to GRIB — Strengths and Limits

NetCDF / CF

- ▶ Widely used in research
- ▶ Human-readable metadata
- ▶ Excellent tool support

Zarr / Cloud-native

- ▶ Chunked, parallel access
- ▶ Designed for object storage
- ▶ Attractive for ML workflows

Why GRIB is not easily replaced

- ▶ Operational ecosystems are huge
- ▶ Encoding efficiency still unmatched
- ▶ Standards evolution is slow by design

Current practice

- ▶ GRIB for production and exchange
- ▶ NetCDF / Zarr for analysis and ML
- ▶ Conversion layers in between

Key message

- ▶ Format choice reflects use case

GPU Access in Practice — Why It Depends on the Platform

- ▶ Most laptops *do have a GPU*
- ▶ But GPU **type and backend** differ strongly
- ▶ Python code must match the backend

Common GPU types

- ▶ **NVIDIA** — typical for HPC clusters
- ▶ **Apple GPU** — Apple Silicon
- ▶ **AMD / Radeon** — some laptops, workstations

Key message

- ▶ The **software interface** matters

GPU backends used in Python

- ▶ CUDA — NVIDIA GPUs (Linux, HPC)
- ▶ MPS / Metal — Apple GPUs (macOS)
- ▶ CPU fallback — always available

Implication for ML workflows

- ▶ Development often on laptops
- ▶ Training often on HPC systems
- ▶ Code must be **portable across backends**

GPU Tests in Python — Minimal Backend Checks

NVIDIA GPU (CUDA)

CUDA test

```
1 import torch
2 print(torch.cuda.is_available())
```

Apple GPU (Metal / MPS)

Apple GPU test

```
1 import torch
2 print(torch.backends.mps.
      is_available())
```

AMD / Radeon GPUs

- ▶ No native PyTorch GPU backend

Interpretation

- ▶ True → GPU usable
- ▶ False → CPU fallback
- ▶ Backend determines capability

GPU Execution in Python — Laptop vs. HPC

Apple GPU (Metal / MPS)

Apple GPU test

```
1 import torch, time
2
3 d = torch.device("mps")
4 x = torch.rand((4000,4000),
     device=d)
5
6 t0 = time.time()
7 y = torch.matmul(x, x)
8 torch.mps.synchronize()
9
10 print("Apple GPU time:",
11      round(time.time()-t0,3))
```

NVIDIA A100 (CUDA)

HPC GPU test

```
1 import torch, time
2
3 d = torch.device("cuda")
4 x = torch.rand((8000,8000),
     device=d)
5
6 t0 = time.time()
7 y = torch.matmul(x, x)
8 torch.cuda.synchronize()
9
10 print("A100 time:",
11      round(time.time()-t0,3))
```

Interactive GPU Access and Verification Workflow

Step 1: Interactive GPU login

Shell alias (interactive GPU)

```
1 alias gpu1='qlogin -q
    gp_inter_dgx \
2   --gpnum-lhost=1 \
3   --cpnum-lhost=16 \
4   -l elapstim_req=6:00:00 \
5   -l memsz_job=240gb'
```

Step 2: Select GPU explicitly

Shell

```
1 export CUDA_VISIBLE_DEVICES=7
2 echo $CUDA_VISIBLE_DEVICES
```

Step 3: Minimal verification

Mini GPU check in Python

```
1 import torch, os
2 print("CUDA_VISIBLE_DEVICES ="
      , os.getenv("CUDA_VISIBLE_DEVICES"))
3 print("Visible GPUs =", torch.
      cuda.device_count())
4 print("GPU name =", torch.cuda.
      get_device_name(0))
```

- ▶ Interactive login allocates resources
- ▶ GPU visibility must be set **explicitly**
- ▶ Always verify from inside Python

CPU vs GPU — Execution Model and Performance: 30sec → 3sec

CPU example

```
1 import torch, time
2 n = 30000
3 x0 = torch.rand((n,n))
4 x1 = torch.rand((n,n))
5 t0 = time.time()
6 y0 = torch.matmul(x0, x0)
7 y1 = torch.matmul(x1, x1)
8 print(time.time()-t0)
```

Two-GPU example

```
1 import torch, time
2 d0 = torch.device("cuda:0")
3 d1 = torch.device("cuda:1")
4 n. = 30000
5 x0 = torch.rand((n,n), device=d0)
6 x1 = torch.rand((n,n), device=d1)
7 t0 = time.time()
8 torch.matmul(x0, x0)
9 torch.matmul(x1, x1)
10 torch.cuda.synchronize()
11 print(time.time()-t0)
```

Observed

- ▶ Operations are **blocking**
- ▶ Second call starts after first ends
- ▶ Parallelism only inside BLAS
- ▶ Operations run **concurrently**

Parallel Matrix Multiplication — What Actually Works

Key idea

- ▶ Naive multi-GPU often shows no speed-up
- ▶ Reason: data transfers dominate computation
- ▶ Solution: true model parallelism

What works

- ▶ Build matrices directly on each GPU
- ▶ Split the computation (block-wise)
- ▶ Avoid CPU–GPU transfers
- ▶ Collect results only at the end

True model-parallel matrix multiplication

```
1 import torch, time
2 n = 30000
3 A0 = torch.rand((n//2,n), device=
                 "cuda:0")
4 A1 = torch.rand((n//2,n), device=
                 "cuda:1")
5 B = torch.rand((n,n),           device
                 ="cuda:0")
6 t0 = time.time()
7 C0 = A0 @ B
8 C1 = A1 @ B.to("cuda:1")
9 torch.cuda.synchronize()
10 print("Time:", round(time.time()-
                         t0,3))
```

Mixed Precision Computing — Why FP16 Matters

Why reduced precision?

- ▶ Modern GPUs are optimized for FP16 / BF16
- ▶ Higher throughput, lower memory traffic
- ▶ Essential for large ML models

Typical ML pattern

- ▶ Dense linear layers
- ▶ Nonlinear activation
- ▶ Large batch sizes

FP16 neural-network-style workload

```
1 import torch, time
2 torch.set_default_dtype(torch.float16)
3 d = torch.device("cuda")
4 x = torch.randn((20000, 1024), device=d)
5 W1 = torch.randn((1024, 4096), device=d)
6 W2 = torch.randn((4096, 1024), device=d)
7 t0 = time.time()
8 y = torch.nn.functional.gelu(x @ W1)
9 z = y @ W2
10 torch.cuda.synchronize()
11 print("FP16 time:", round(time.time() - t0, 3))
```