

# Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀  
playing 🤖 with  
Python 🐍 and AI/ML!  
 

## Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	<b>Lab:</b> Python Startup: Basics	<b>Lab:</b> Feed-forward and Graph NNs	<b>Lab:</b> Graph Learning with PyTorch	<b>Lab:</b> Containers and Reproducibility	<b>Lab:</b> Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	<b>Lab:</b> Work environments, Python everywhere	<b>Lab:</b> Simple Transformer and LLM Use	<b>Lab:</b> Agent Frameworks	<b>Lab:</b> CI/CD Pipelines	<b>Lab:</b> Graph-based Assimilation
12:45–13:30	<b>Lunch Break</b>				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	<b>Lab:</b> GRIB, NetCDF and Obs Visualisation	<b>Lab:</b> RAG Pipeline	<b>Lab:</b> DAWID Exploration	<b>Lab:</b> Anemoi Training Pipeline	<b>Lab:</b> Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	<b>The AI Transformation</b>	Learning from Observations Only
16:15–17:00	<b>Lab:</b> Torch Tensors and First Neural Net	<b>Lab:</b> Radar, SAT and Multimodal Data	<b>Lab:</b> MLflow Hands-on	<b>Lab:</b> How work style could change	<b>Lab:</b> ORIGEN and Open Discussion
17:00–20:00	Joint Dinner				

## AI and ML — A Problem-Solving Perspective

### Classical approach

- ▶ Explicit equations
- ▶ Physical laws and dynamics
- ▶ Expert-designed structure
- ▶ Limited by model assumptions

$$\partial_t x = F(x, \theta)$$

The temporal change of  $x$  is calculated based on  $x$  and parameters  $\theta$ .

### AI / ML approach

- ▶ Learn mappings  $x \rightarrow y$  from data
- ▶ High-dimensional, nonlinear relations
- ▶ Neural nets as universal approximators
- ▶ Used in natural sciences as well as language

$$\hat{z} = f_\theta(x)$$

Some quantity  $z$  is estimated from input  $x$

## AI and ML — A Set of Tools

### Core ML frameworks

- ▶ PyTorch, TensorFlow
- ▶ scikit-learn
- ▶ Automatic differentiation
- ▶ GPU acceleration

$$\min_{\theta} L(y, \hat{y}(x; \theta))$$

**Training** = minimizing a loss function by adjusting parameters  $\theta$ .

### AI as a service

- ▶ **LLM APIs** (OpenAI, Mistral, Anthropic, Google, Meta, ...)
- ▶ Pre-trained foundation models
- ▶ On-premise models (Llama, Mistral)
- ▶ Cloud, local, or hybrid use

API call → model inference

Models are used without training from scratch.

**What is my own role in this?**

## AI and ML — A New Paradigm for Interactivity

### Human–AI interaction

- ▶ Code assistants
- ▶ Natural language interfaces
- ▶ Interactive problem solving
- ▶ Rapid prototyping

### AI in research workflows

- ▶ Data exploration
- ▶ Hypothesis support
- ▶ Equation generation and review
- ▶ Support Reasoning

Prompt → Response

**Iterative dialogue** replaces  
static interfaces.

Human ↔ AI

Collaboration , not replacement.

**AI as Partner for Reasoning.**

## Critical Evaluation I — Reliability and Limits

### Strengths

- ▶ Fast pattern recognition
- ▶ Handles high-dimensional data
- ▶ Automates repetitive tasks
- ▶ Strong empirical performance

### Limitations

- ▶ No physical understanding
- ▶ Hallucinations possible
- ▶ Sensitive to data bias
- ▶ Weak extrapolation

$$\hat{y} = f_{\theta}(x)$$

Works well within  
the learned data regime.

$$f_{\theta}(x) \neq \text{truth}$$

Prediction is not validation.

## Critical Evaluation II — Trust, Oversight, Responsibility

### Why human oversight matters

- ▶ AI outputs look convincing
- ▶ Errors are often non-obvious
- ▶ No built-in notion of consequences
- ▶ Responsibility remains human

### Key risk dimensions

- ▶ Transparency and explainability
- ▶ Bias and unfairness
- ▶ Reproducibility
- ▶ Accountability

Decision = AI + Human

AI supports, it does not decide.

Confidence  $\neq$  Correctness

Trust must be earned, not assumed.

**Be careful, AI makes many mistakes!**

## Torch Tensors — The Core Data Structure

### What is a tensor?

- ▶ Similar to NumPy arrays
- ▶ Supports CPU and GPU
- ▶ Tracks operations for gradients
- ▶ Basis of all learning

### Key properties

- ▶ Shape and dtype
- ▶ Device awareness
- ▶ `requires_grad=True`

#### Basic tensor example

```
1 import torch
2
3 x = torch.tensor([2.,3.],
4                  requires_grad=True)
5
6 y = x[0]**2 + x[1]**2
7 y.backward()
8 print(x.grad)
```

$$\nabla_x (x_1^2 + x_2^2)$$

## Automatic Differentiation (Autograd)

### Core idea

- ▶ Gradients computed automatically
- ▶ No manual derivative formulas
- ▶ Works for arbitrary computation graphs
- ▶ Enabled by dynamic graphs

$$\frac{\partial \mathcal{L}}{\partial \theta}$$

Gradients drive parameter updates.

### Why this matters

- ▶ Learning = optimization
- ▶ Backpropagation at scale
- ▶ Essential for deep networks
- ▶ Same mechanism on CPU and GPU

$$\theta_{k+1} = \theta_k - \eta \nabla_{\theta} \mathcal{L}$$

Gradient-based learning step.

## Data Handling in PyTorch: Dataset and DataLoader

### Why data loaders exist

- ▶ Datasets often too large for memory
- ▶ Training uses **mini-batches**
- ▶ Data order matters for optimization
- ▶ Separation of data and model logic

$$(x_i, y_i) \rightarrow (X_B, Y_B)$$

Samples are grouped into batches.

### What DataLoader provides

- ▶ Batching
- ▶ Optional shuffling
- ▶ Parallel loading (CPU workers)
- ▶ Consistent interface for training loops

$$(X_B, Y_B) \rightarrow \mathcal{L}(f_\theta(X_B), Y_B)$$

Each batch produces one loss value.

## Batches Explained: What Comes Out of the DataLoader

### Features and labels

- ▶ Features  $x$ : input quantities
- ▶ Labels  $y$ : measured target values
- ▶ Learning means fitting  $x \rightarrow y$

### Structure of the data

- ▶  $N$  = number of samples
- ▶  $d$  = number of features per sample
- ▶ One row = one  $(x, y)$  pair

$$X \in \mathbb{R}^{N \times d} \Rightarrow X_B \in \mathbb{R}^{B \times d}$$

### Labels and targets

- ▶ Labels collected in  $Y$
  - ▶ One target per input sample
  - ▶ Same batching as for features
- $$Y \in \mathbb{R}^{N \times k} \Rightarrow Y_B \in \mathbb{R}^{B \times k}$$

### Why mini-batches help

- ▶ Memory-efficient processing
- ▶ Faster parameter updates
- ▶ Noise improves generalization

$$\nabla_{\theta} \mathcal{L}(X_B, Y_B) \approx \nabla_{\theta} \mathcal{L}(X, Y)$$

## Defining a Simple Neural Network

### Neural network idea

- ▶ Learn a mapping  $f_\theta : x \rightarrow \hat{y}$
- ▶ Parameters  $\theta$  are trainable
- ▶ Composition of simple operations

### Basic building blocks

- ▶ Linear transformation
- ▶ Nonlinear activation
- ▶ Output layer

$$\hat{y} = f_\theta(x)$$

#### Minimal PyTorch model

```
1 import torch.nn as nn
2
3 class SimpleNN(nn.Module):
4     def __init__(self):
5         super().__init__()
6         self.fc1 = nn.Linear(1,16)
7         self.relu = nn.ReLU()
8         self.fc2 = nn.Linear(16,1)
9
10    def forward(self,x):
11        x = self.fc1(x)
12        x = self.relu(x)
13        return self.fc2(x)
```

## What the Model Actually Represents

### Model as a function

- ▶ Neural network defines a parametric function
- ▶ Parameters = weights and biases
- ▶ Training adjusts these parameters

$$\hat{y} = W_2 \sigma(W_1 x + b_1) + b_2$$

Nonlinearity  $\sigma$  enables complex mappings.

### Trainable parameters

- ▶ Layer 1:  $1 \rightarrow 16$
- ▶ Layer 2:  $16 \rightarrow 1$

fc1: 16 weights + 16 biases = 32

fc2: 16 weights + 1 bias = 17

**Total:**  $32 + 17 = 49$  parameters

## Loss Functions — Measuring Error

### What is a loss?

- ▶ Quantifies model error
- ▶ Single scalar value
- ▶ Compares prediction vs label

The loss defines the objective that learning tries to minimize .

### Example: Mean Squared Error

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- ▶  $y$ : true label
- ▶  $\hat{y}$ : model prediction
- ▶  $N$ : number of samples

Lower loss ⇒ better fit .

## The Adam Optimizer — Adaptive Gradient Descent

### Why optimizers matter

- ▶ Loss defines what to minimize
- ▶ Optimizer defines how
- ▶ Controls stability and speed

Goal: update parameters to reduce loss  
efficiently.

Adam combines momentum and scaling:

$$\delta\theta \propto -\frac{\text{average current gradient}}{\text{typical gradient size}}$$

### Adam in a nutshell

- ▶ Uses gradients
- ▶ Tracks first moment (mean)
- ▶ Tracks second moment (variance)
- ▶ Adaptive step size per parameter

Parameter update (conceptually):

$$\theta_{t+1} = \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

## End-to-End Example — Learning a Sine Function

### Goal of the example

- ▶ Approximate a known function:  
 $\sin(x)$
- ▶ Learn from sampled input–output pairs
- ▶ Demonstrate full ML workflow

Target mapping:

$$x \mapsto \sin(x)$$

### Components involved

- ▶ Synthetic dataset  $(x, y)$
- ▶ Neural network model
- ▶ Loss function (error measure)
- ▶ Optimizer updating parameters

Training objective:

$$\min_{\theta} \sum_i \|f_{\theta}(x_i) - y_i\|^2$$

## Sine Example — Data and DataLoader

### Dataset construction

- ▶ Sample input values  $x$
- ▶ Compute labels  $y = \sin(x)$
- ▶ Supervised learning setup

Each sample:

$$x_i \rightarrow y_i$$

#### Creating dataset and loader

```
1 x = np.linspace(0, 2*np.pi, 1000)
2 y = np.sin(x)
3
4 x_t = torch.tensor(x).float().
    unsqueeze(1)
5 y_t = torch.tensor(y).float().
    unsqueeze(1)
6
7 data = TensorDataset(x_t, y_t)
8 loader = DataLoader(data,
9                         batch_size=32,
10                        shuffle=True)
```

## Model and Training Loop

### Model idea

- ▶ Input: scalar  $x$
- ▶ Output: scalar  $\hat{y}$
- ▶ Learn nonlinear mapping

### Training

- ▶ Compare  $\hat{y}$  and  $y$
- ▶ Minimize prediction error
- ▶ Update model parameters

### Model and training loop

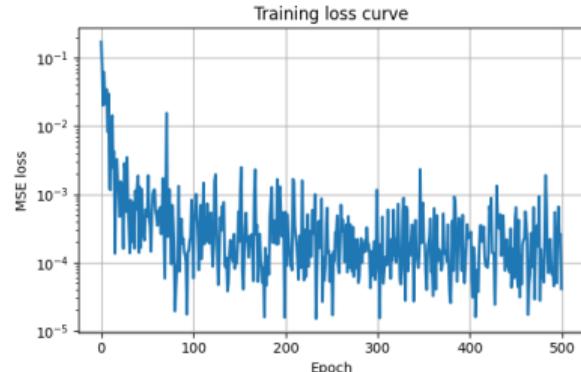
```
1 model = nn.Sequential(  
2     nn.Linear(1,16), nn.ReLU(),  
3     nn.Linear(16,16), nn.ReLU(),  
4     nn.Linear(16,1)  
5 )  
6  
7 loss_fn = nn.MSELoss()  
8 opt = torch.optim.Adam(  
9     model.parameters(), lr=0.01)  
10  
11 for x_b,y_b in loader:  
12     opt.zero_grad()  
13     y_p = model(x_b)  
14     loss = loss_fn(y_p, y_b)  
15     loss.backward()  
16     opt.step()
```

## Sine Example — Training Outcome

### What happened during training?

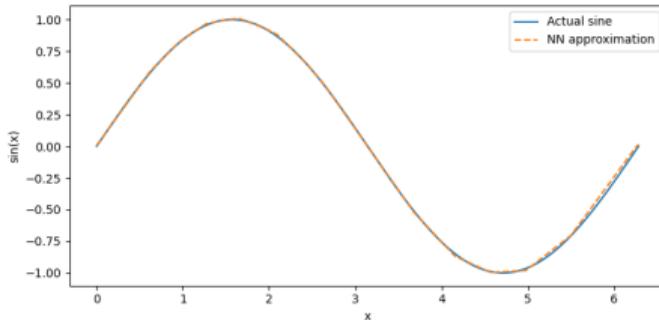
- ▶ Loss decreases over epochs
- ▶ Network parameters adapt
- ▶ Mapping  $x \rightarrow \hat{y}$  improves

Loss  
during  
training



### Interpretation

- ▶ Model learned a smooth function
- ▶ No explicit sine formula given
- ▶ Learning via gradient descent



Actual sine (solid) vs. model prediction  
(dashed)

## DataLoader Without Shuffling

### What happens?

- ▶ Samples returned in fixed order
- ▶ Batches follow dataset sequence
- ▶ Same batches every epoch

### Why this can be problematic

- ▶ Correlated samples in one batch
- ▶ Biased gradient estimates
- ▶ Slower or unstable learning

#### Ordered batches

```
1 loader = DataLoader(  
2     dataset,  
3     batch_size=4,  
4     shuffle=False  
5 )  
6  
7 for x_b, y_b in loader:  
8     print(y_b.squeeze())  
9     break
```

First batch always contains the same labels.

## DataLoader With Shuffling

### What changes?

- ▶ Samples randomly permuted
- ▶ Different batches every epoch
- ▶ Decorrelated gradients

### Why this helps

- ▶ More robust optimization
- ▶ Better generalization
- ▶ Standard practice in ML

==== DataLoader WITHOUT shuffling ===

Batch 1: [1, 2, 3, 4, 5, 6]

Batch 2: [7, 8, 9, 10, 11, 12]

Batch 3: [13, 14, 15, 16, 17, 18]

Batch 4: [19, 20]

### Shuffled batches

```
1 loader = DataLoader(  
2     dataset,  
3     batch_size=4,  
4     shuffle=True  
5 )  
6  
7 for x_b, y_b in loader:  
8     print(y_b.squeeze())  
9     break
```

First batch changes every run.

==== DataLoader WITH shuffling ===

Batch 1: [3, 4, 15, 7, 18, 6]

Batch 2: [11, 16, 13, 17, 9, 1]

Batch 3: [14, 19, 12, 5, 8, 2]

Batch 4: [20, 10]

## From Prediction to Understanding

### Prediction is not enough

- ▶ Low error  $\neq$  understanding
- ▶ Correct output may hide fragile behavior
- ▶ Especially risky outside training range

Models can be accurate yet **misleading**.

### Why gradients matter

- ▶ Sensitivity of output to input
- ▶ Reveal decision boundaries
- ▶ Identify unstable regions

$$\nabla_x f(x)$$

Measures how predictions change locally.

## Defining a Classifier — Simple Version

### Binary classification setup

- ▶ Input: 2D feature vector  
 $x = (x_1, x_2)$
- ▶ Output: probability  $\hat{y} \in [0, 1]$
- ▶ Decision via threshold

$$\hat{y} = f_{\theta}(x)$$

Class label inferred from  $\hat{y}$ .

#### Simple classifier model

```
1 class SimpleClassifier(nn.Module):  
2     def __init__(self):  
3         super().__init__()  
4         self.net = nn.Sequential(  
5             nn.Linear(2,1),  
6             nn.Sigmoid()  
7         )  
8     def forward(self,x):  
9         return self.net(x)
```

Minimal nonlinear decision model.

## Improving the Classifier — Nonlinear Model

### Why improve the model?

- ▶ Linear boundary often insufficient
- ▶ Real data is nonlinear
- ▶ Need higher expressive power

$$\hat{y} = \sigma(W_2 \phi(W_1 x))$$

Hidden layer enables nonlinear separation.

#### Better classifier model

```
1 class BetterClassifier(nn.Module):  
2     def __init__(self):  
3         super().__init__()  
4         self.net = nn.Sequential(  
5             nn.Linear(2, 32),  
6             nn.ReLU(),  
7             nn.Linear(32, 1),  
8             nn.Sigmoid()  
9         )  
10    def forward(self, x):  
11        return self.net(x)
```

Nonlinear decision boundary.

## Applying the Classifier to Data

### Training the classifier

- ▶ Input: feature vectors  $(x_1, x_2)$
- ▶ Output: class probability  $\hat{y}$
- ▶ Supervised binary classification

### Goal

- ▶ Separate two classes
- ▶ Learn a decision boundary
- ▶ Minimize classification error

#### Training loop

```
1 loss_fn = nn.BCELoss()  
2 opt = optim.Adam(model.parameters(),  
                  lr=0.01)  
3  
4 for epoch in range(epochs):  
5     opt.zero_grad()  
6     y_p = model(X)  
7     loss = loss_fn(y_p, y_true)  
8     loss.backward()  
9     opt.step()
```

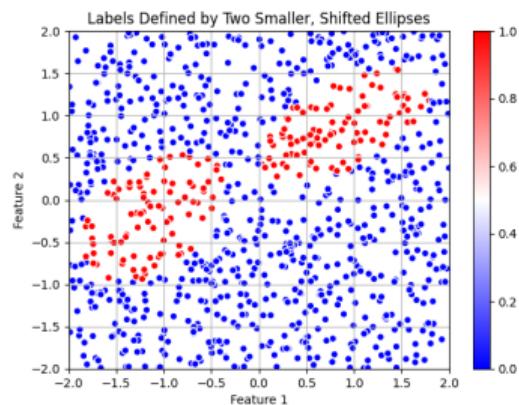
Prediction  $\hat{y} \in [0, 1]$  interpreted as probability.

```
grid_points = torch.stack([xx.flatten(), yy.flatten()], dim=1)  
grid_points.requires_grad = True  
grid_grads = grid_points.grad.detach().numpy()
```

## Labels vs. Learned Classification

### Ground truth labels

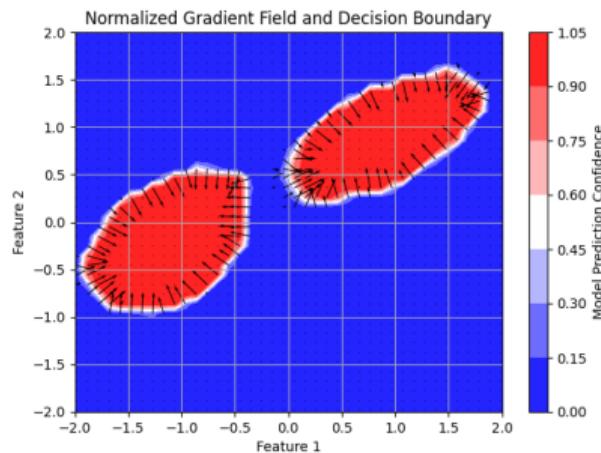
- ▶ Classes defined by geometry
- ▶ Two shifted, rotated ellipses
- ▶ Labels fixed before learning



Blue / red = predefined classes

### Model prediction

- ▶ Network output  $\hat{y} \in [0, 1]$
- ▶ Nonlinear decision boundary
- ▶ Smooth transition between classes



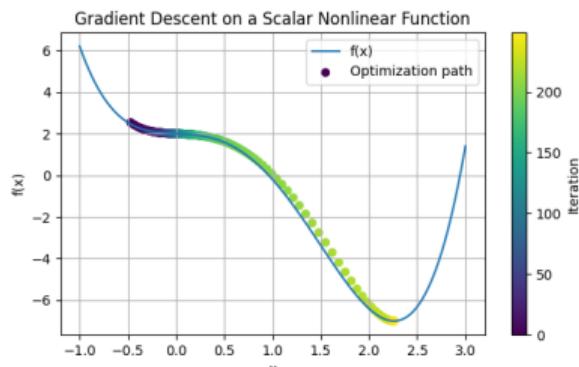
Colors = prediction, arrows = gradients

## Chapter 4 — Take-Home Messages

### What AI/ML really is

- ▶ Differentiable function approximation
- ▶ Learned from data, not hard-coded
- ▶ Optimized via gradients

Learn  $f_\theta$  s.t.  $f_\theta(x) \approx y$



### What really matters

- ▶ Tensors + autograd are the core
- ▶ Loss defines what is learned
- ▶ Data handling controls stability
- ▶ **Domain knowledge remains essential**

AI supports decisions — it does not replace responsibility.