

Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 
playing  with
Python  and AI/ML!
 

Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	Lab: Python Startup: Basics	Lab: Feed-forward and Graph NNs	Lab: Graph Learning with PyTorch	Lab: Containers and Reproducibility	Lab: Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	Lab: Work environments, Python everywhere	Lab: Simple Transformer and LLM Use	Lab: Agent Frameworks	Lab: CI/CD Pipelines	Lab: Graph-based Assimilation
12:45–13:30	Lunch Break				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	Lab: GRIB, NetCDF and Obs Visualisation	Lab: RAG Pipeline	Lab: DAWID Exploration	Lab: Anemoi Training Pipeline	Lab: Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	The AI Transformation	Learning from Observations Only
16:15–17:00	Lab: Torch Tensors and First Neural Net	Lab: Radar, SAT and Multimodal Data	Lab: MLflow Hands-on	Lab: How work style could change	Lab: ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

Lecture 9 — Diffusion and Flexible Graph Networks

From prediction to distributions

- ▶ Classical ML: predict a single output
- ▶ Reality: many plausible outcomes
- ▶ We need models that *sample*

Weather and climate are inherently:

- ▶ stochastic
- ▶ high-dimensional
- ▶ uncertain

Two complementary ideas

- ▶ Diffusion models learn how to turn noise into structure
- ▶ Graph neural networks learn from sparse, irregular observations

Key message

- ▶ Learn *distributions*, not just functions
- ▶ Learn *structure*, not just grids


Learning to Sample a Distribution

Regression is not enough

- ▶ Regression predicts a single value
- ▶ Many systems have multiple valid outcomes
- ▶ We need to model and sample the uncertainty !
- ▶ We want samples: $x \sim p(x)$

Key idea

- ▶ Learn a *mapping from noise to data*
- ▶ Noise represents uncertainty



Sampling via a Neural Network

Neural sampler: noise → data

Neural generator mapping noise to samples

```
1 class Generator(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.net = nn.Sequential(
5             nn.Linear(1, 64),
6             nn.ReLU(),
7             nn.Linear(64, 64),
8             nn.ReLU(),
9             nn.Linear(64, 1) )
10
11     def forward(self, z):
12         return self.net(z)
```

The network transforms white noise into samples.

What is different from regression?

- ▶ No target output for a given input
- ▶ Input is random noise
- ▶ Only the *distribution* matters

Interpretation

- ▶ $z \sim \mathcal{N}(0, 1)$
- ▶ $x = f_\theta(z) \sim p_\theta(x)$

```
z = torch.randn(N, 1) # noise
x = G(z)           # samples
```

Training without Targets: Distribution Matching

Key problem

There is no correct output x for a given noise input z .

Solution: compare distributions

Differentiable distribution loss (1D Wasserstein)

```
1 def wasserstein_1d(x_gen, x_data):  
2     xg, _ = torch.sort(x_gen.view(-1))  
3     xd, _ = torch.sort(x_data.view(-1))  
4     return torch.mean((xg - xd)**2)
```

What happens here?

- ▶ Draw samples from both distributions
- ▶ Sort them by value
- ▶ Compare *quantiles*

Interpretation

- ▶ We do not match samples
- ▶ We match ranks / transport

This turns sampling into a learnable problem.


Learning to Sample a Distribution

Training result

- ▶ Network maps noise to samples
- ▶ Generated samples follow $p(x)$
- ▶ Multimodal structure is recovered

Key point

- ▶ No regression target was used
- ▶ Only a distribution loss



Target distribution and generated samples

This already enables stochastic modeling.

Why One-Step Sampling Is Not Enough

What we did so far

- ▶ Learn a direct map: noise → data
- ▶ Works well in low dimensions
- ▶ Works for simple distributions

But this has limits

- ▶ High-dimensional distributions are complex
- ▶ Direct mapping becomes unstable
- ▶ Hard to represent fine structure

Key idea

- ▶ Do not sample in one step
 - ▶
- Use many small, simple steps

Intuition

- ▶ Large transport is hard
- ▶ Small corrections are easy

This leads to diffusion models .

Forward Diffusion: Adding Noise Step by Step

Forward process

We gradually destroy structure by adding noise:

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, 1)$$

- ▶ $t = 0$: clean data
- ▶ $t \uparrow$: increasing noise
- ▶ $t = T$: pure noise

Important properties

- ▶ Process is fixed
- ▶ No learning involved
- ▶ Fully known statistics

Why do this?

- ▶ Generates training data
- ▶ Connects data to noise

This defines the learning problem.

Reverse Diffusion: Learning to Denoise

Idea: Train a network to

remove a *small amount of noise*:

$$x_t \longrightarrow x_{t-1}$$

Denoising step (learned reverse process)

```
1 # x_t : noisy signal at step t
2 # model learns to predict a less
   noisy version
3
4 for t in reversed(range(1, T)):
5     x = model(x, t) # denoise one step
```

Training principle



- ▶ Start from clean signal x_0
- ▶ Add noise step by step
- ▶ Train the model to undo each step

Key point

- ▶ Same network for all steps
- ▶ Only local corrections

Each step is simple and stable.

Diffusion in Action: Example Result



What we see

- ▶ Start from pure noise
- ▶ Apply learned denoising steps
- ▶ End up with valid samples

Key observation

- ▶ Samples follow the target distribution
- ▶ Stochastic but structured

This is sampling by refinement , not regression.

Forward Diffusion Process

Let $x_0 \sim p_{\text{data}}(x)$ be a data sample.

The forward diffusion process is a Markov chain:

$$q(x_{1:T} | x_0) = \prod_{t=1}^T q(x_t | x_{t-1})$$

with Gaussian transitions

$$q(x_t | x_{t-1}) = \mathcal{N}\left(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I\right),$$

with $\beta_t \in (0, 1)$.

Equivalently, each step can be written as

$$x_t = \sqrt{1 - \beta_t} x_{t-1} + \sqrt{\beta_t} \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, I).$$

Consequences

- ▶ Mean is damped at each step
- ▶ Variance increases monotonically
- ▶ Signal-to-noise ratio decreases

After many steps:

$$x_T \approx \mathcal{N}(0, I)$$

This connects data to pure noise .

Closed-Form Forward Diffusion

Because all forward steps are Gaussian, the marginal distribution of x_t given x_0 has a closed form:

$$q(x_t | x_0) = \mathcal{N}\left(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t)I\right)$$

with cumulative noise factor

$$\bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s).$$

Equivalently, sampling can be written as

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I).$$

Interpretation

- ▶ Signal amplitude decays with $\sqrt{\bar{\alpha}_t}$
- ▶ Noise amplitude grows with $\sqrt{1 - \bar{\alpha}_t}$
- ▶ t directly controls noise level

As $t \rightarrow T$:

$$\bar{\alpha}_t \rightarrow 0 \quad \Rightarrow \quad x_t \sim \mathcal{N}(0, I).$$

This matches the example exactly.

Reverse Diffusion Process

The reverse process aims to invert the forward diffusion:

$$p_{\theta}(x_{0:T}) = p(x_T) \prod_{t=1}^T p_{\theta}(x_{t-1} | x_t),$$

with $p(x_T) = \mathcal{N}(0, I)$.

Each reverse step is modeled as a Gaussian:

$$p_{\theta}(x_{t-1} | x_t) = \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_t),$$

where μ_{θ} is predicted by a neural network.

The variance Σ_t is usually fixed or predefined.

What is learned?

- ▶ Only the conditional mean
- ▶ One network for all t
- ▶ Local denoising steps

Key idea

Each reverse step removes
a small amount of noise .

This turns sampling into a
sequence of simple corrections.

Why Diffusion Sampling Is Correct

Assume a scalar Gaussian data distribution:

$$x_0 \sim \mathcal{N}(\mu_0, \sigma_0^2), \quad x_1 = x_0 + \epsilon,$$

with $\epsilon \sim \mathcal{N}(0, \sigma^2)$.

The optimal denoiser (MSE sense) is:

$$f^*(x_1) = \mathbb{E}[x_0 | x_1] = \frac{\sigma^2 \mu_0 + \sigma_0^2 x_1}{\sigma_0^2 + \sigma^2}.$$

This pulls the noisy sample toward μ_0 , but reduces variance.

Key observation

- ▶ Denoising alone shrinks variance
- ▶ Mean is biased toward μ_0

Therefore

To recover the full distribution, we must sample:

$$x_0 = f^*(x_1) + \sqrt{\sigma_{\text{post}}^2} \xi, \quad \xi \sim \mathcal{N}(0, 1).$$

This is exactly what diffusion does at every step.

Why Graph Networks?

The problem

- ▶ Data is sparse, irregular, incomplete
- ▶ Classical grids assume full coverage
- ▶ Interpolation rules are often fixed

In many applications, we only know values at

- ▶ a few locations,
- ▶ irregular positions,
- ▶ changing resolutions.

The graph perspective

- ▶ Nodes represent locations
- ▶ Edges represent influence / neighborhood
- ▶ Features store values and metadata

Key idea

- ▶ Geometry is encoded in the graph
- ▶ Learning happens locally
- ▶ Same model works on different domains

Graph Example: Reconstructing a Function from Sparse Data


Setup

- ▶ 1D spatial domain discretized into nodes
- ▶ True function $f(x)$ defined on all nodes
- ▶ Observations available only at few locations

Each node carries:

- ▶ observed value (or zero if missing)
- ▶ a binary observation mask

The task is to reconstruct $f(x)$ at *all* nodes.



Sparse observations and GNN reconstruction

Message Passing: How Information Flows

Local updates

Each node updates its state by combining:

- ▶ its own current value,
- ▶ information from neighboring nodes.

This update is *learned*, not prescribed.

Key mechanism

- ▶ Same update rule for all nodes
- ▶ Shared parameters across the graph
- ▶ Repeated over multiple layers

Interpretation

- ▶ Observations act as sources
- ▶ Information spreads gradually
- ▶ Missing values are inferred

Why this matters

- ▶ No fixed interpolation stencil
- ▶ Flexible neighborhood influence
- ▶ Scales to new domains

Message Passing: Minimal Code Example

One GNN message-passing layer (conceptual)

```
1 class GNNLayer(nn.Module):
2     def __init__(self, dim):
3         super().__init__()
4         self.self_lin = nn.Linear(dim, dim)
5         self.neigh_lin = nn.Linear(dim, dim)
6
7     def forward(self, x, edge_index):
8         row, col = edge_index
9         agg = torch.zeros_like(x)
10        agg.index_add_(0, row, x[col])
11        return F.relu(
12            self.self_lin(x) +
13            self.neigh_lin(agg) )
```

What this shows

- ▶ Neighbor features are summed
- ▶ Same weights used everywhere
- ▶ Graph defines information flow

Key insight

- ▶ No coordinates needed
- ▶ No stencil prescribed
- ▶ Structure comes from edges

Generalization: Same Network, Different Domain

What changes


- ▶ Number of nodes N
- ▶ Physical domain length
- ▶ Resolution of the grid

What stays the same


- ▶ Network architecture
- ▶ Trainable parameters
- ▶ Message-passing rule

Trained on $[0, \pi]$, applied on $[0, 4\pi]$.

Only the *graph* is rebuilt.



Failed reconstruction (absolute coordinates)



Successful (coordinate-free, graph-based)

Why PyTorch Lightning?

Problem with raw PyTorch

- ▶ Training loops get long and repetitive
- ▶ Device handling (CPU/GPU) is manual
- ▶ Logging and checkpoints are ad hoc

Code quickly becomes hard to read and maintain.

```
from torch_geometric.loader import DataLoader
```

```
# Create dataset
dataset = [generate_sample_graph(n_nodes=64, dn=8, k=3) for _ in range(200)]
loader = DataLoader(dataset, batch_size=16, shuffle=True)
```

```
# Train
```

```
model = GNNInterpolator()
trainer = pl.Trainer(max_epochs=200, logger=False, enable_checkpointing=False)
```

LightningModule: Clean Structure

Minimal Lightning module structure

```
1 class Model(pl.LightningModule):
2     def forward(self, x):
3         ...
4
5     def training_step(self, batch,
batch_idx):
6         ...
7         return loss
8
9     def configure_optimizers(self):
10        return torch.optim.Adam(self.
parameters())
```

Key idea

- ▶ No explicit training loop
- ▶ No device checks
- ▶ No boilerplate code

Focus stays on the model and the loss.

What Lightning Handles for You

Automatically

- ▶ CPU / GPU / multi-GPU
- ▶ Epoch and batch loops
- ▶ Gradient handling
- ▶ Logging

```
class GNNInterpolator(pl.LightningModule):  
    def __init__(self):  
        super().__init__()  
        self.gcn1 = GCNConv(2, n1) # 2 input features  
        self.gcn2 = GCNConv(n1, n2)  
        self.out = GCNConv(n2, 1)  
  
    def forward(self, data):  
        x = data.x_feat  
        x = F.relu(self.gcn1(x, data.edge_index))  
        x = self.gcn2(x, data.edge_index)
```

You still control

- ▶ Model architecture
- ▶ Loss functions
- ▶ Optimizers
- ▶ Data handling

Lightning enforces structure — not constraints.

```
def training_step(self, batch, batch_idx):  
    pred = self(batch)  
    mask = ~torch.isnan(batch.y_obs)  
    loss = F.mse_loss(pred[~mask], batch.y[~mask])  
    self.log("train_loss", loss)  
    return loss
```

Why PyTorch Geometric?

Graph-specific challenges

- ▶ Variable graph sizes
- ▶ Sparse connectivity
- ▶ Efficient message passing

Helps implementing this efficiently.

PyTorch Geometric

- ▶ Native graph data structures
- ▶ Optimized message passing
- ▶ Many standard GNN layers

Built directly on PyTorch tensors.

Graphs in PyTorch Geometric

Basic PyG data object

```
1 from torch_geometric.data import Data
2
3 data = Data(
4     x = node_features,
5     edge_index = edge_index,
6     y = targets
7 )
```

Key concept

- ▶ Nodes = rows of x
- ▶ Edges define message flow
- ▶ Graph size is flexible

No grid assumptions required.

PyTorch Geometric + Lightning

Why combine them?

- ▶ PyG: graph operations
- ▶ Lightning: clean training


Each library does one thing well.

Result

- ▶ Clean, readable code
- ▶ Scales to large graphs
- ▶ Easy experimentation

Ideal for research and teaching.

Graph Neural Network Structure



Example structure with latent dimension 64 times number of nodes.

Graph Representation: Data vs Structure

Node data

$$X \in \mathbb{R}^{N \times d}$$

$$\mathbf{x}_i \in \mathbb{R}^d \quad (i = 1, \dots, N)$$

- ▶ N : number of nodes
- ▶ d : features per node
- ▶ One feature vector per node

Interpretation

- ▶ Rows correspond to nodes
- ▶ Columns correspond to features
- ▶ Node order is arbitrary

Graph structure

$$\text{edge_index} \in \mathbb{N}^{2 \times E}$$

Each column defines one directed edge :

$$(j \rightarrow i)$$

- ▶ Encodes connectivity only
- ▶ No numerical weights
- ▶ No trainable parameters

Separation of roles

- ▶ Data: X
- ▶ Structure: edge_index
- ▶ Learning: weights (later)

Node-wise Feature Transformation

Input

$$X \in \mathbb{R}^{N \times d}, \quad \mathbf{x}_i \in \mathbb{R}^d$$

Learned feature map

$$\mathbf{h}_i = \sigma(W\mathbf{x}_i + \mathbf{b})$$

$$W \in \mathbb{R}^{h \times d}, \quad \mathbf{b} \in \mathbb{R}^h$$

$$H \in \mathbb{R}^{N \times h}$$

Key property

Same transformation for every node.

Node-wise linear layer

```
1 self.lin = nn.Linear(d, h)
2 h = F.relu(self.lin(x))
```

What happens here

- ▶ Feature extraction
- ▶ Dimensionality change
- ▶ No neighbor interaction

What does *not* happen

- ▶ No graph usage
- ▶ No message passing
- ▶ No dependence on N

Message Passing: Injecting Graph Structure

Latent node features

$$H \in \mathbb{R}^{N \times h},$$

$$\mathbf{h}_i \in \mathbb{R}^h$$

Neighborhood aggregation

$$\mathbf{a}_i = \sum_{j \in \mathcal{N}(i)} \mathbf{h}_j,$$

$$A \in \mathbb{R}^{N \times h}$$

Graph-aware update

$$\mathbf{h}'_i = \sigma(W_{\text{self}}\mathbf{h}_i + W_{\text{neigh}}\mathbf{a}_i)$$

Graph affects *which terms are summed*, not the number of parameters.

Edge-based aggregation

```
1 row, col = edge_index
2 agg = torch.zeros_like(h)
3 agg.index_add_(0, row, h[col])
```

- ▶ Uses `edge_index` only
- ▶ No learnable graph parameters
- ▶ Permutation invariant

Effect

- ▶ Information exchange
- ▶ Local propagation
- ▶ One-hop interaction

Stacking Message Passing and Output Projection

Layer stacking

$$H^{(0)} \in \mathbb{R}^{N \times d},$$

$$H^{(1)} = \Phi(H^{(0)}, \text{edge_index}),$$

$$H^{(2)} = \Phi(H^{(1)}, \text{edge_index})$$

Each layer expands the receptive field:

- ▶ 1 layer: 1-hop neighbors
- ▶ 2 layers: 2-hop neighbors

Final node representation

$$H^{(L)} \in \mathbb{R}^{N \times h}$$

Output projection

$$\hat{y}_i = w^\top h_i^{(L)} + b,$$

$$w \in \mathbb{R}^h, \quad b \in \mathbb{R}$$

$$\hat{y} \in \mathbb{R}^N$$

Final linear layer

```
1 self.out = nn.Linear(h, 1)
2 y_hat = self.out(h).squeeze(-1)
```

Key properties

- ▶ Same output map for all nodes
- ▶ Independent of graph size
- ▶ Fully permutation invariant

Lecture 9 — Big Picture

What we learned

- ▶ **Sampling** instead of regression
- ▶ Explicit modeling of **uncertainty**
- ▶ Generating distributions, not point estimates

Diffusion networks

- ▶ Noise → data via iterative refinement
- ▶ Correct sampling through stochastic reverse steps
- ▶ Strong theoretical foundation

Graph networks

- ▶ Flexible representation of structure
- ▶ Learning from sparse, irregular data
- ▶ Strong generalization across domains

Order, structure, efficiency

- ▶ Inductive bias through graphs and locality
- ▶ Separation of geometry and learning
- ▶ Scalable implementations with
 - ▶ PyTorch Lightning
 - ▶ PyTorch Geometric

Modern ML combines **sampling**, **structure**, and **efficient tooling** to build flexible, generalizable models.