

Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 
 playing  with
 Python  and AI/ML!
 

Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	Lab: Python Startup: Basics	Lab: Feed-forward and Graph NNs	Lab: Graph Learning with PyTorch	Lab: Containers and Reproducibility	Lab: Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	Lab: Work environments, Python everywhere	Lab: Simple Transformer and LLM Use	Lab: Agent Frameworks	Lab: CI/CD Pipelines	Lab: Graph-based Assimilation
12:45–13:30	Lunch Break				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	Lab: GRIB, NetCDF and Obs Visualisation	Lab: RAG Pipeline	Lab: DAWID Exploration	Lab: Anemoi Training Pipeline	Lab: Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	The AI Transformation	Learning from Observations Only
16:15–17:00	Lab: Torch Tensors and First Neural Net	Lab: Radar, SAT and Multimodal Data	Lab: MLflow Hands-on	Lab: How work style could change	Lab: ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

MLOps: From Machine Learning to Operations

Why this lecture

- ▶ Machine learning models are **no longer research prototypes**
- ▶ AI systems increasingly enter **operational environments**
- ▶ Reliability, reproducibility and traceability become critical

Traditional ML focuses on *training*.

Operations focus on *stability*.

MLOps connects both worlds.

Key question

- ▶ How do we turn ML code into an operational system?

This includes

- ▶ controlled software environments
- ▶ automated build and deployment
- ▶ clear separation of roles and responsibilities
- ▶ safe execution on HPC and production systems

This lecture focuses on *processes*, not algorithms.
But also lots of **continuous integration (CI)**.

Why Machine Learning Needs Operations

Classical software assumptions

- ▶ Code fully defines behavior
- ▶ Same input → same output
- ▶ Changes are explicit and infrequent

Once deployed, behavior is largely predictable and stable.



Data Improvement

More & better training data



Bug Fixes

Fix errors in code/pipeline

Machine learning reality

- ▶ Behavior emerges from **data**
- ▶ Outputs are **statistical, not deterministic**
- ▶ Models **degrade as data distributions change**

Operational consequences

- ▶ Retraining becomes part of operations
- ▶ Monitoring must include **model performance**
- ▶ Reproducibility extends **beyond source code**



Feature Engineering
New features & transformations



Architecture
Improve model architecture



Hyperparameters
Optimize training parameters

Machine learning systems require operational control from the very beginning.

What Is MLOps?

Basic idea

- ▶ MLOps extends DevOps principles to machine learning
- ▶ Focus shifts from code alone to data and models
- ▶ The full model lifecycle becomes operational

MLOps treats ML systems as *long-lived services*, not one-off experiments.

What MLOps manages

- ▶ Data pipelines and preprocessing
- ▶ Training and validation workflows
- ▶ Model versions and metadata
- ▶ Deployment, monitoring and rollback

Why this is necessary

- ▶ ML behavior changes even when code does not
- ▶ Manual processes do not scale
- ▶ Missing control leads to silent failures

MLOps is not a tool — it is an engineering discipline.

DevOps vs MLOps: What Really Changes

DevOps perspective

- ▶ Software behavior defined by **source code**
- ▶ Releases are **explicit and controlled**
- ▶ Bugs are fixed by changing code

How do we monitor systems, and control changes, fixes, improvements?



MLOps perspective

- ▶ Behavior emerges from **data and training**
- ▶ Models evolve through **retraining**
- ▶ Performance can change **without code changes**

Operational impact

- ▶ Data becomes a **first-class artifact**
- ▶ Validation must be **continuous**
- ▶ Rollback applies to **models**, not just code

MLOps extends DevOps — it does not replace it.

The Classical MLOps Cycle

From experiment to operation

- ▶ Data ingestion and preprocessing
- ▶ Model training and validation
- ▶ Deployment into an operational environment

Operational feedback loop

- ▶ Monitoring of system and model performance
- ▶ Detection of drift and degradation
- ▶ Triggering retraining or rollback

This cycle extends the classical DevOps loop by explicitly including data and models.

Key property

- ▶ The cycle is continuous, not linear
- ▶ Operations actively influence development

Modern MLOps has many elements of typical NWP development cycles. NWP also integrates data, e.g. orography, canopy layers, and real data through observations.

MLOps is a closed loop, not a one-time deployment.

How Automation Is Achieved in MLOps

Automation principles

- ▶ Each step is explicitly defined and scripted
- ▶ Execution is trigger-based, not manual
- ▶ Results are logged and versioned automatically

Automation replaces informal procedures by repeatable execution rules.

Technical mechanisms

- ▶ CI/CD pipelines trigger builds and checks
- ▶ Version control tracks code, configs, and metadata
- ▶ Registries store models and runtime artifacts
- ▶ Containers freeze execution environments

Key effect

- ▶ Humans decide what should happen
- ▶ Systems enforce how it happens

Automation makes MLOps scalable, auditable, and safe.

CI and CD in MLOps

Continuous Integration (CI)

CI = Every change is automatically built, checked, and validated

- ▶ Triggered by commits to code or configuration
- ▶ Executes tests, linters, and consistency checks
- ▶ Detects problems early and reproducibly

Continuous Delivery / Deployment (CD)

CD = Validated artifacts are automatically prepared for operation

- ▶ Packages software, models, and environments
- ▶ Produces versioned, deployable artifacts
- ▶ Enables controlled rollout or rollback

CI/CD in MLOps

- ▶ CI validates logic and structure
- ▶ CD delivers operational artifacts

Important distinction

- ▶ CI/CD does not decide when models are used
- ▶ Operations retain execution control

Automation without loss of control.

MLOps as a Role-Based Pipeline

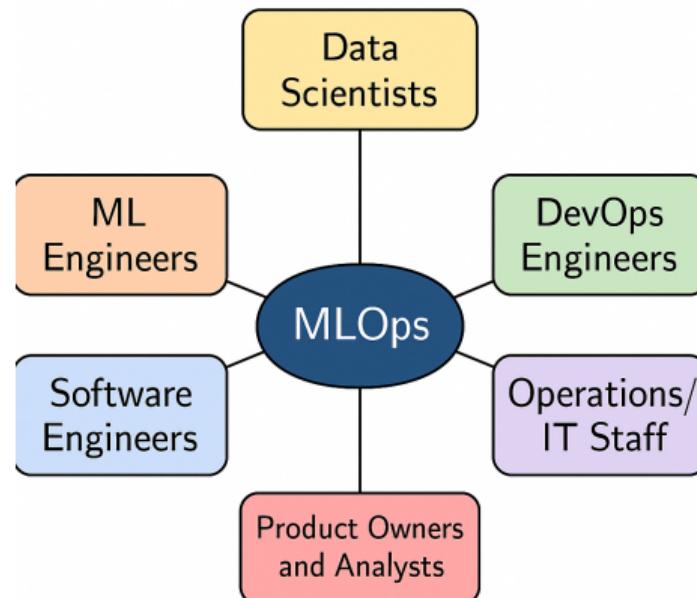
It is a **sequence of responsibilities** that together form a **complete lifecycle**:

- ▶ data preparation and feature design
- ▶ model development and validation
- ▶ packaging and deployment
- ▶ monitoring, feedback and retraining

Each of these steps:

- ▶ requires **different skills**,
- ▶ has **different risks**,
- ▶ and implies **different ownership**.

A useful way to understand MLOps is therefore as a **role game**, where different actors take responsibility for different parts of the pipeline.



Roles in MLOps Environments

Who Defines What “Good” Means?

In MLOps, quality is not a technical concept .

It is defined by:

- ▶ physical constraints,
- ▶ domain knowledge,
- ▶ operational requirements,
- ▶ user expectations.

These aspects cannot be learned from data alone.

Therefore:

- ▶ Domain experts define what “good” means ,
- ▶ ML systems optimize *towards* this definition,
- ▶ Operations validate it against reality.

Without a domain expertise, metrics are meaningless.

Domain Expert

- ▶ Defines quality and success criteria
- ▶ Specifies relevant metrics and constraints
- ▶ Interprets model behaviour in context
- ▶ Does not implement models or pipelines

Roles in the MLOps Pipeline

Data Scientist

- ▶ Develops models to optimize domain criteria
- ▶ Explores data and hypotheses
- ▶ Translates questions into ML problems
- ▶ **Does not decide operational relevance**

ML Engineer

- ▶ Turns models into software artifacts
- ▶ Builds training and inference pipelines
- ▶ **Ensures reproducibility and versioning**
- ▶ Bridges research and operations

DevOps / Platform Engineer

- ▶ Provides CI/CD and build infrastructure
- ▶ Manages containers and registries
- ▶ Ensures security and scalability
- ▶ **Does not tune models**

Domain Expert / Operations / Users

- ▶ Run models in production
- ▶ Monitor behaviour and failures
- ▶ Validate outputs against reality
- ▶ **Do not modify code or models**

Roles should be clear, people may be the same.

Starting Point: A Baseline ML Application

We start with a minimal working ML example.

The goal at this stage is simple:

- ▶ load a pretrained model,
- ▶ apply it to input data,
- ▶ inspect the output.

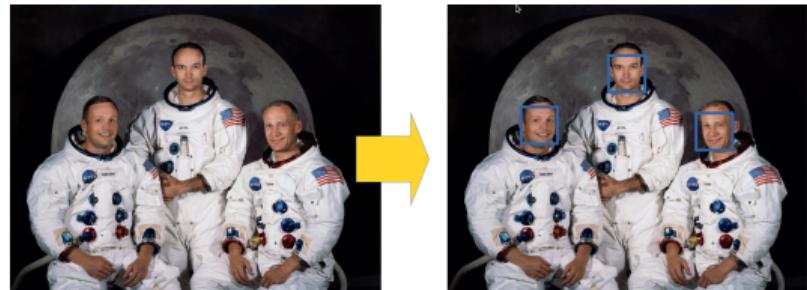
This corresponds to running the notebook:

- ▶ 00_face-detection-onnx.py

At this point:

- ▶ the method works,
- ▶ results can be inspected visually.

Nothing here is operational yet.



The notebook instantiates a pretrained face-detection model and applies it directly to an input image.

```
def mark_faces(image_filename):  
    """Mark all faces recognized in the image"""\n    image = PIL.Image.open(image_filename)  
  
    faces = detect_faces(image)  
  
    render_data = detections_to_render_data(  
        faces, bounds_color=Colors.GREEN, line_w
```



Input image

Raw input data, loaded via PIL. No preprocessing, no metadata, no assumptions beyond file availability.



Model output

Bounding boxes produced by a pretrained face-detection model and rendered directly onto the image.

This is pure inference: load → detect → render.

Docker: Installation and Input Files

If you have **admin rights** on your computer, installation of docker is easy.

```
brew install --cask docker
```

Step 1: Install Docker

- ▶ Install **Docker Desktop** (Mac)
- ▶ Start Docker service
- ▶ Verify installation:

Check Docker

```
1 docker --version
2 docker run hello-world
```

If this works, Docker is **ready to use**.

Step 2: Python code to run in Docker

02_code_to_execute_in_docker.py

```
1 from pathlib import Path
2
3 print("Hello from Docker")
4
5 out = Path("docker_output.txt")
6 out.write_text("Generated inside
    Docker\n")
7
8 print("Wrote:", out)
```

This is **ordinary Python**. **No Docker logic** inside the code.

Dockerfile, Image Build, and Execution

Step 4: Build and run

Step 3: Define the container

```
FROM python:3.11-slim
WORKDIR /app
COPY 02_code_to_execute_in_docker.py /app/
CMD ["python", "02_code_to_execute_in_docker.py"]
```

Docker **freezes the runtime environment**:

- ▶ operating system (Debian slim)
- ▶ Python version (3.11.x)
- ▶ system libraries and pip

A Docker image already acts as a virtual environment.

No venv activation is required; packages installed in the image are isolated by default.

Build image

```
1 docker build -t \
2   python-hello-docker \
3   -f 02_dockerfile.txt .
```

Run container

```
1 docker run --rm \
2   -v "$(pwd):/app" \
3   python-hello-docker
```

Volume mount exposes results. Without **-v**, files disappear.

Building Docker Containers on GitHub

General principle

- ▶ Docker containers do **not need** to be built locally
- ▶ GitHub provides **CI runners** with Docker installed
- ▶ A container can be built **directly on the platform**

The build is moved from the laptop to the platform.

Concrete example (this lecture)

- ▶ Python script generates a plot
- ▶ Dockerfile defines the runtime
- ▶ GitHub Actions builds the image
- ▶ Image is pushed to a **container registry**

Result

- ▶ Reproducible container image
- ▶ Independent of local setup
- ▶ Ready for download and execution

Code lives in Git.

Images live in the registry.

What is required

- ▶ Source code
- ▶ A **Dockerfile**
- ▶ A CI workflow definition

CI builds containers as first-class artifacts.

Downloading and Running a Container Image

Pulling a container image

- ▶ Images are downloaded via Docker CLI
- ▶ Private images require authentication
- ▶ Access is granted via registry tokens

Download image

```
1 docker pull ghcr.io/eumetnet-e  
-ai/docker-plot-03:latest
```

Running the container

```
docker run --rm \  
-v "$(pwd)/03_docker_plot:/app" \  
ghcr.io/eumetnet-e-ai/docker-plot-03:latest
```

Why the volume mount matters

- ▶ Container files are ephemeral
- ▶ Volume mounts expose results
- ▶ Without -v, outputs disappear

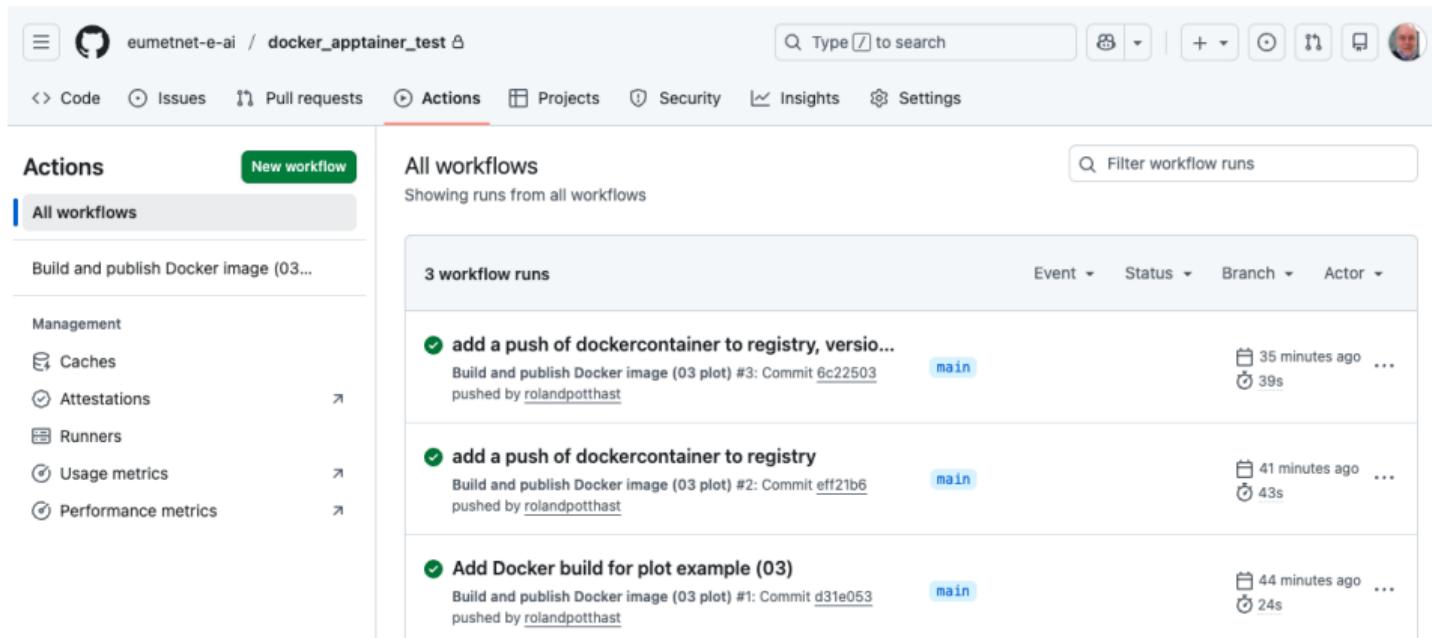
Execution is decoupled from development.

Important distinction

- ▶ Git access \neq registry access
- ▶ Separate permissions for code and containers

Build once, run anywhere — with controlled environments.

Docker Containers Built on GitHub



The screenshot shows a GitHub repository page for 'eumetnet-e-ai / docker_apptainer_test'. The 'Actions' tab is selected. On the left, there's a sidebar with 'Actions' and 'Management' sections. The 'Management' section includes 'Caches', 'Attestations', 'Runners', 'Usage metrics', and 'Performance metrics'. The main area displays 'All workflows' with three workflow runs listed:

Event	Status	Branch	Actor
add a push of dockercontainer to registry, versio...	main	35 minutes ago	...
add a push of dockercontainer to registry	main	41 minutes ago	...
Add Docker build for plot example (03)	main	44 minutes ago	...

Each row shows a green checkmark, the commit message, the branch ('main'), the time of the event, and a '...' button.

Source code in Git → CI build on GitHub → Container registry → Pull and run anywhere

GitHub Container Registry (ghcr.io)

What is ghcr.io?

- ▶ GitHub's container registry service
- ▶ Stores Docker / OCI container images
- ▶ Integrated with GitHub repositories and CI

How it works

- ▶ CI builds an image from a Dockerfile
- ▶ Image is pushed to ghcr.io
- ▶ Image is identified by name, tag, and SHA

Containers are artifacts, not source code.

Access and usage

- ▶ Images are pulled via Docker or Apptainer
- ▶ Private images require authentication
- ▶ Access control is separate from Git

Browser access

- ▶ Containers can be inspected in the browser
- ▶ Tags, SHAs, and metadata are visible
- ▶ No direct download button

Click to inspect. Pull to execute.

ghcr.io stores images; runtimes execute them.

Apptainer on HPC: Build vs Run

What happens when pulling Docker images

- ▶ Apptainer converts Docker images into local containers
- ▶ This involves SquashFS compression
- ▶ Compression is memory- and thread-intensive

Typical failure on login nodes

- ▶ Limited memory and process counts
- ▶ SquashFS build cannot create threads

Out of memory (frag_thrd)

Failed to create thread

Important distinction

- ▶ Running containers is lightweight
- ▶ Building containers is a compute task

HPC policy (typical)

- ▶ Login nodes: editing, submitting jobs
- ▶ Compute nodes: builds, compression, heavy work

Apptainer pull = build step, not runtime.

Disk quota is irrelevant — memory and threads decide.

Correct AppTainer Workflow on HPC

Recommended solution

- ▶ Run container pulls as **batch jobs**
- ▶ Request sufficient **memory**
- ▶ Avoid compressed SIF builds on login nodes

Use sandbox format

- ▶ Uncompressed directory container
- ▶ Much lower memory pressure
- ▶ Fully usable with AppTainer

Batch job (concept)

```
1 #PBS -q gp_norm_all
2 #PBS -l memsz_job=128gb
3 #PBS -l cpunum_job=2
4
5 apptainer build --sandbox \
6 my-container \
7 docker://ghcr.io/ORG/IMAGE:TAG
```

Running the container

- ▶ No root privileges
- ▶ Current directory mounted by default
- ▶ Same performance as SIF

HPC rule: build elsewhere, run everywhere.

For production: build SIFs in CI, not on the cluster.

Building Apptainer Containers on GitHub

Motivation

- ▶ HPC systems often **cannot build containers locally**
- ▶ Docker images must be converted to **Apptainer (SIF)**
- ▶ Conversion is **memory-intensive**

Key idea

- ▶ Move the **build step off the HPC**
- ▶ Use **GitHub Actions** as build infrastructure
- ▶ Produce ready-to-run **SIF artifacts**

What GitHub provides

- ▶ Linux build nodes with Docker
- ▶ Sufficient memory for SquashFS
- ▶ Integrated access to **ghcr.io**

Required inputs

- ▶ Docker image in a registry
- ▶ Apptainer installed in CI
- ▶ Authentication token (`read:packages`)

Result: A portable **single-file container**.

CI builds containers, HPC only runs them.

CI Workflow: Docker Image to Apptainer SIF

Automated workflow

- ▶ Triggered after successful Docker build
- ▶ Runs on GitHub Linux runner
- ▶ Converts image to **SIF** format

Main steps within the Github environment

1. Install Apptainer
2. Authenticate to ghcr.io
3. Pull Docker image
4. Convert to **SIF**
5. Upload as artifact

This is **CI/CD** for runtime environments.

Core CI command on GitHub

```
1 apptainer pull \
2   docker-plot-03.sif \
3   docker://ghcr.io/ORG/IMAGE:TAG
```

Distribution

- ▶ SIF stored as CI artifact
- ▶ Downloadable via browser
- ▶ Transfer to HPC via scp

On the cluster

- ▶ **No build step**
- ▶ No Docker required
- ▶ Immediate execution

Triggering the SIF Build via GitHub Actions, YAML

What this workflow does

- ▶ Listens for a successful Docker build
- ▶ Runs only if the image exists
- ▶ Converts Docker → Apptainer

Why this design

- ▶ Avoids duplicate builds
- ▶ Enforces a clean dependency chain
- ▶ CI controls the runtime environment

Docker build is a prerequisite.

This makes the SIF build:

- ▶ deterministic, reproducible
- ▶ traceable to a commit

```
name: Export Apptainer SIF (Example 03)
on:
  workflow_run:
    workflows: ["Build and publish Docker image (03)"]
    types: [completed]
jobs:
  build-sif:
    if: ${{ github.event.workflow_run.conclusion == 'success' }}
    runs-on: ubuntu-latest
    permissions:
      contents: read
      packages: read
    steps:
      - name: Install Apptainer
        run: |
          sudo apt-get update
          sudo apt-get install -y \
```

Running the Apptainer Container on the HPC

On the HPC system

- ▶ No Docker available
- ▶ Apptainer pre-installed
- ▶ No build permissions required

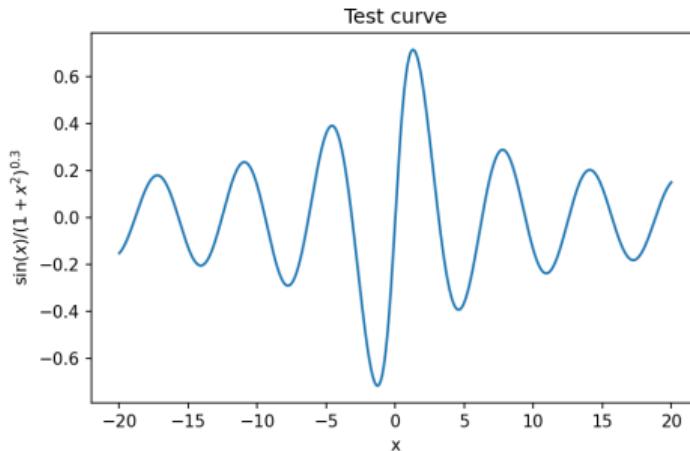
Run the container

HPC execution

```
1 apptainer exec \
2 docker-plot-03.sif python \
3 -B 03_plot_curve.py
```

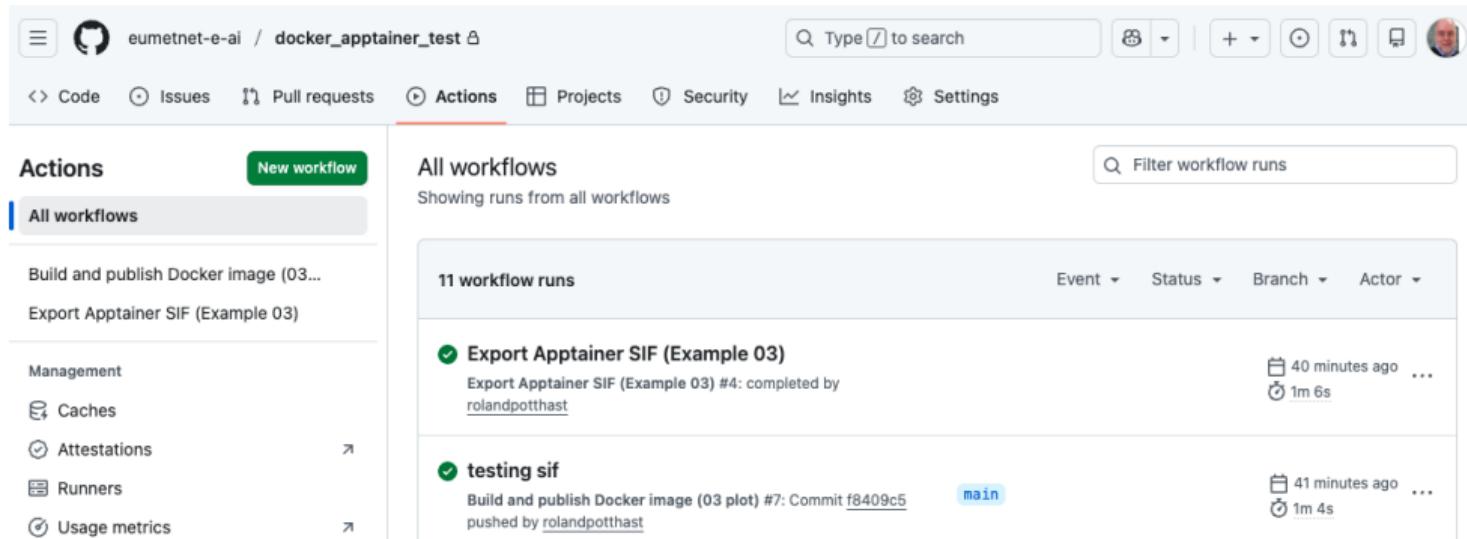
What happens

- ▶ Python code runs inside container
- ▶ Output file written to working directory
- ▶ No external dependencies needed



Same container, different system.

From Docker Image to Apptainer SIF on GitHub



The screenshot shows a GitHub repository page for `eumetnet-e-ai / docker_apptainer_test`. The `Actions` tab is selected. On the left, a sidebar lists management options: `All workflows`, `Caches`, `Attestations`, `Runners`, and `Usage metrics`. The main area displays the `All workflows` section, which lists 11 workflow runs. Two runs are highlighted:

- Export Apptainer SIF (Example 03)**: Completed by `rolandpotthast` 40 minutes ago, 1m 6s ago.
- testing sif**: Completed by `rolandpotthast` 41 minutes ago, 1m 4s ago. This run is associated with a commit `f8409c5` on the `main` branch.

Why Containers are Essential in ML Frameworks

The core problem in ML

- ▶ ML code depends on:
 - ▶ specific library versions
 - ▶ CUDA / CPU features
 - ▶ system-level dependencies
- ▶ These dependencies **change over time**
- ▶ Results become hard to reproduce

ML models are not standalone artifacts.

What containers provide

- ▶ Encapsulation of:
 - ▶ code
 - ▶ libraries
 - ▶ runtime environment
- ▶ Identical execution on:
 - ▶ laptops
 - ▶ CI systems
 - ▶ HPC clusters

Containers turn models into executable artifacts.

Reproducibility, portability, and controlled execution are **non-negotiable** in ML.