

# Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀  
playing 🤖 with  
Python 🐍 and AI/ML!  
 

## Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	<b>Lab:</b> Python Startup: Basics	<b>Lab:</b> Feed-forward and Graph NNs	<b>Lab:</b> Graph Learning with PyTorch	<b>Lab:</b> Containers and Reproducibility	<b>Lab:</b> Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	<b>CI/CD for Machine Learning</b>	AI-based Data Assimilation
12:00–12:45	<b>Lab:</b> Work environments, Python everywhere	<b>Lab:</b> Simple Transformer and LLM Use	<b>Lab:</b> Agent Frameworks	<b>Lab:</b> CI/CD Pipelines	<b>Lab:</b> Graph-based Assimilation
12:45–13:30	<b>Lunch Break</b>				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	<b>Lab:</b> GRIB, NetCDF and Obs Visualisation	<b>Lab:</b> RAG Pipeline	<b>Lab:</b> DAWID Exploration	<b>Lab:</b> Anemoi Training Pipeline	<b>Lab:</b> Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	<b>The AI Transformation</b>	Learning from Observations Only
16:15–17:00	<b>Lab:</b> Torch Tensors and First Neural Net	<b>Lab:</b> Radar, SAT and Multimodal Data	<b>Lab:</b> MLflow Hands-on	<b>Lab:</b> How work style could change	<b>Lab:</b> ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

## Why CI/CD Is Not Optional for AI/ML

### What continuously changes in AI/ML

- ▶ Model parameters through retraining
- ▶ Training and validation data
- ▶ Feature engineering and preprocessing
- ▶ Hyperparameters and runtime configuration

As a consequence:

- ▶ System behavior is **not fixed**
- ▶ Outputs depend on  
code, data, and environment
- ▶ Small changes can have **large effects**

### Why manual workflows break down

- ▶ Experiments cannot be reproduced reliably
- ▶ Results depend on undocumented environments or training data
- ▶ Errors surface **late or not at all**
- ▶ Deployment decisions become guesswork

Without automation:

- ▶ Models cannot be trusted operationally
- ▶ Debugging becomes **forensic work**

**CI/CD is the mechanism that makes AI/ML systems controllable and trustworthy.**

## AI/ML Is Not Special — But Updating Becomes Complex

### From a software engineering perspective

- ▶ AI/ML systems are still **software systems**
- ▶ They use standard languages, libraries and toolchains
- ▶ The same DevOps principles apply

There is **no special engineering magic** in AI/ML:

- ▶ version control
- ▶ testing
- ▶ packaging
- ▶ deployment

**CI/CD is needed to manage updates, not to handle “AI magic”.**

### Where the real complexity comes from

- ▶ Frequent retraining with **new data**
- ▶ Multiple preprocessing and feature pipelines
- ▶ Changing **model architectures**
- ▶ Different **loss functions** and objectives
- ▶ Many runtime and training configurations

As a result:

- ▶ Updates are **continuous**
- ▶ Reproducibility becomes **non-trivial**
- ▶ Manual tracking no longer works

# AI/ML CI/CD Pipeline

Key components for building robust MLOps workflows

## Training Data

- Data Versioning
- Quality Checks
- Labeling Pipeline
- Data Drift Detection

## Model Architecture

- Model Registry
- Architecture Versioning
- Hyperparameters
- Code Repository

## MLflow

- Experiment Tracking
- Model Registry
- Artifact Storage
- Metrics Logging

## CI/CD Pipeline

- Automated Testing
- Model Training
- Build & Package
- Deployment Automation

## Verification

- Model Validation
- Performance Tests
- A/B Testing
- Shadow Deployment

## Applications

- API Endpoints
- Batch Inference
- Real-time Serving
- Model Integration

## Monitoring

- Performance Metrics
- Model Drift
- System Health
- Alert Management

## Starting Locally: A Concrete Git Hook Example

Black is an automatic code formatter for Python.

*It rewrites Python source code into a single, consistent style, without asking questions.*

### Minimal pre-commit hook

.git/hooks/pre-commit

```
1 #!/bin/sh
2 pytest || exit 1
3 black .
```

This hook is executed automatically when running:

- ▶ git commit

If tests fail, the commit is blocked.

**Git hooks enforce local discipline — not global policy.**

### What this enforces locally

- ▶ Code must be syntactically correct
- ▶ Tests must pass
- ▶ Code formatting is consistent

Key properties:

- ▶ Runs before code leaves the laptop
- ▶ No CI server involved
- ▶ Immediate feedback to the developer

## When Are Git Hooks Executed?

### Git commands trigger hooks

Git automatically executes hooks at well-defined points in its workflow.

Common examples:

- ▶ `pre-commit` — before a commit is created
- ▶ `commit-msg` — to validate commit messages
- ▶ `pre-push` — before pushing to a remote

Hooks run **before** Git completes the command.

### Effect on the workflow

- ▶ Hook succeeds → Git continues
- ▶ Hook fails → Git aborts the command

This means:

- ▶ Invalid code never enters the repository
- ▶ Errors are caught **immediately**
- ▶ No manual checks are required

Hooks are **deterministic**: same input → same outcome.

**Hooks turn Git commands into enforced quality gates.**

## A Git Hook in Action: What Actually Happens

### Before committing

- ▶ Python test file is **poorly formatted**
- ▶ Code is syntactically valid
- ▶ Tests pass when run manually

### After running black:

- ▶ Formatting is rewritten automatically
- ▶ No code logic is changed

### During git commit

- ▶ pytest is executed automatically
- ▶ black is executed automatically
- ▶ Both run via the **pre-commit hook**

The hook turns a manual checklist into an automatic guarantee.

### Outcome:

- ▶ Tests pass
- ▶ Formatting is consistent
- ▶ Commit is **accepted**

```
def add( a ,b ):  
    return a+b  
def test_answer( ):  
    print("Testing add(1, 3) == 4")  
    assert add(1,3)==4  
  
def add(a, b):  
    return a + b  
def test_answer():  
    print("Testing add(1, 3) == 4")  
    assert add(1, 3) == 4
```

## Strengths and Limits of Git Hooks

### Strengths

- ▶ Extremely fast feedback . Helps you!!
- ▶ No external infrastructure required
- ▶ Works offline
- ▶ Integrated directly into Git commands

### Limitations

- ▶ Run only on the developer machine
- ▶ Not enforced across a team
- ▶ Can be bypassed or disabled
- ▶ No neutral execution environment

Git hooks are ideal for:

- ▶ formatting checks
- ▶ unit tests
- ▶ catching trivial mistakes early

As a consequence:

- ▶ Hooks improve discipline
- ▶ But they do **not guarantee quality**

**Git hooks accelerate development — CI platforms enforce standards.**

## Formatting vs Behavior: Two Different Checks

### Code formatting (Black)

- ▶ Enforces a consistent style
- ▶ Removes whitespace and layout differences
- ▶ Does **not** change program logic

Example:

black reformats code

```
1 def add( a ,b ):  
2     return a+b
```

### Behavior testing (pytest)

- ▶ Checks whether code does the right thing
- ▶ Executes functions and validates results
- ▶ Fails if behavior changes unexpectedly

Formatting makes code readable. Testing makes code **correct**.

becomes:

```
1 def add(a, b):  
2     return a + b
```

**Style and correctness are independent concerns.**

## Code and Test: Defining and Enforcing Behavior

### Formatted Code

```
test_example.py
```

```
1 def add(a, b):  
2     return a + b
```

This code:

- ▶ is syntactically correct
- ▶ is well formatted
- ▶ may still be **logically wrong**

Formatting alone cannot guarantee correctness.

**Correct behavior is defined by function tests,  
not by appearance.**

### Test that enforces behavior

```
test_example.py
```

```
1 def test_add():  
2     assert add(2, 3) == 5
```

This test:

- ▶ executes the function
- ▶ checks the expected result
- ▶ fails automatically if behavior changes

The test can run:

- ▶ locally
- ▶ in Git hooks
- ▶ in CI pipelines

## CI/CD Tools: Local Discipline and Correctness

### Version control

- ▶ **Git**

Tracks changes to source code and configuration files, enabling branching, merging, and full reconstruction of development history.

- ▶ **GitHub / GitLab**

Hosting platforms for Git repositories, providing collaboration, access control, and CI integration.

### Local enforcement

- ▶ **Git hooks**

Local scripts executed automatically on Git events (e.g. pre-commit) to enforce rules before code is committed.

- ▶ **pre-commit framework**

Version-controlled hook manager that ensures identical checks across developers and CI environments.

### Code quality and correctness

- ▶ **Black**

Automatic Python formatter that enforces a single, deterministic code style without configuration decisions.

- ▶ **pytest**

Python testing framework that executes test functions and fails automatically when expected behavior is violated.

These tools:

- ▶ run fast
- ▶ provide immediate feedback
- ▶ prevent trivial errors from propagating

**Local discipline reduces error rates before CI even starts.**

# CI/CD Tools: Automation and Execution

## CI orchestration

### ▶ GitHub Actions

Event-driven CI system integrated into GitHub, executing workflows defined in YAML on managed runners.

### ▶ GitLab CI

Pipeline-based CI system configured via `.gitlab-ci.yml`, supporting self-hosted and specialized runners (e.g. HPC, GPU).

### ▶ Jenkins

Standalone automation server using scripted pipelines, common in legacy and enterprise environments.

## Execution environments

### ▶ Python virtual environments

Isolate Python dependencies to ensure reproducible runtime behavior.

### ▶ Containers (Docker, Apptainer)

Package applications and dependencies into portable, reproducible execution units across systems.

## Artifacts

### ▶ CI artifacts and registries

Store outputs such as logs, test reports, trained models, and container images produced during pipelines.

**CI platforms enforce rules; environments make them reproducible.**

## Why Data Formats Matter in AI/ML

### Characteristics of training data

- ▶ Very large (GB–TB scale)
- ▶ Multi-dimensional (time, space, channels)
- ▶ Often produced continuously

In contrast to classical workflows:

- ▶ Data rarely fits into memory
- ▶ Full sequential reads are uncommon

### Typical ML access patterns

- ▶ Repeated sampling of small subsets
- ▶ Random or structured access
- ▶ Parallel reading by many workers

As a result, the **data format** directly affects:

- ▶ I/O performance
- ▶ Scalability of training
- ▶ Feasibility of distributed ML

**Data formats are part of the ML infrastructure, not just storage.**

## Zarr: Why It Is Attractive for AI/ML

### Design principles

- ▶ Chunked, array-based storage
- ▶ Designed for cloud and HPC
- ▶ Partial reads without full downloads

### Why this helps ML training

- ▶ Efficient random sampling
- ▶ Scales to distributed workers
- ▶ Reduces I/O bottlenecks

### Performance features

- ▶ Parallel reads of independent chunks
- ▶ Flexible chunk layout
- ▶ Compression per chunk

### Typical use cases

- ▶ Images and video
- ▶ Satellite and geospatial data
- ▶ Scientific simulation output

**Zarr optimizes data access patterns — not the ML logic itself.**

## Zarr: Limitations and When to Use Alternatives

### Limitations

- ▶ Limited native integration with ML frameworks
- ▶ Metadata overhead for many small arrays
- ▶ Not well suited for tabular data
- ▶ Care needed for concurrent writes

Zarr adds complexity that is not always justified.

### When to choose alternatives

- ▶ Tabular ML data: Parquet , Arrow
- ▶ TensorFlow pipelines: TFRecord
- ▶ PyTorch streaming: WebDataset
- ▶ Small datasets: NetCDF, HDF5, NPY

### Guiding principle

- ▶ Choose formats based on data structure
- ▶ Optimize for access pattern, not fashion

**Zarr is powerful when the problem matches the format.**

## Demo Forecast Dataset (Europe, 5 variables)

### Dataset design

We generate a small **synthetic forecast archive** over Europe:

- ▶ grid: lat × lon (regular lat/lon)
- ▶ time axis: `valid_time` for lead times (e.g. 0–120h)
- ▶ variables:
  - ▶ `t2m` (2m temperature)
  - ▶ `u10, v10` (10m wind)
  - ▶ `mslp` (mean sea-level pressure)
  - ▶ `tp` (precipitation)

We mimic the **access patterns** of real NWP:

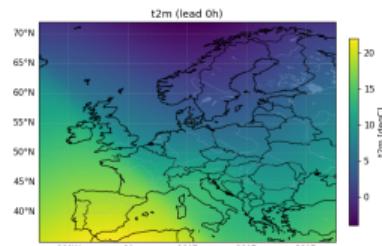
- ▶ map snapshot for one lead time
- ▶ town / station time series extraction
- ▶ local patches for ML training

### Notebook outputs

- ▶ Zarr store:  
`data/demo_eu_forecast.zarr`
- ▶ first plots: Europe maps
- ▶ metadata: coordinates + attrs

### Next

- ▶ generate fields with moving patterns
- ▶ write to Zarr with chunking



Lead Time: 0h

## Synthetic Forecast Generation (moving patterns)

### Idea

We generate physically plausible  
spatio-temporal structure :

- ▶ coherent “weather patterns” move eastward with time
- ▶ correlated variables:
  - ▶ pressure wave → wind via spatial gradients
  - ▶ temperature advected + noise
  - ▶ precipitation linked to fronts (thresholded patterns)

```
# grid (Europe)
lon2d, lat2d = np.meshgrid(lon, lat)
# moving phase (eastward shift with lead time)
phase = kx * lon2d + ky * lat2d - omega * lead_hours[t]

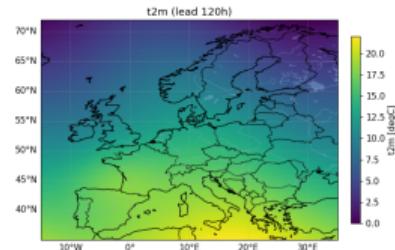
# temperature: advected pattern
t2m[t] = 10.0 + 6.0*np.cos(phase + 0.7) + 0.8*rng.normal(size=lon2d.shape)
```

### Result

Fields are realistic enough to demonstrate:

- ▶ map snapshots
- ▶ point extraction
- ▶ patch-based ML sampling

### Example snapshot



Lead Time: 120h

## Writing the Zarr Archive (chunking for point extraction)

### Zarr write: key decisions

We write the dataset as a **chunked archive**:

- ▶ **time chunk = 1:** fast town time series extraction
- ▶ **lat/lon chunks = 60:** moderate patch size for ML sampling

### Chunk layout

- ▶ chunks stored as independent files
- ▶ town extraction reads **few chunks only**
- ▶ patches match ML mini-batches

We generate chunks that contain exactly **one lead time** and a  **$60 \times 60$  spatial patch**.

### Example:

`(valid_time, lat, lon) = (1, 60, 60)`

```
# -----
# Zarr write
# -----
OUTDIR = "data"
os.makedirs(OUTDIR, exist_ok=True)

ZARR_PATH = os.path.join(OUTDIR, "demo_eu_forecast.zarr")

# best practice: chunk so point extraction is fast
# - valid_time chunk = 1 (fast timeseries reading)
# - lat/lon chunks ~ moderate
chunked = ds.chunk({"valid_time": 1, "lat": 60, "lon": 60})

# overwrite if exists
if os.path.exists(ZARR_PATH):
    import shutil
    shutil.rmtree(ZARR_PATH)

chunked.to_zarr(ZARR_PATH, mode="w", consolidated=True)
print("Wrote:", ZARR_PATH)
```

# Open Zarr Archive and Inspect Dataset

## Open archive

We open the Zarr store as an `xarray.Dataset`:

- ▶ lazy loading via `dask.array`
- ▶ chunk layout is preserved:  
`(valid_time, lat, lon) = (1, 60, 60)`

## Dataset structure

- ▶ dimensions: `valid_time=41, lat=121, lon=181`
- ▶ variables: `t2m, u10, v10, mslp, tp`
- ▶ coordinate: `lead_time` as `timedelta`

## Key observation

Even though the archive is chunked, `xarray` provides a `single logical dataset` interface.

## Example:

```
ZARR_PATH = "data/demo_eu_forecast.zarr"  
ds = xr.open_zarr(ZARR_PATH, consolidated=True)
```

```
xarray.Dataset  
-> Dimensions: (valid_time: 41, lat: 121, lon: 181)  
- Coordinates:  
  valid_time (valid_time) datetime64[ns] 2026-01-10 ... 2026-01-15  
  lat (lat) float64 35.0 35.31 35.62 ... 71.69 72.0  
  lon (lon) float64 -15.0 -14.72 -14.44 ... 34.72 35.0  
  init_time () datetime64[ns] ...  
  lead_time (valid_time) timedelta64[s] dask.array<chunksize={1}, meta=np....  
- Data variables:  
  mslp (valid_time, lat, lon) float32 dask.array<chunksize={1, 60, 60}, m...  
  t2m (valid_time, lat, lon) float32 dask.array<chunksize={1, 60, 60}, m...  
  tp (valid_time, lat, lon) float32 dask.array<chunksize={1, 60, 60}, m...  
  u10 (valid_time, lat, lon) float32 dask.array<chunksize={1, 60, 60}, m...  
  v10 (valid_time, lat, lon) float32 dask.array<chunksize={1, 60, 60}, m...  
- Attributes:  
  title : Synthetic European Forecast Demo  
  institution : Demo dataset for Zarr tutorial  
  history : Created in Notebook 04
```

Dataset preview: dimensions, coords, variables, and Dask chunks.

## Town Forecast Extraction (nearest grid point)

### Offline town database

We use a small built-in town table:

- ▶ no internet required
- ▶ town → (lat,lon)

### Point extraction

From the full dataset we extract a single point forecast :

- ▶ use nearest neighbor selection:  
`ds.sel(..., method="nearest")`
- ▶ result: dataset with only dimension `valid_time`

### Example:

```
TOWNS = {  
    "Berlin": (52.5200, 13.4050), ... }  
  
pt = ds.sel(lat=lat0, lon=lon0, method="nearest")  
  
xarray.Dataset  
  
Dimensions: (valid_time: 41)  
Coordinates:  
  valid_time (valid_time) datetime64[ns] 2026-01-10 ... 2026-01-15  
  init_time () datetime64[ns] ...  
  lat (0) float64 52.58  
  lead_time (valid_time) timedelta64[1s] dask.array<chunksize=(1,), meta=np.ndarray>  
  lon (0) float64 13.33  
Data variables:  
  mslp (valid_time) float32 dask.array<chunksize=(1,), meta=np.ndarray>  
  t2m (valid_time) float32 dask.array<chunksize=(1,), meta=np.ndarray>  
  tp (valid_time) float32 dask.array<chunksize=(1,), meta=np.ndarray>  
  u10 (valid_time) float32 dask.array<chunksize=(1,), meta=np.ndarray>  
  v10 (valid_time) float32 dask.array<chunksize=(1,), meta=np.ndarray>  
Attributes:  
  title : Synthetic European Forecast Demo  
  institution : Demo dataset for Zarr tutorial  
  history : Created In Notebook 04
```

### Output

A compact dataset with time series for: `t2m`, `u10`, `v10`, `mslp`, `tp`

Point forecast: only dimension `valid_time` remains.

## Town Forecast Dynamics (build time series table)

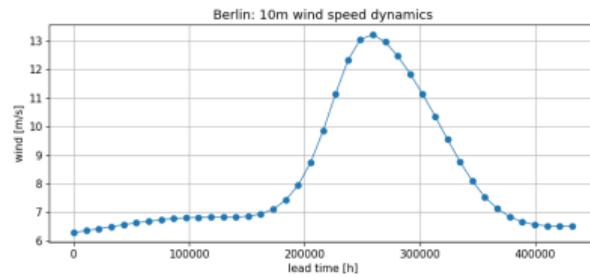
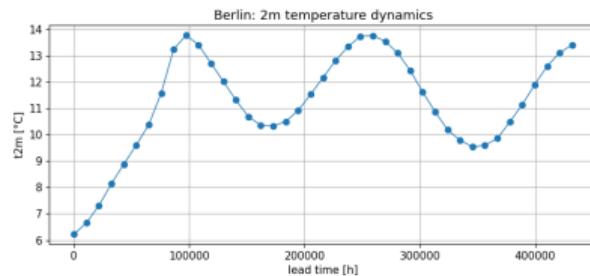
### Prepare time series data

We convert the extracted point dataset into a clean table:

- ▶ convert times to datetime
- ▶ convert lead time to hours
- ▶ compute derived variables (e.g. wind speed)

```
# -----
# Create a clean timeseries table
#
lead_h = (pt["lead_time"].values / np.timedelta64(1, "h")).astype(int)

df = pd.DataFrame({
    "valid_time": pd.to_datetime(pt["valid_time"].values),
    "lead_h": lead_h,
    "t2m_C": pt["t2m"].values,
    "wind_ms": np.sqrt(pt["u10"].values**2 + pt["v10"].values**2),
    "mslp_hPa": pt["mslp"].values,
    "tp_mm": pt["tp"].values,
})
```



## Local Map Plot (OSM basemap + 100 km temperature overlay)

Visualize a local patch around the town:

- ▶ **OpenStreetMap (OSM)** background tiles
- ▶ overlay  $t_{2m}$  for one lead time
- ▶ show only a 100 km box for clarity and speed

### Key steps

- ▶ select lead time index (e.g. 24h)
- ▶ slice  $t_{2m}$  in a local lon/lat window
- ▶ reproject to WebMercator (EPSG:3857)

```
t2m_local = ds["t2m"].isel(valid_time=i).sel(  
    lat=slice(latc-dlat, latc+dlat),  
    lon=slice(lonc-dlon, lonc+dlon))  
)  
  
# plot on OSM background (contextily)  
ctx.add_basemap(ax, source=ctx.providers.OpenStreetMap.Mapnik, zoom=10)
```



Example: Berlin,  $t_{2m}$  at lead 24h  
(local 100 km window).

## Training a CNN: field(t) -> field(t+1)

### Learning task

We train a CNN to predict next-step temperature:

- ▶ input at time t:  $t_{2m}$ ,  $u10$ ,  $v10$ ,  $mslp$
- ▶ target:  $t_{2m}(t+1)$

### Patch training

Instead of full fields, we train on random patches:

- ▶ sample 64x64 patches from random positions
- ▶ efficient mini-batches on laptop/GPU
- ▶ same idea as Zarr chunks: local spatial tiles

### Normalization

Channel-wise mean/std over time and space:

- ▶ stabilize training
- ▶ same scaling for training and evaluation

### Example:

```
# inputs and target in RAM
t2m = ds["t2m"].astype("float32").values
u10 = ds["u10"].astype("float32").values
v10 = ds["v10"].astype("float32").values
mslp = ds["mslp"].astype("float32").values

# normalization
Xin = np.stack([t2m, u10, v10, mslp], axis=1)
X_mean = Xin.mean(axis=(0,2,3), keepdims=True)
X_std = Xin.std(axis=(0,2,3), keepdims=True) + 1e-6

y_mean = t2m.mean()
y_std = t2m.std() + 1e-6

# patch sampler: X(t) -> y(t+1)
Xt, yt = sample_patch_batch(batch_size=8, patch=64)

# small CNN: 4ch -> 1ch
model = SmallCNN(in_ch=4, out_ch=1, hidden=32)
loss_fn = nn.MSELoss()
opt = optim.Adam(model.parameters(), lr=2e-3)
```

## Training Loop and Outcome (loss curve)

Example:

### Training loop

We train on randomly sampled spatial patches:

- ▶ each step draws a new batch: location + time
- ▶ objective: MSE between predicted and true  $t2m(t+1)$
- ▶ optimizer: Adam

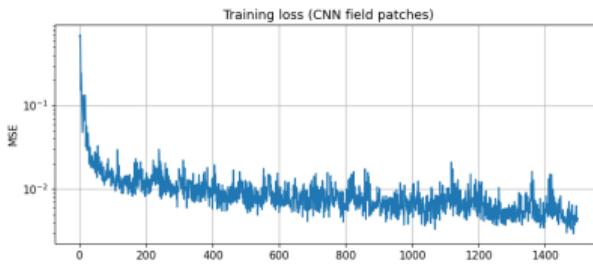
```
rng = np.random.default_rng(123)
n_steps = 800
batch_size = 8

loss_hist = []

model.train()
for step in range(1, n_steps + 1):
    Xt, yt = sample_patch_batch(batch_size=batch_size, rng=rng)
    pred = model(Xt)
    loss = loss_fn(pred, yt)

    opt.zero_grad()
    loss.backward()
    opt.step()

    loss_hist.append(float(loss.item()))
```



## Evaluation: rollout forecast from initial time

### One-step skill

We test the CNN on random samples:

- ▶ predict  $t2m(t+1)$  from  $X(t)$
- ▶ compare to truth with RMSE
- ▶ compare to persistence baseline

### Rollout forecast

Starting from the initial forecast field:

- ▶ use CNN repeatedly:  $\hat{x}_{t+1} = f_\theta(\hat{x}_t)$
- ▶ build a full forecast trajectory
- ▶ then extract town series for comparison

This gives a realistic "mini NWP" demonstration:

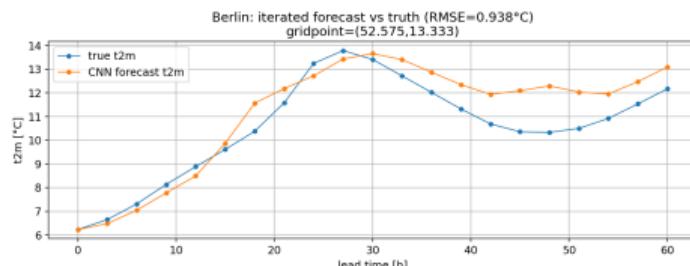
Zarr archive -> ML training -> forecast rollout .

### Example (rollout):

```
# rollout: start from true initial field at t0
t2m_hat = t2m[0].copy()

pred_series = [t2m_hat]
for k in range(0, nt-1):
    Xk = build_features(t2m_hat, u10[k], v10[k], mslp[k])
    t2m_hat = predict_next_field(model, Xk)
    pred_series.append(t2m_hat)

pred_series = np.stack(pred_series, axis=0) # (time, lat, lon)
```



## Outcome: Town Forecasts from ML Rollout

### Town-level evaluation

After rollout, we extract time series at town locations:

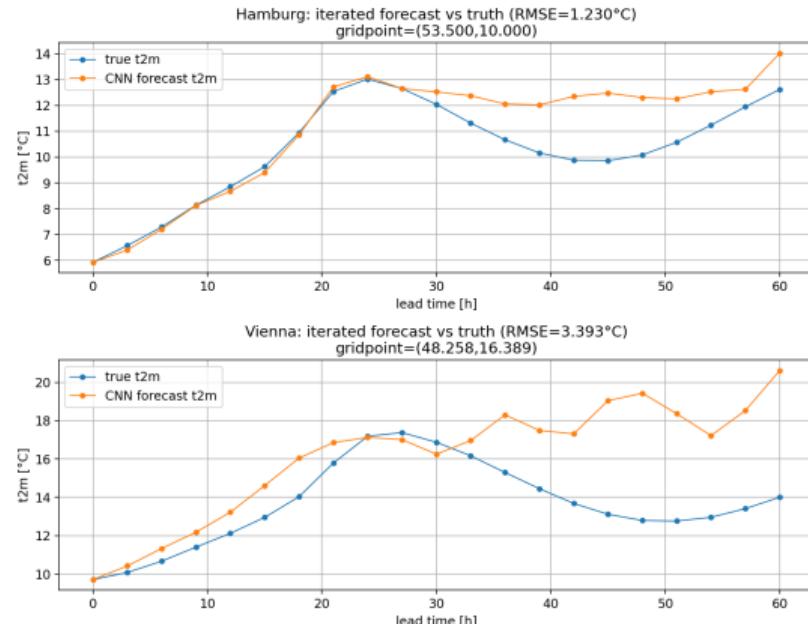
- ▶ nearest neighbor on the lat/lon grid
- ▶ compare truth vs ML forecast
- ▶ generate one plot per town

### Saved outputs

- ▶ figures: `fc_<town>.png`
- ▶ fast way to compare skill across cities

```
for town, (lat0, lon0) in TOWNS.items():
    # nearest grid point
    pt_true = ds.sel(lat=lat0, lon=lon0, method="nearest")

    # predicted series at same grid point
    y_true = pt_true["t2m"].values[:n_forecast+1]
    y_pred = pred_series[:, iy, ix]  # from rollout
```



## Zarr Chunk Decomposition (patch-only visualization)

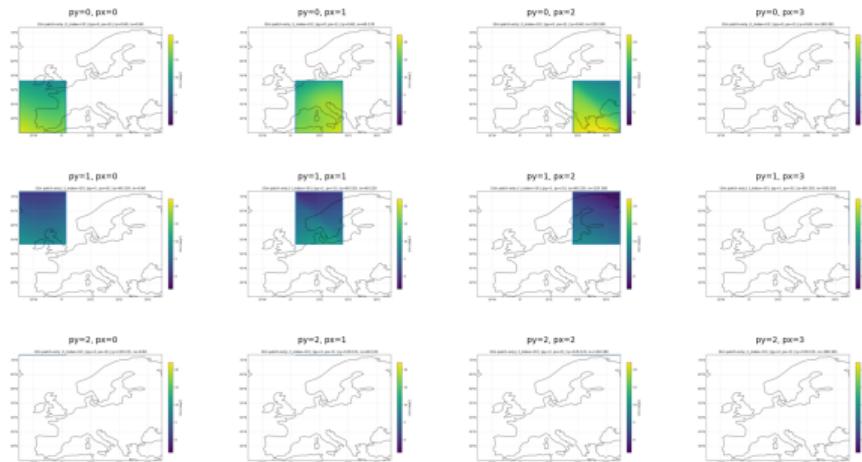
### Goal

Visualize the Zarr chunk structure :

- ▶ each chunk corresponds to a spatial patch (60x60)
- ▶ plot patch values only (no full-field replot)
- ▶ use one global color scale (vmin/vmax of full field)

### Procedure

- ▶ load full field once (for vmin/vmax)
- ▶ loop over patch indices ( $py, px$ )
- ▶ save one image per patch



Patch-only plots:  $py=0..2$ ,  $px=0..3$  (Europe extent, global color scale).