

# Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀  
playing 🤖 with  
Python 🐍 and AI/ML!  
 

## Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	<b>Lab:</b> Python Startup: Basics	<b>Lab:</b> Feed-forward and Graph NNs	<b>Lab:</b> Graph Learning with PyTorch	<b>Lab:</b> Containers and Reproducibility	<b>Lab:</b> Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	<b>Lab:</b> Work environments, Python everywhere	<b>Lab:</b> Simple Transformer and LLM Use	<b>Lab:</b> Agent Frameworks	<b>Lab:</b> CI/CD Pipelines	<b>Lab:</b> Graph-based Assimilation
12:45–13:30	<b>Lunch Break</b>				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	<b>Anemoi:</b> AI-based Weather Modelling	AI and Physics
14:30–15:30	<b>Lab:</b> GRIB, NetCDF and Obs Visualisation	<b>Lab:</b> RAG Pipeline	<b>Lab:</b> DAWID Exploration	<b>Lab:</b> Anemoi Training Pipeline	<b>Lab:</b> Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	<b>The AI Transformation</b>	Learning from Observations Only
16:15–17:00	<b>Lab:</b> Torch Tensors and First Neural Net	<b>Lab:</b> Radar, SAT and Multimodal Data	<b>Lab:</b> MLflow Hands-on	<b>Lab:</b> How work style could change	<b>Lab:</b> ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

## Why Anemoi?

### Weather forecasting is a large-scale problem

- ▶ modern forecasts rely on very large state vectors
- ▶ global and regional fields with millions of grid points
- ▶ high temporal resolution and long time series

### Training data is massive

- ▶ global reanalyses such as ERA5, ICON-DREAM
- ▶ convection-permitting simulations (e.g. Arome, ICON-Force)
- ▶ multi-variable, multi-level, multi-year datasets

These datasets quickly reach terabyte scale.

### What this implies for ML

- ▶ training must be parallel and distributed
- ▶ data access must be chunked and efficient
- ▶ geometry and grid structure must be respected

### Role of Anemoi

- ▶ parallelizes fields and samples
- ▶ integrates data, graphs, models, and training
- ▶ supports the full ML lifecycle for weather models

Anemoi is not a model — it is a framework.

## Design Philosophy of Anemoi

### Core principles

- ▶ configuration over hard-coded logic
- ▶ clear separation of responsibilities
- ▶ scalable by construction

Anemoi is designed such that:

- ▶ data handling,
- ▶ graph construction,
- ▶ model definition,
- ▶ training orchestration

are **independent but composable** components.

### Declarative workflows

- ▶ experiments are defined in **YAML**
- ▶ components are selected, not programmed
- ▶ changes are traceable and reproducible

### Operational mindset

- ▶ same setup works on laptop, HPC, and cloud
- ▶ parallelism is explicit, not accidental
- ▶ full provenance of data and models

**Anemoi treats ML experiments as engineering systems.**

## YAML: The Configuration Backbone

### Why configuration matters at scale

- ▶ experiments involve many interacting components
- ▶ parameters change more often than code
- ▶ reproducibility depends on explicit configuration

In Anemoi, YAML files define :

- ▶ datasets and preprocessing
- ▶ model architectures and hyperparameters
- ▶ training strategies and resources

### Why YAML?

- ▶ human-readable and versionable
- ▶ hierarchical and structured
- ▶ easy to override and compose

### Design choice

- ▶ no Python code for experiment logic
- ▶ no hidden defaults in scripts
- ▶ configuration becomes a **first-class artifact**

**In Anemoi, changing the experiment means changing YAML.**

## YAML in Practice: Declaring an Experiment

A single YAML file describes an experiment

- ▶ what data is used
- ▶ which model is trained
- ▶ how training is executed

No experiment logic is hidden in Python code.

The configuration is the experiment.

This YAML file fully specifies one training run.

Example: training configuration (simplified)

```
1 dataset:  
2   name: era5  
3   variables: [t2m, u10, v10]  
4   resolution: 0.25  
5  
6 model:  
7   name: graph_transformer  
8   hidden_dim: 256  
9   num_layers: 6  
10  
11 training:  
12   batch_size: 4  
13   max_epochs: 50  
14   accelerator: gpu
```

## OmegaConf: From YAML to Runtime Objects

### The problem

- ▶ YAML files are static text
- ▶ training code needs structured objects
- ▶ configuration must be inspectable at runtime

### Properties of DictConfig

- ▶ hierarchical (mirrors YAML structure)
- ▶ accessed via attributes or keys
- ▶ supports type checking

### OmegaConf solves this

- ▶ represents configuration as Python objects
- ▶ preserves hierarchy and structure
- ▶ supports interpolation and defaults

### Why this matters

- ▶ training code stays generic
- ▶ behaviour is fully configuration-driven
- ▶ experiments become reproducible by construction

The result is a `DictConfig` object.

**Configuration becomes part of the program state.**

# OmegaConf in Practice and in Anemoi

## Using OmegaConf in Python

- ▶ configuration is passed into the training code
- ▶ parameters are accessed programmatically
- ▶ no hard-coded values are needed

### Accessing the configuration

```
1 # cfg provided by Hydra
2 print(cfg.model.layer_sizes)
3 print(cfg.training.max_epochs)
```

### Resolving and converting

```
1 from omegaconf import
    OmegaConf
2
3 OmegaConf.resolve(cfg)
4 cfg_dict = OmegaConf.
    to_object(cfg)
```

### Role in Anemoi

- ▶ schemas validate configs before training
- ▶ resolved configs are logged with results
- ▶ the full configuration defines the experiment

**OmegaConf links configuration and execution.**

## Hydra: Managing Variants of the Same Experiment

### Motivation

- ▶ same training code
- ▶ same task (e.g. learning  $\sin(x)$ )
- ▶ different model architectures

Hydra allows:

- ▶ selecting model variants via YAML
- ▶ switching architectures without code changes
- ▶ keeping experiments comparable

### Base configuration

- ▶ defines which components are active
- ▶ serves as the experiment entry point

#### config.yaml

```
1 defaults:  
2   - model: mlp  
3   - training: simple  
4   - _self_
```

#### model/mlp.yaml

```
1 layer_sizes: [1, 64, 1]  
2 activation: relu
```

#### model/deep.yaml

```
1 layer_sizes: [1, 128, 64, 32,  
               1]  
2 activation: tanh
```

The experiment is composed, not rewritten.

## Hydra in Practice: Same Code, Different Outcomes

### What changes between these runs

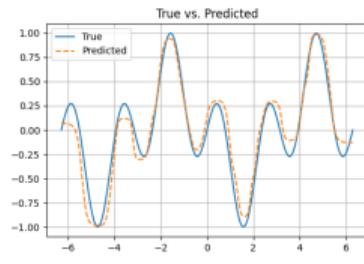
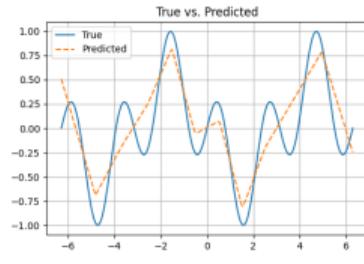
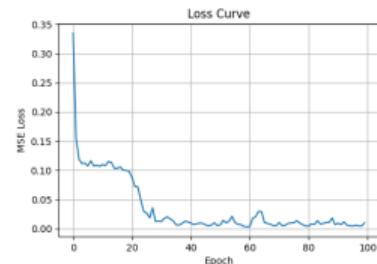
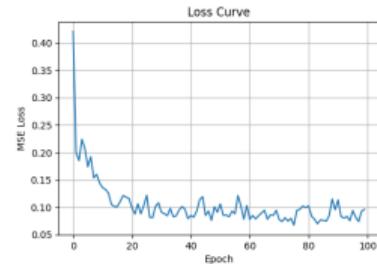
- ▶ only the **Hydra configuration**
- ▶ different model definitions (`mlp` vs `deep`)
- ▶ same training loop, same data

### Hydra controls:

- ▶ which YAML files are composed
- ▶ which model architecture is instantiated
- ▶ which hyperparameters are active

### Key observation

- ▶ shallow model converges smoothly
- ▶ deeper model shows slower, less stable training



**Hydra makes architectural choices explicit, comparable, and reproducible.**

## Why Graphs in Weather Machine Learning?

### Weather data is not Cartesian

- ▶ global models use spherical geometry
- ▶ grids are often **non-uniform and unstructured**
- ▶ classical CNN assumptions break down

Examples:

- ▶ ICON triangular grid
- ▶ reduced Gaussian grids
- ▶ observation networks

### Why graphs are a natural abstraction

- ▶ nodes represent spatial locations
- ▶ edges represent physical neighbourhoods
- ▶ locality is explicit, not implicit

### Key idea

- ▶ move from **array indices** to **connectivity**
- ▶ geometry becomes part of the model

**Graphs decouple geometry from resolution.**

## Icosahedral and Geodesic Graphs

### ICON as a guiding principle

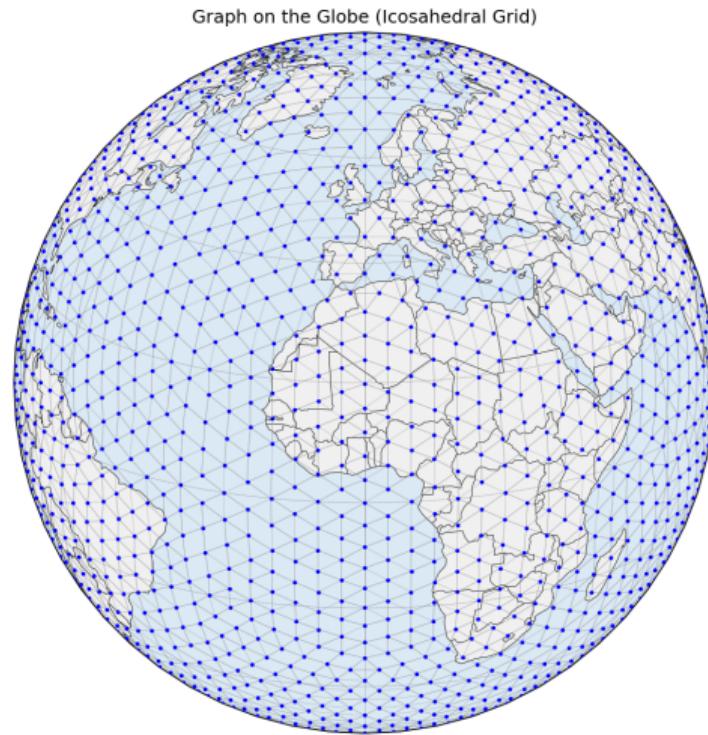
- ▶ refined icosahedral grids
- ▶ nearly uniform cell areas
- ▶ no singularities at the poles

Anemoi adopts the same idea:

- ▶ nodes placed on a refined icosahedron
- ▶ spherical geometry is explicit
- ▶ resolution controlled by refinement level

### Key trade-off

- ▶ higher resolution  $\Rightarrow$  more nodes
- ▶ increased cost, but better spatial fidelity



Icosahedral graph on the sphere (orthographic projection). Edges follow great-circle distances.

## Edges, Neighbours, and Graph Topology

### Defining neighbourhoods

- ▶ edges encode spatial interaction
- ▶ most common choice:  
k-nearest neighbours (kNN)
- ▶ distance measured on the sphere

Each node exchanges information only with its neighbours:

- ▶ local interactions
- ▶ scalable message passing

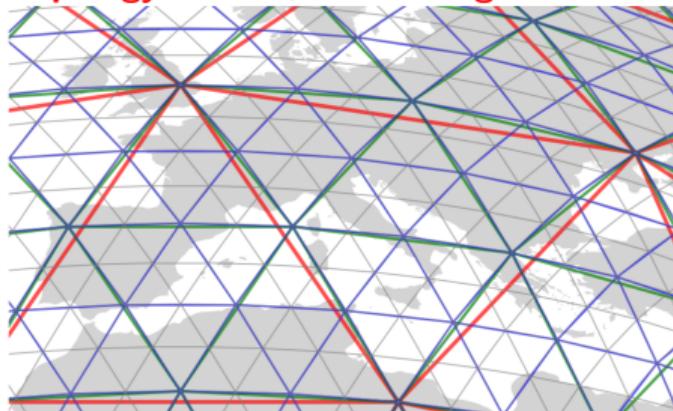
### Edge index representation

- ▶ graph stored as `edge_index`
- ▶ integer pairs  $(i, j)$  define connections
- ▶ independent of data values

### Reuse across experiments

- ▶ graph built once
- ▶ reused for training, validation, inference
- ▶ ensures geometric consistency

**Topology is fixed; data changes.**



## Anemoi: A Modular ML Framework for Weather

### Design philosophy

- ▶ separate concerns cleanly
- ▶ make each component replaceable
- ▶ scale from experiments to operations

Anemoi is not a single package, but a coordinated ecosystem .

Each package addresses one layer:

- ▶ data handling
- ▶ graph construction
- ▶ model definition
- ▶ training orchestration

### Core Anemoi packages

- ▶ `anemoi-datasets`
- ▶ `anemoi-graphs` in `anemoi-core` repo
- ▶ `anemoi-models` in `anemoi-core` repo
- ▶ `anemoi-training` in `anemoi-core` repo

All packages:

- ▶ use YAML + Hydra + OmegaConf
- ▶ integrate via clearly defined interfaces
- ▶ are developed in a shared monorepo

**Graphs connect geometry; packages connect workflows.**

## anemoi-datasets: Data and Metadata

### Purpose

- ▶ prepare large meteorological datasets for ML
- ▶ unify metadata, statistics, and structure
- ▶ decouple raw data formats from training

### Key features

- ▶ ingestion from GRIB / NetCDF
- ▶ conversion to Zarr
- ▶ automatic statistics (mean, std, min, max)
- ▶ validation and consistency checks

<https://anemoi.readthedocs.io/projects/datasets/en/latest/>

### Why this matters

- ▶ scalable I/O for HPC and cloud
- ▶ reproducible data pipelines
- ▶ identical datasets across experiments

**Data becomes a first-class, validated object.**

	shortName	Name	Levels
0	P	Pressure	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]
1	QV	Specific Humidity	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]
2	T	Temperature	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]
3	U	U-Component of Wind	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]
4	V	V-Component of Wind	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]
5	W	Vertical Velocity (Geometric) (w)	[101, 108, 112, 119, 120, 49, 57, 64, 70, 75, 79, 86, 91, 96]

## anemoi-graphs: Geometry and Connectivity

### Purpose

- ▶ represent spatial geometry explicitly
- ▶ construct reusable graph topologies
- ▶ support non-Cartesian grids

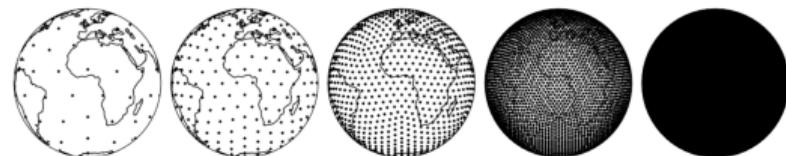
### Design principle

- ▶ graph built once
- ▶ reused for training and inference
- ▶ geometry independent of data values

### Key components

- ▶ TriNodes from refined icosahedra
- ▶ neighbour definitions via kNN
- ▶ edge construction on the sphere

Topology is fixed; learning happens on it.



Docs: <https://anemoi.readthedocs.io/projects/graphs/en/latest/>

## anemoi-models: Neural Architectures

## Purpose

- ▶ define ML models for weather prediction
  - ▶ separate architecture from training logic
  - ▶ support graph-based learning

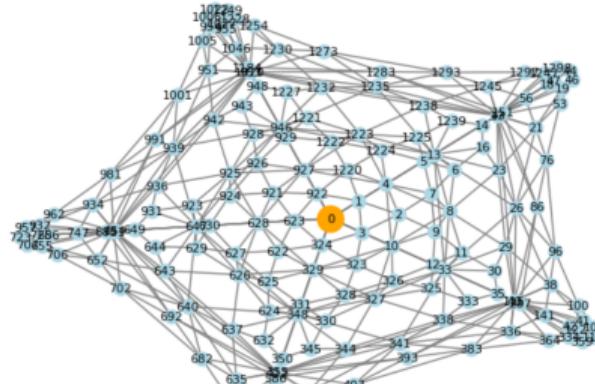
## Model families

- ▶ graph neural networks
  - ▶ transformer-based architectures
  - ▶ hybrid models for spatio-temporal data

Interfaces

- ▶ standardized input/output conventions
  - ▶ integration with PyTorch Lightning
  - ▶ instantiated via Hydra configs

## Models are components, not scripts.



## anemoi-training: Orchestrating Experiments

### Purpose

- ▶ orchestrate the full training lifecycle
- ▶ integrate data, graphs, and models
- ▶ scale to HPC environments

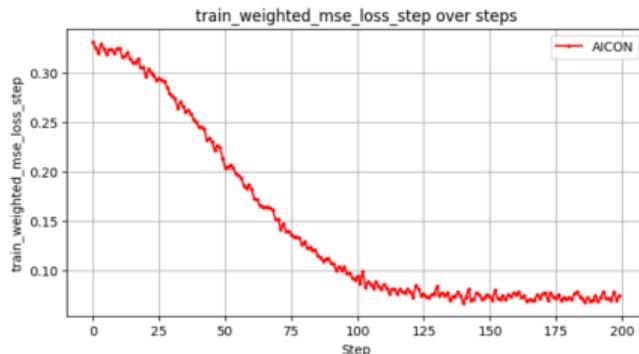
### Key elements

- ▶ AnemoiTrainer as central controller
- ▶ PyTorch Lightning backend
- ▶ distributed and multi-node training

### Experiment control

- ▶ driven entirely by configuration
- ▶ automatic logging and checkpointing
- ▶ reproducible, restartable runs

Training becomes declarative and scalable.



Docs: <https://anemoi.readthedocs.io/projects/training/en/latest/>

## Creating Zarr Datasets with anemoi-datasets

- ▶ convert raw meteorological data into ML-ready format
- ▶ standardize structure, metadata, and statistics
- ▶ enable scalable training and validation

### Key idea

- ▶ dataset creation is fully configuration-driven
- ▶ no preprocessing logic is hard-coded
- ▶ every dataset is reproducible

### Typical processing steps:

- ▶ read GRIB or NetCDF input
- ▶ reshape and flatten grids if required
- ▶ compute global statistics
- ▶ write chunked Zarr archive

<https://anemoi.readthedocs.io/projects/datasets/en/latest/>

### Command-line workflow

```
1 anemoi-datasets create \
2   dataset.yaml \
3   era5.zarr
```

### Result

- ▶ validated Zarr dataset
- ▶ embedded metadata and statistics
- ▶ directly usable by Anemoi training

**Data preparation becomes a reproducible pipeline.**

## Validating and Inspecting Zarr Datasets

### Why validation matters

- ▶ training assumes consistent metadata
- ▶ missing values or shapes cause silent errors
- ▶ errors should be caught **before** training

Anemoi provides built-in tools to:

- ▶ verify temporal coverage
- ▶ check spatial resolution and shapes
- ▶ ensure statistics are present

Validation is part of the dataset lifecycle, not an afterthought.

Docs: [anemoi-datasets](#)

### Inspecting a dataset

```
1 anemoi-datasets inspect  
      era5.zarr
```

Typical output includes:

- ▶ time range and frequency
- ▶ variables and dimensions
- ▶ min / max / mean / std
- ▶ total size and chunking

**Only validated datasets enter training.**

## ERA Data Stored in Zarr

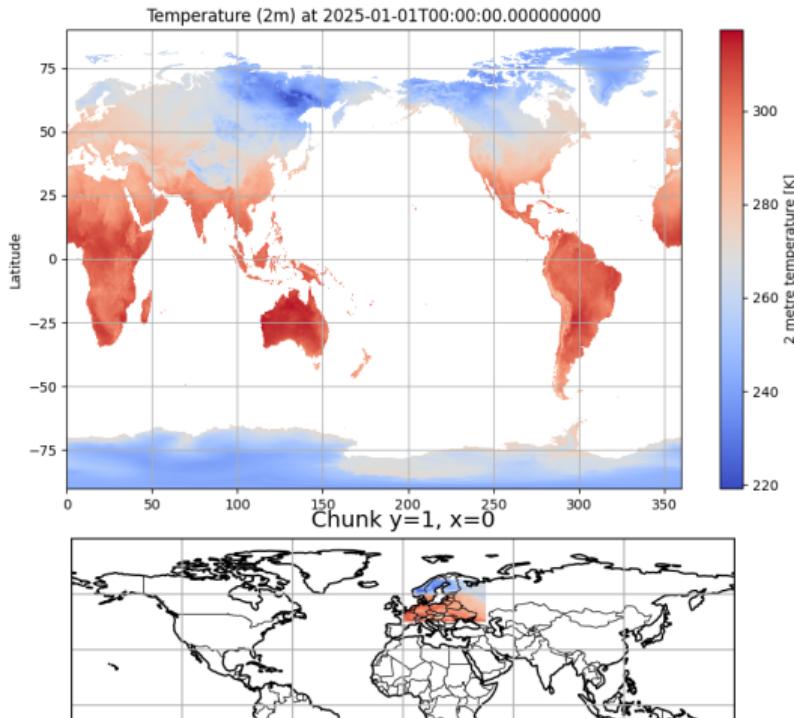
### What these images show

- ▶ ERA reanalysis data after Zarr conversion
- ▶ accessed chunk-wise, not as monolithic files
- ▶ identical physical content, different storage logic

Zarr enables:

- ▶ parallel access to spatial subdomains
- ▶ efficient mini-batching for ML
- ▶ scalable training on HPC systems

**Zarr is an ML-optimized view  
of reanalysis data.**



Top: Global ERA5 2 m temperature field.  
Bottom: Spatial chunk extracted from Zarr archive.

## Application Example: Learning CBCF Warning Polygons from Forecast Fields

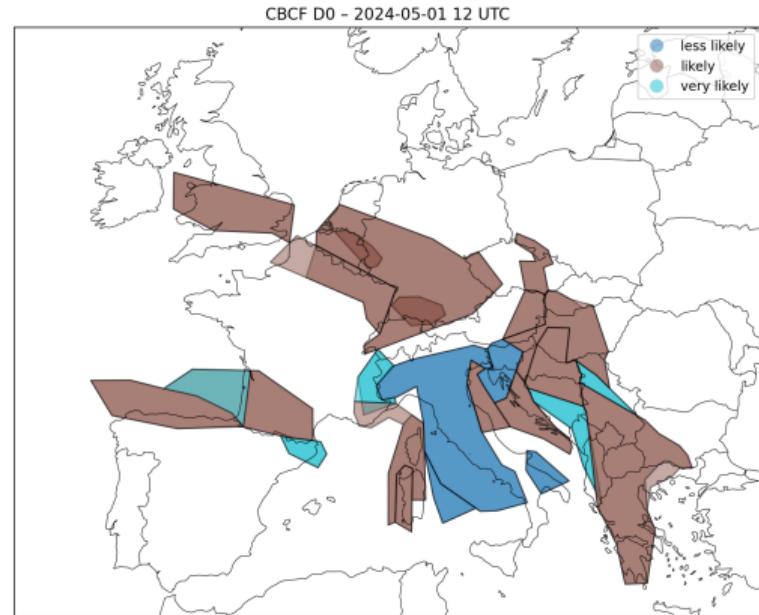
Learn an AI model that predicts operational warning regions (CBCF polygons) from model forecast fields.

### Training data

- ▶ multi-year archive of forecast fields (e.g. wind/gust)
- ▶ forecaster-issued warning polygons

### Core idea

- ▶ polygons are a human product representation
- ▶ neural nets learn on grids / tensors
- ▶ therefore we convert polygons → mask labels



Example CBCF polygons (likelihood categories).

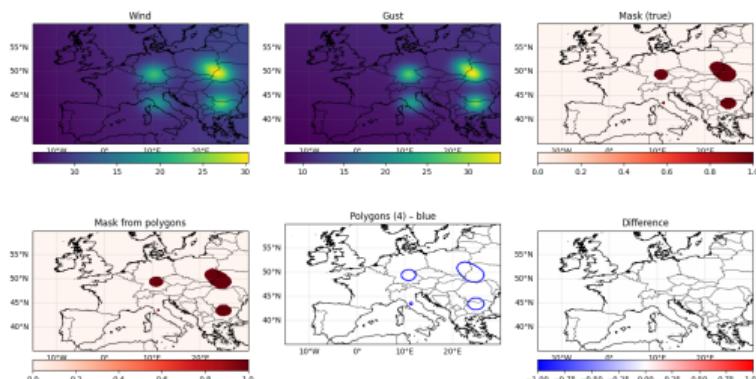
We train on masks — and recover polygons afterwards.

## From Polygons to Masks: Training Labels on the Model Grid

### Why convert polygons to masks?

- ▶ operational warnings are defined as **polygons**
- ▶ neural nets require **grid-aligned targets**
- ▶ rasterization yields a binary/soft **mask**  
 $m(x) \in [0, 1]$

- ▶ predictors: **wind** + **gust** fields
- ▶ target: warning mask derived from polygons



### Two-way conversion

- ▶ **Polygon → mask:** rasterize onto model grid
- ▶ **Mask → polygon:** contour / region extraction

### Quality control

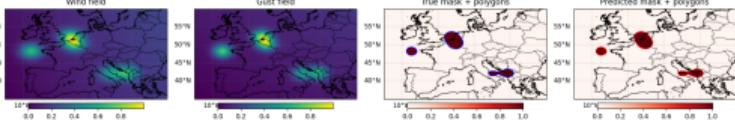
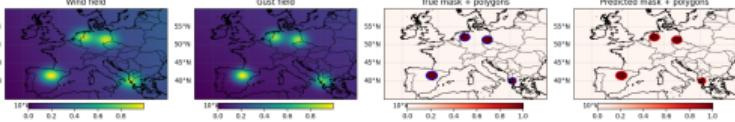
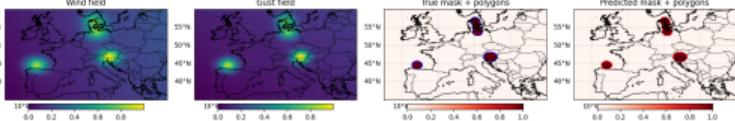
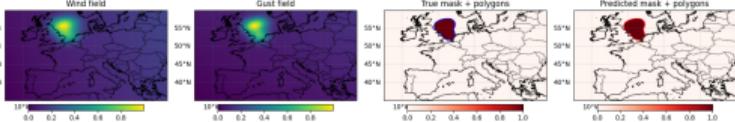
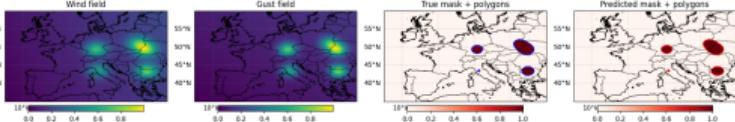
- ▶ check **mask(polygons)** against reference labels
- ▶ difference highlights discretization artefacts

Polygon-to-mask conversion and consistency check.

## Learning Demo: Inputs (Wind/Gust) → True vs Predicted Mask & Polygons

This figure shows one training/evaluation example

- ▶ predictors are gridded forecast fields
- ▶ target is a warning mask derived from forecaster polygons
- ▶ predicted mask is transformed back into polygons



4-column result: inputs, target, prediction.

## Real Demo: CBCF Training with Anemoi (Repository Setup)

We ran the CBCF demo using an Anemoi-based training stack.

### Local repo structure

- ▶ anemoi-datasets (data handling, zarr)
- ▶ anemoi-transform (preprocessing)
- ▶ anemoi-core (graphs, models, training)
- ▶ training\_config/ (yaml configs)
- ▶ batch\_train.sh (run script)

### Key point

- ▶ this is not a standalone notebook hack
- ▶ it is a reproducible training workflow (configs + checkpoints)

### Dependencies (uv workspace)

- ▶ anemoi-training,
- anemoi-models, anemoi-graphs,
- ...
- ▶ aicon-catalog[create]

### Commands (as executed)

```
cd cbcf
git pull
ls training_config
# run: batch_train.sh
```

### Outcome

- ▶ trained checkpoints
- ▶ inference + plots

## Configuration-Driven Workflow: Training & Diagnosis YAML

- ▶ training is **fully reproducible** and versionable
- ▶ separation of concerns:
  - ▶ dataset / resolution (R3B5, R3B8, ...)
  - ▶ variables (inputs/targets)
  - ▶ model architecture choice
  - ▶ training schedule / precision / hardware

### Conceptual config blocks

```
dataset:  
  resolution: R03B05  
model:  
  type: graph/diagnoser  
training:  
  precision: 16-mixed  
outputs:  
  checkpoint: inference-last.ckpt
```

(Shown here only conceptually — actual files live in `training_config/` and include many more details.)

### Example config files

- ▶ `training_config/cbcf_graph_R3B5.yaml`
- ▶ `training_config/cbcf_diagnose_R3B5.yaml`

### Demo message

We can train CBCF models for  
**different ICON grids and resolutions** without changing  
code — only by switching configs.

## Inference Pipeline: Load Checkpoint → Predict → Plot

### What happens after training?

- ▶ load **checkpoint** (`inference-last.ckpt`)
- ▶ load forecast data from **zarr archive**
- ▶ build input tensor (multi-step input supported)
- ▶ run `model.predict_step(...)`

### Outputs

- ▶ predicted CBCF fields / masks
- ▶ comparison against targets
- ▶ saved plots per date + variable

### Code sketch (from `main.py`)

```
model = torch.load(ckpt)
model.eval()

ds = zarr.open("...CBCF.zarr")

input = build_tensor(...)
y_pred = model.predict_step(input)
```

`plot(y_pred); plot(target)`

The demo uses multi-step inputs and produces predicted/target plots.

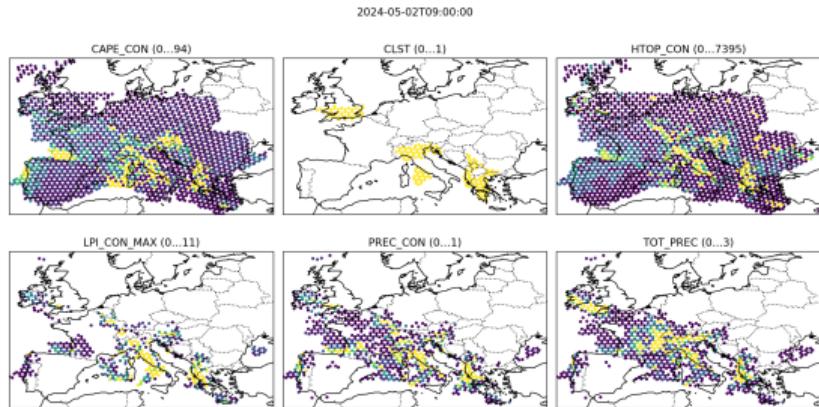
### Take-away

- ▶ same pipeline works for long archives
- ▶ basis for **operational verification and monitoring**

## CBCF Results: Anemoi Training Demo (Examples)

### What we show

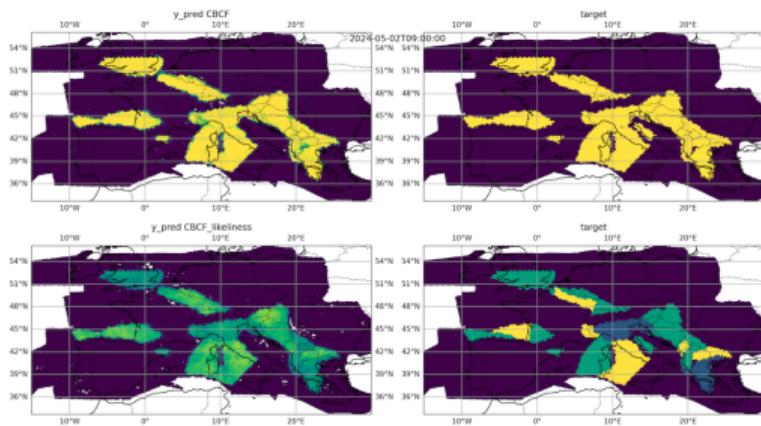
- ▶ results from the Anemoi-based training workflow
- ▶ prediction of CBCF target fields / masks from forecast inputs
- ▶ comparison against the corresponding targets



### Interpretation

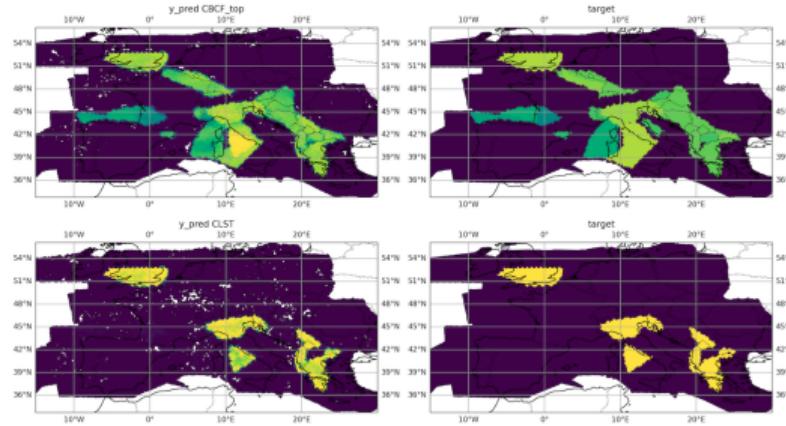
- ▶ spatial patterns are learned coherently
- ▶ overall structure matches operational products
- ▶ remaining differences motivate the next steps:
  - ▶ improved architectures
  - ▶ longer training / larger datasets
  - ▶ calibration and polygon-space verification

## CBCF Results: Predicted vs Target Fields (Selected Outputs)



Fields shown (left = prediction, right = target):

- ▶ CBCF: main warning mask/product
- ▶ CBCF\_likeliness: continuous likelihood / confidence field



Fields shown (left = prediction, right = target):

- ▶ CBCF\_top: top-level / dominant CBCF intensity field
- ▶ CLST: additional label channel (sparser warning type)