

Python and AI/ML for Weather, Climate and Environmental Applications

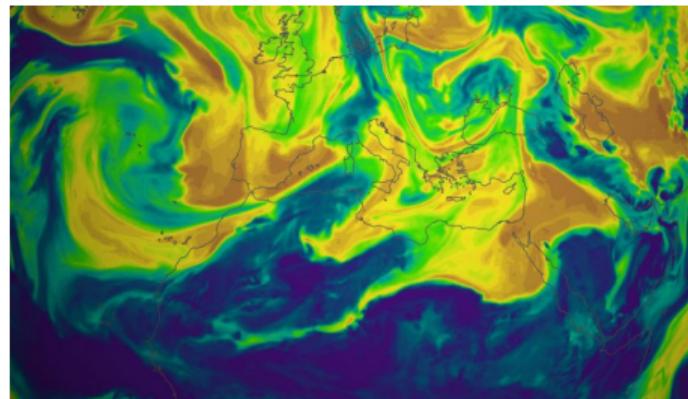


Let us enjoy 
 playing  with
 Python  and AI/ML!
 

Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	Lab: Python Startup: Basics	Lab: Feed-forward and Graph NNs	Lab: Graph Learning with PyTorch	Lab: Containers and Reproducibility	Lab: Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	Lab: Work environments, Python everywhere	Lab: Simple Transformer and LLM Use	Lab: Agent Frameworks	Lab: CI/CD Pipelines	Lab: Graph-based Assimilation
12:45–13:30	Lunch Break				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	Lab: GRIB, NetCDF and Obs Visualisation	Lab: RAG Pipeline	Lab: DAWID Exploration	Lab: Anemoi Training Pipeline	Lab: Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	The AI Transformation	Learning from Observations Only
16:15–17:00	Lab: Torch Tensors and First Neural Net	Lab: Radar, SAT and Multimodal Data	Lab: MLflow Hands-on	Lab: How work style could change	Lab: ORIGEN and Open Discussion
17:00–20:00					Joint Dinner

Why Emulators in Numerical Weather Prediction?



- ▶ High-resolution NWP is computationally expensive
- ▶ Many applications do not require full physical fidelity
- ▶ Fast forecasts enable new applications and scale

Emulator concept

- ▶ Learns the forecast step: state → state
- ▶ Operates within an existing NWP ecosystem
- ▶ Preserves grids, variables, and semantics

Goal of this lecture

- ▶ Understand emulator-based AI systems
- ▶ Compare AIFS and AICON
- ▶ Walk through AICON step by step

Emulators complement numerical models — they do not replace physical understanding.

Big-Tech AI Weather Models: Architectural Landscape

GraphCast (Google DeepMind)

- ▶ Graph Neural Network (GNN)
- ▶ Icosahedral grid, $\sim 0.25^\circ$
- ▶ Multi-mesh message passing
- ▶ Deterministic medium-range forecasts

Pangu-Weather (Huawei)

- ▶ 3D Transformer architecture
- ▶ Regular lat–lon grid, $\sim 0.25^\circ$
- ▶ Attention over space and vertical levels
- ▶ Image-like representation of atmosphere

GenCast (Google DeepMind)

- ▶ Diffusion model on graphs
- ▶ Icosahedral grid, $\sim 1^\circ\text{--}0.25^\circ$
- ▶ Probabilistic forecasting
- ▶ Explicit uncertainty representation

FourCastNet (NVIDIA)

- ▶ Fourier Neural Operator (FNO)
- ▶ Regular lat–lon grid, $\sim 0.25^\circ$
- ▶ Spectral convolution in Fourier space
- ▶ Extremely fast inference

All models replace the forecast step — none are embedded in classical NWP systems.

Research and Operational AI Weather Systems

Aurora (Microsoft)

- ▶ Transformer-based foundation model
- ▶ Regular lat–lon grids
- ▶ Multi-task Earth-system scope
- ▶ Research-driven, not NWP-native

AIFS (ECMWF)

- ▶ Graph Neural Network (Anemoi)
- ▶ ECMWF Flagship ML Model
- ▶ Trained on ECMWF reanalysis ERA
- ▶ Operational global forecast system

BRIS (Met Norway)

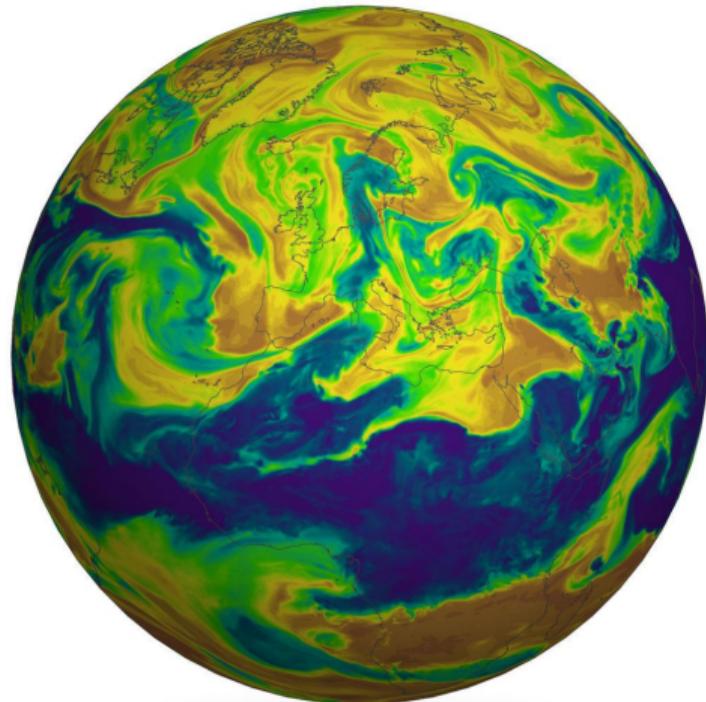
- ▶ Anemoi-based national system
- ▶ Nested Grid Approach, high resolution
- ▶ Independent training, AIFS-aligned
- ▶ Operational national usage

AICON (DWD)

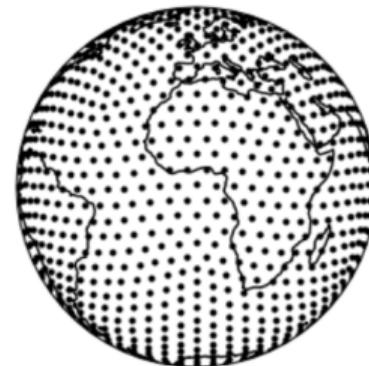
- ▶ Graph Neural Network (Anemoi)
- ▶ Native ICON triangular grid
- ▶ Trained on ICON-DREAM reanalysis
- ▶ Operational ICON emulator

European Meteorological Infrastructure is carrying out exciting development.

Anemoi based Emulators: Fields and Grids



Icosahedral Grid Representation



Triangular ICON mesh used as the computational graph for message passing and inference.

AICON learns a mapping on the ICON mesh, not on a Cartesian grid.

ICON-DREAM Reanalysis: High-quality training basis for AI emulators

Motivation

- ▶ reanalyses are key for **climate services** and many applications
- ▶ now also essential as **training basis** for AI-based NWP emulators

ICON reanalysis framework at DWD

- ▶ **ICON-DREAM:** global-to-regional ensemble reanalysis
- ▶ **ICON-FORCE:** 2.1 km ensemble reanalysis for Central Europe
- ▶ focus on **Europe** and on **continuous updates** (2010–present)

ICON-DREAM delivers **consistent, high-resolution data**

ICON-DREAM: key characteristics

- ▶ global ICON at **13 km**, 120 levels
- ▶ two-way nest over Europe at **6.5 km**
- ▶ DA: **LETKF** with **20-member ensemble** for *B*-covariances
- ▶ ensemble at 40 km (20 km over Europe)

ICON-FORCE (2 km):

- ▶ hourly LETKF (KENDA), incl. radar + SEVIRI
- ▶ 2-moment microphysics, snow / SST analyses

Outlook: backward extension to the 1980s using rescued obs + CDR products.

AICON Walkthrough: Environment Setup

Thanks to Florian Prill for a great Notebook!

Core ML stack

- ▶ Python ≥ 3.10
- ▶ PyTorch (CPU or CUDA)
- ▶ PyTorch Lightning
- ▶ PyTorch Geometric (PyG)

Anemoi framework

- ▶ anemoi-core (graphs, models, training)
- ▶ anemoi-datasets (Zarr + YAML I/O)
- ▶ anemoi-inference (generic inference)
- ▶ Supporting packages (e.g. transforms)

Data and configuration

- ▶ Zarr (chunked training data)
- ▶ Xarray (inspection and analysis)
- ▶ eccodes + earthkit (GRIB2 handling)
- ▶ Hydra / OmegaConf (YAML configuration)

Recommended setup

- ▶ Pinned requirements.txt
- ▶ Virtual environment or container
- ▶ Consistent library versions

Reproducibility depends on a stable, fully controlled software environment.

AICON Walkthrough: Manual Setup (Packages and ecCodes)

Install required Python packages

Pinned Python environment

```
1 # Install all required
   packages (version!)
2 %pip install --no-cache-
   dir -U -r ../
   requirements.txt
3
4 # Inspect installed
   versions
5 !pip freeze
```

Exact package versions are critical for PyTorch, PyG, Lightning, and Anemoi compatibility.

Configure ecCodes definitions

```
import os
from pathlib import Path

# Path to ICON/DWD ecCodes definitions
# (provided with the repo)
edzw_defs = Path("eccodes/definitions.edzw").resolve()

# Standard ecCodes definitions (system install)
std_defs = "/path/to/eccodes/definitions"

os.environ["ECCODES_DEFINITION_PATH"] = \
    f"{edzw_defs}:{std_defs}"
```

ICON GRIB decoding requires extended definition tables beyond standard ecCodes.

AICON Walkthrough: Notebook Initialization

Jupyter and runtime setup

Notebook basics

```
1 % Enable autoreload for
    iterative development
2 %load_ext autoreload
3 %autoreload 2
4
5 # Standard imports
6 import os
7 import numpy as np
8 import torch
```

Ensures code changes are picked up and the runtime environment is visible.

Initialization is not boilerplate — it defines the experimental contract.

Determinism and sanity checks

Reproducibility checks

```
1 # Fix random seeds
2 torch.manual_seed(42)
3 np.random.seed(42)
4
5 # Check PyTorch device
6 device = torch.device("cuda" if
    torch.cuda.is_available()
    else "cpu")
7 print("Using device:", device)
```

Deterministic behavior is essential for debugging and scientific comparison.

AICON training setup (Anemoi): config-driven workflow

Main idea

- ▶ training is
fully controlled by YAML configs
- ▶ Hydra composes config blocks:
 - ▶ data (Zarr), variables,
normalization
 - ▶ model architecture (AICON /
GNN)
 - ▶ trainer settings (epochs, batch,
strategy)

Workflow

- ▶ choose config (start from integration test)
- ▶ instantiate trainer
- ▶ run training (checkpoints + logger)

Notebook pattern (core lines)

```
config_filename = "test_aicon_01.yaml"  
  
# Hydra loads and resolves config  
config = load_config(config_filename)  
  
# Anemoi trainer wraps Lightning  
trainer = AnemoiTrainer(config)  
  
# run training  
trainer.train()
```

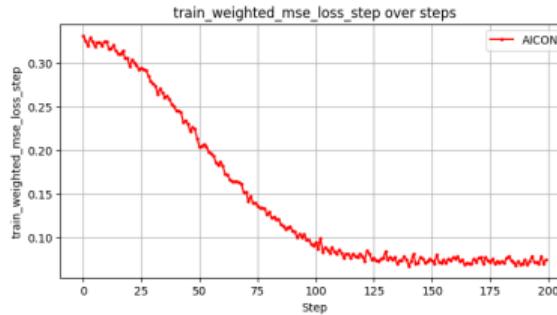
Outputs

- ▶ checkpoints: last.ckpt,
inference-last.ckpt
- ▶ logs: TensorBoard / MLflow-style metrics

Training diagnostics: loss curve and experiment logger

What we monitor

- ▶ training loss vs epoch/step
- ▶ validation loss (generalization)
- ▶ learning rate schedule
- ▶ walltime + throughput



Why it matters

- ▶ early detection of instability
- ▶ reproducible experiments
- ▶ compare model variants

Take-away: training is observable and traceable.

```
class GraphicalLiveLogger(Logger):
    def __init__(self, interval=10):
        super().__init__()
        self.metric_name = "train_weighted_mse_loss_step"
        self.metrics_history = []

    @property
    def name(self): return "LiveLogger"
    @property
    def version(self): return "0.1"

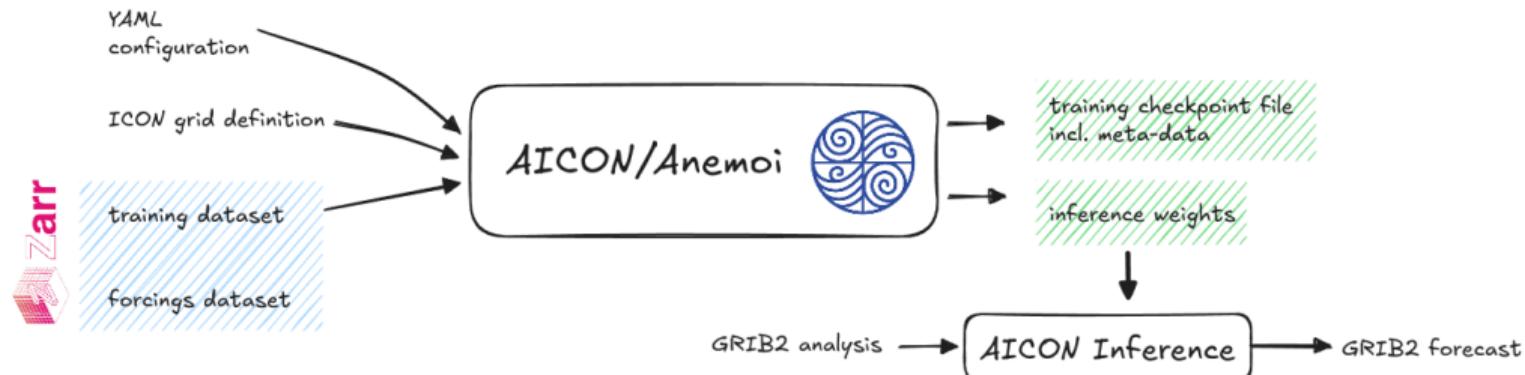
    def log_hyperparams(self, *args, **kwargs): pass

    def log_metrics(self, metrics, step):
        self.metrics_history.append((step, metrics))

    def finalize(self, status):
        steps, values = zip(*[s,m[self.metric_name]] for s,m in self.metrics_history if self.metric_name in m)
        plt.figure(figsize=(8,4))
        plt.plot(steps, values, 'o-', ms=2, label="AICON", color="red")
        plt.title(f"{self.metric_name} over steps"); plt.xlabel("Step"); plt.ylabel(self.metric_name)
        plt.grid(True); plt.legend(); plt.show()

trainer.model.logger_enabled = True
trainer.loggers = GraphicalLiveLogger()
```

AICON training workflow (overview)



Pipeline: ICON data → graph dataset → training config → trainer → checkpoints + logs

Training configuration: input and outputs

Input data (from config)

```
text
1 ICON mesh (NetCDF):
2   <.../icon_mesh.nc>
3
4 Training interval:
5   <YYYY-MM-DD HH> ... <YYYY-MM-
6   DD HH>
7 Data files:
8   <...> (joined list)
```

Output directories

```
text
1 Checkpoints:
2   <.../checkpoints/...>
3
4 Logs (incl. MLFlow):
5   <.../logs/...>
```

Take-away: config \Rightarrow reproducible paths and artifacts.

Key point: a single config fixes both data geometry
and time range .

Training dataset (Zarr) inspected as Xarray

Key idea

- ▶ dataset is a **4D tensor** :
(time, variable, ensemble, cell)
- ▶ efficient I/O via **chunking**
(Dask)
- ▶ includes metadata:
 - ▶ latitudes / longitudes
 - ▶ variable-wise statistics
(mean, std, min/max)

**Take-away: this is ML-ready
data with rich metadata.**

▶ Dimensions: (variable: 84, time: 724, ensemble: 1, cell: 11520)

▶ Coordinates: (0)

▼ Data variables:

count	(variable)	float64	dask.array<chunksize=(84,), me...		
data	(time, variable, ensemble, cell)	float32	dask.array<chunksize=(1, 84, 1,...		
dates	(time)	datetime64[s]	dask.array<chunksize=(724,), m...		
has_nans	(variable)	object	dask.array<chunksize=(84,), me...		
latitudes	(cell)	float64	dask.array<chunksize=(11520,),...		
longitudes	(cell)	float64	dask.array<chunksize=(11520,),...		
maximum	(variable)	float64	dask.array<chunksize=(84,), me...		
mean	(variable)	float64	dask.array<chunksize=(84,), me...		
minimum	(variable)	float64	dask.array<chunksize=(84,), me...		
squares	(variable)	float64	dask.array<chunksize=(84,), me...		
stdev	(variable)	float64	dask.array<chunksize=(84,), me...		
sums	(variable)	float64	dask.array<chunksize=(84,), me...		

Vertical levels: selected ICON layers in the dataset

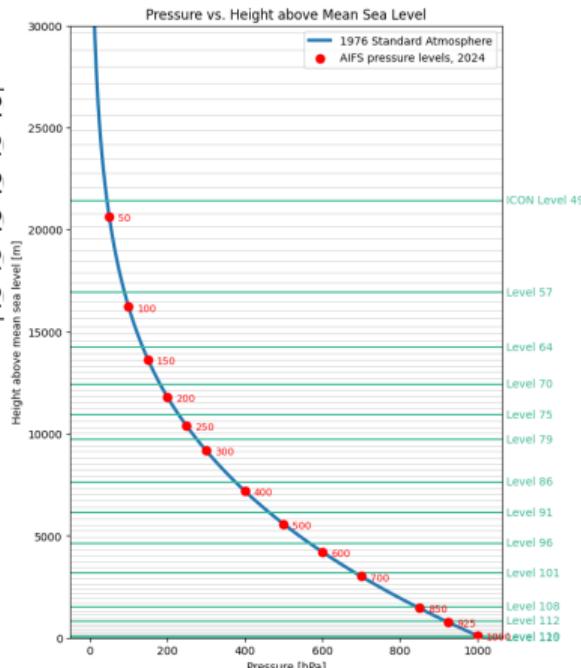
Variables on model levels

shortName	Name	Levels
P	Pressure	49,57,64,70,75,79,86,...,12
QV	Specific humidity	49,57,64,70,75,79,86,...,12
T	Temperature	49,57,64,70,75,79,86,...,12
U	U-component of wind	49,57,64,70,75,79,86,...,12
V	V-component of wind	49,57,64,70,75,79,86,...,12
W	Vertical velocity (geom.)	49,57,64,70,75,79,86,...,12

Forcings / static features (examples):

EMIS_RAD, FR_LAND, FR_LAKE, HSURF, Z0,
 insolation, SSO_*, sin/cos(latitude, longitude, local_time,
 julian_day)

Idea: multi-level state + forcings \Rightarrow complete ML input for AICON.



ICON level structure: full model vertical grid (selection highlighted).

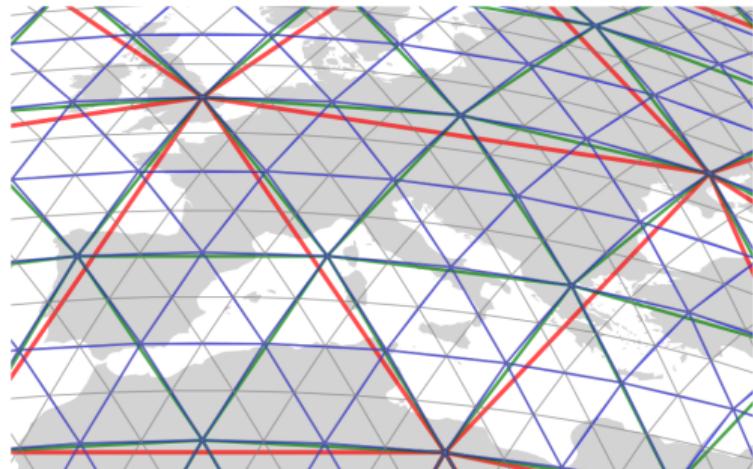
ICON multimesh: the hidden graph behind AICON

Key idea

- ▶ AICON does not run on a regular lat-lon grid
- ▶ it uses ICON's hierarchical triangular mesh
- ▶ the GNN operates on a hidden multi-mesh

Multi-mesh principle

- ▶ union of coarse-to-fine subgraphs
 $RnB0 \cup RnB1 \cup \dots \cup RnBk$
- ▶ mixes short-range + long-range edges
- ▶ similar spirit as GraphCast multi-mesh



Hidden mesh edges shown level-wise: coarse levels enable long-range interaction.

Take-away: the multi-mesh gives the GNN both local physics and global context .

AICON model: encoder – graph processor – decoder

Graph-to-graph forecasting

- ▶ input: state + forcings on ICON cells (nodes)
- ▶ output: next-step state (multi-var, multi-level)

Architecture blocks

- ▶ **Encoder:** maps raw variables to latent node/edge features
- ▶ **Processor:** message passing on the mesh graph (several layers)
- ▶ **Decoder:** projects latent features to physical output variables

Take-away: encoder/decoder makes the network variable-agnostic in latent space.

```
encoder:  
    _target_: anemoi.models.layers.mapper.GraphTransformerForwardMapper  
    _convert_: all  
    trainable_size: ${model.trainable_parameters.data2hidden}  
    sub_graph_edge_attributes: ${model.attributes.edges}  
    num_chunks: 2  
    cpu_offload: ${model.cpu_offload}  
    mlp_hidden_ratio: 4  
    num_heads: 16  
    qk_norm: false  
    layer_kernels:  
        Linear:  
            _target_: torch.nn.Linear  
            _partial_: true  
        Activation:  
            _target_: torch.nn.GELU  
decoder:  
    _target_: anemoi.models.layers.mapper.GraphTransformerBackwardMapper  
    _convert_: all  
    trainable_size: ${model.trainable_parameters.hidden2data}  
    sub_graph_edge_attributes: ${model.attributes.edges}  
    num_chunks: 1  
    cpu_offload: ${model.cpu_offload}  
    mlp_hidden_ratio: 4  
    num_heads: 16  
    qk_norm: false  
    initialise_data_extractor_zero: false  
    layer_kernels:  
        Linear:  
            _target_: torch.nn.Linear  
            _partial_: true
```

Encoder maps physics variables → latent graph; decoder maps latent → output fields.

Graph structure: local neighborhoods on the ICON mesh

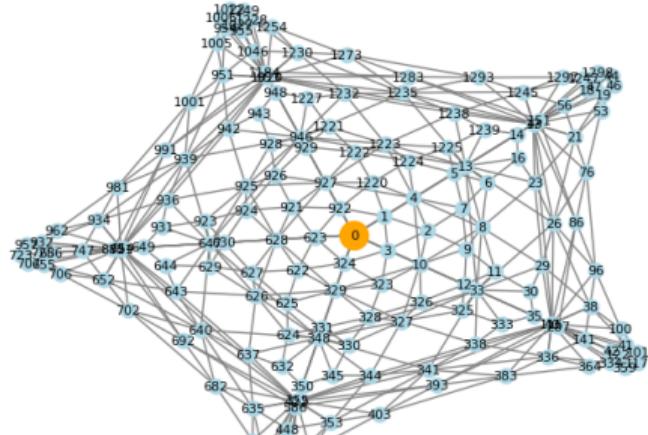
Why this matters

- ▶ AICON runs on the ICON grid as graph neural network (GNN)
- ▶ each ICON cell \Rightarrow one graph node
- ▶ edges connect mesh neighbors

Message passing view

- ▶ prediction at one node uses information from:
 - ▶ **1st neighbors** (directly connected)
 - ▶ **2nd neighbors**
(neighbors-of-neighbors)
- ▶ this defines the local receptive field of the GNN

Take-away: GNNs learn transport + interaction patterns through local connectivity .



Example: subgraph around one node (orange) with 1st and 2nd neighbors.

Graph Transformer attention: which neighbors matter?

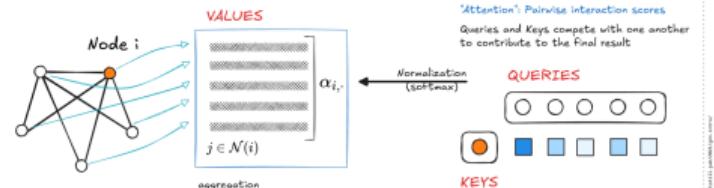
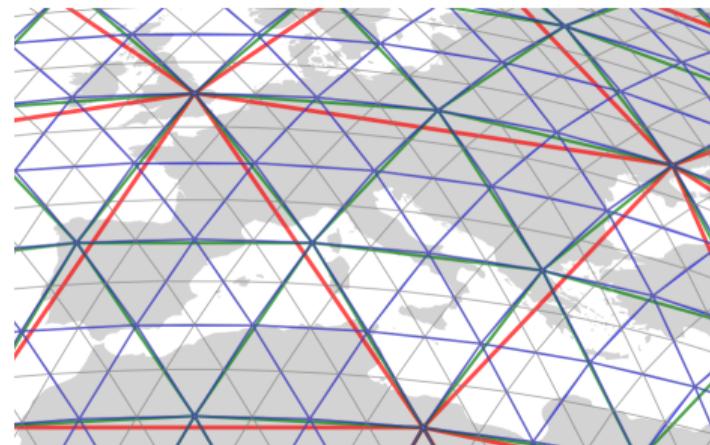
Graph Transformer vs. plain message passing

- ▶ classic GNN: neighbors contribute via fixed averaging / sum
- ▶ GraphTransformer: each edge gets a learned attention weight

Interpretation

- ▶ for one receiver node i , all incoming edges ($j \rightarrow i$)
- ▶ attention weights $\alpha_{i,j}$ sum to 1 (per head)
- ▶ thick edges \Rightarrow high influence of neighbor j

Take-away: the model learns where to look on the graph depending on flow regime and structure.



Example: (top) ICON multi-mesh graph (coarse + fine connections); (bottom) attention setup

Transfer learning: reuse weights on a finer hidden mesh

Motivation

- ▶ training on the finest multi-mesh is expensive
- ▶ idea: **train on coarse mesh** and transfer to finer mesh

What transfers well?

- ▶ most encoder/processor/decoder weights
- ▶ learned graph attention patterns are largely mesh-agnostic

What must be adapted?

- ▶ graph-dependent parameters:
 - ▶ hidden node coordinates / embeddings
 - ▶ extra trainable node/edge attributes
- ▶ these are re-initialized and fine-tuned on the new mesh

Example (AICON walkthrough)

- ▶ change hidden mesh resolution:
 - ▶ `max_level_multimesh: 3 → 4`
- ▶ load weights only + enable transfer learning

Config idea (YAML)

```
training:  
  load_weights_only: true  
  transfer_learning: true  
  run_id: <pretrained_run_id>
```

```
graph:  
  nodes:  
    icon_mesh:  
      node_builder:
```

Operational inference at DWD: why a dedicated tool?

Goal

- ▶ run AICON forecasts **operationally** and **reproducibly**
- ▶ integrate into existing NWP infrastructure (GRIB2, workflows, monitoring)

Key ingredients

- ▶ **checkpoint:** inference-last.ckpt
- ▶ **input:** ICON / reanalysis data streams
- ▶ **output:** GRIB2 forecast products

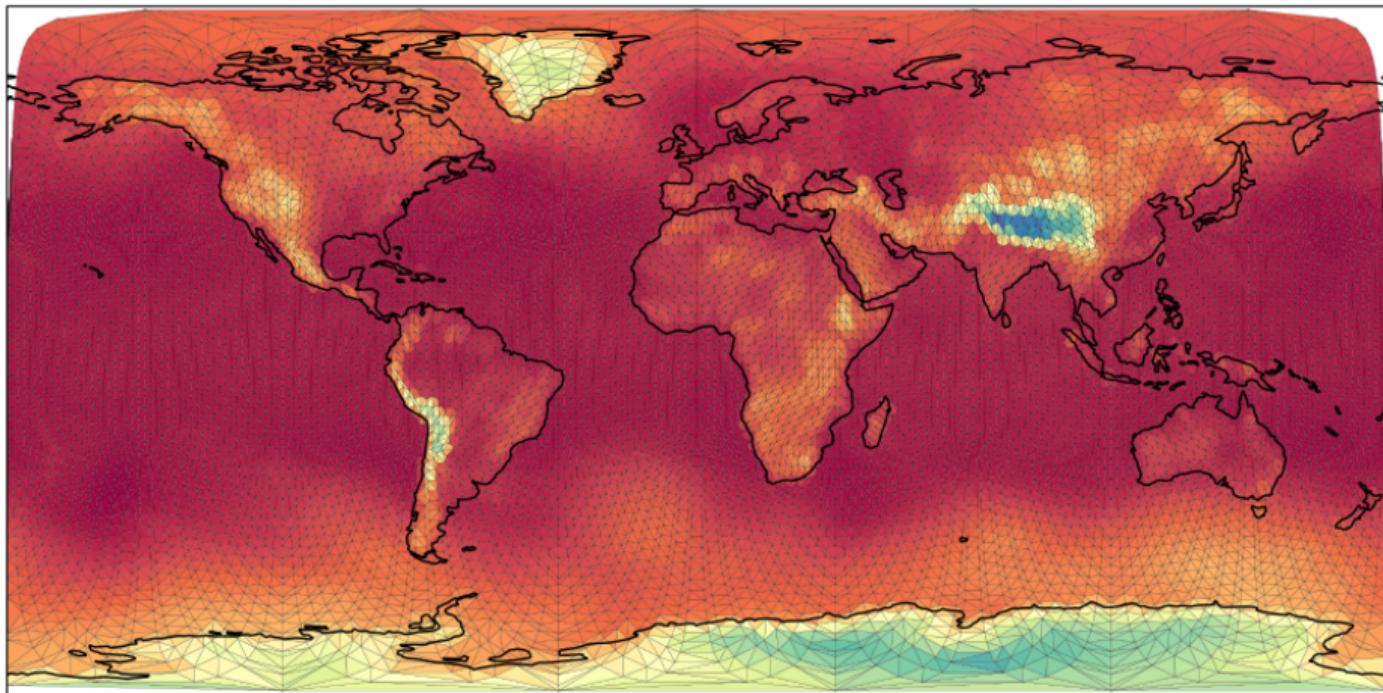
Operational packaging

- ▶ DWD-specific I/O and preprocessing (GRIB2 conventions)
- ▶ additional input engineering (e.g. soil moisture index etc.)
- ▶ packaging as stable runtime for ops (container)
- ▶ hooks for monitoring and verification pipelines

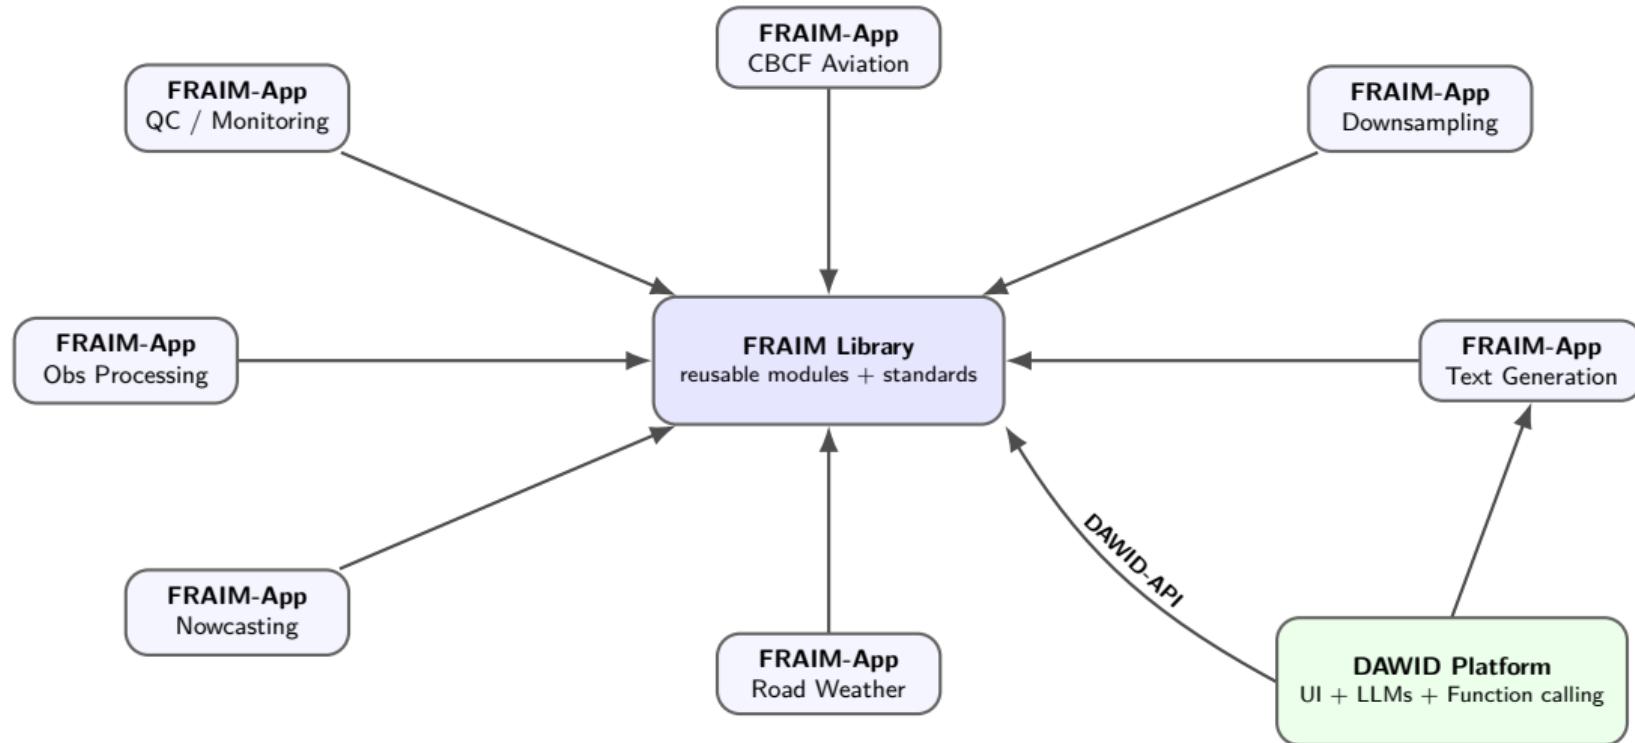
- ▶ Apptainer / Singularity container
- ▶ fixed paths to constants + weights
- ▶ versioning: container + config + checkpoint

**Take-away: DWD inference tool =
bridge from ML model to ops .**

Demo: AICON forecast from operational inference



FRAIM Ecosystem: Central Library, FRAIM-Apps, and DAWID Platform



Anemoi vs. FRAIM — What are they about?

Anemoi (ECMWF & Partners)

- ▶ international collaboration to build ML-based forecast models
- ▶ end-to-end framework for training & inference
- ▶ well-defined pipeline:
datasets → graphs/models → training → inference/deploy
- ▶ Open Source, pan-European community

Typical question:

How do we train and operate an ML forecasting model operationally?

FRAIM (E-AI / DWD / Partners)

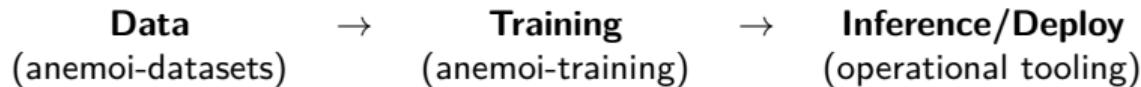
- ▶ international collaboration focused on many AI applications
- ▶ method toolbox / modular framework across a broad portfolio: products, services, and smaller AI components
- ▶ platform/integration view: standards, building blocks, reuse
- ▶ umbrella framework (forecasting is only one use case)

Typical question:

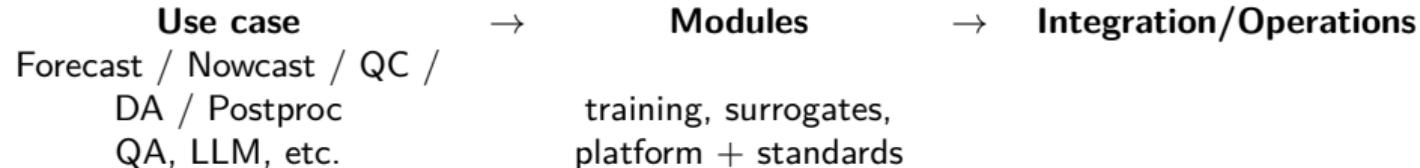
How do we build a modular AI ecosystem for many meteorological products and services?

Pipeline view: where does each framework sit?

Anemoi: lifecycle of an ML forecast model



FRAIM: system / toolbox view across multiple AI components



Take-away:

Anemoi is **forecast-model-centric**, FRAIM is **architecture- & reuse-centric**.

Concrete differences and a good integration story

Differences (short & concrete)

- ▶ **Scope:** Anemoi = end-to-end ML forecasting
FRAIM = modular AI toolbox for many meteorological products
- ▶ **Core artifacts:**
 - ▶ **Anemoi:** datasets, graphs, model weights, training pipelines
 - ▶ **FRAIM:** central library + reusable modules + standards
- ▶ **Organization principle:**
 - ▶ Anemoi: one coherent forecasting stack
 - ▶ FRAIM: many use-case driven FRAIM-Apps (each in its own repo)

How does this combine well?

- ▶ FRAIM can integrate mfa or Anemoi as a forecasting or processing engine
- ▶ FRAIM then provides:
 - ▶ product-specific pipelines (QC, DA, verification, monitoring, etc.)
 - ▶ deployment patterns and operational standards

One-liner:

Anemoi = forecast-model factory;
FRAIM = modular product ecosystem.

DAWID — LLM Platform, Tools, and FRAIM Integration

DAWID capabilities

- ▶ **LLM user interface**
 - ▶ chat + document context (RAG)
 - ▶ role-based workflows (science / dev / ops)
- ▶ **Function calling**
 - ▶ controlled execution of domain functions
 - ▶ structured I/O and provenance
- ▶ **Agents (multi-step)**
 - ▶ plan → call tools → validate

DAWID-API integration

- ▶ **Unified API to LLMs and tools**
 - ▶ multi-LLM backends
 - ▶ tool/function endpoints
- ▶ **Link to FRAIM-Apps**
 - ▶ FRAIM-App → DAWID-API: reasoning + tooling
 - ▶ consistent interface across products

**Take-away: DAWID connects
LLMs + tools + FRAIM-Apps .**

Key idea: DAWID is the **interactive AI cockpit** .