

Python and AI/ML for Weather, Climate and Environmental Applications



Let us enjoy 🚀
playing 🤖 with
Python 🐍 and AI/ML!
 

Five-Day Schedule Overview

Time	Day 1	Day 2	Day 3	Day 4	Day 5
09:00–10:00	Opening by ECMWF DG, Start: Coding & Science in the Age of AI	Neural Network Architectures	Diffusion and Graph Networks	MLOps Foundations	Model Emulation, AIFS and AICON
10:00–11:00	Lab: Python Startup: Basics	Lab: Feed-forward and Graph NNs	Lab: Graph Learning with PyTorch	Lab: Containers and Reproducibility	Lab: Emulation Case Studies
11:00–12:00	Python, Jupyter and APIs	Large Language Models	Agents and Coding with LLMs	CI/CD for Machine Learning	AI-based Data Assimilation
12:00–12:45	Lab: Work environments, Python everywhere	Lab: Simple Transformer and LLM Use	Lab: Agent Frameworks	Lab: CI/CD Pipelines	Lab: Graph-based Assimilation
12:45–13:30	Lunch Break				
13:30–14:30	Visualising Fields and Observations	Retrieval-Augmented Generation (RAG)	DAWID System and Feature Detection	Anemoi: AI-based Weather Modelling	AI and Physics
14:30–15:30	Lab: GRIB, NetCDF and Obs Visualisation	Lab: RAG Pipeline	Lab: DAWID Exploration	Lab: Anemoi Training Pipeline	Lab: Physics-informed Neural Networks
15:30–16:15	Introduction to AI and Machine Learning	Multimodal Large Language Models	MLflow: Managing Experiments	The AI Transformation	Learning from Observations Only
16:15–17:00	Lab: Torch Tensors and First Neural Net	Lab: Radar, SAT and Multimodal Data	Lab: MLflow Hands-on	Lab: How work style could change	Lab: ORIGEN and Open Discussion
17:00–20:00				Joint Dinner	

Lecture 20 — Obs-to-Obs Learning on a 2D Toy Atmosphere

Goal of this lecture

- ▶ Build a simple 2D dynamical system $\phi(x, z, t)$ with transport + diffusion + source
- ▶ Generate two observation types
 - ▶ **Radiosondes (RS):** sparse vertical profiles at a few columns
 - ▶ **Satellite (SAT):** integrated vertical weighted observations for every column
- ▶ Train a neural network to predict next-step observations

$$(y_t^{sat}, y_t^{rs}) \mapsto (y_{t+1}^{sat}, y_{t+1}^{rs})$$

- ▶ Reconstruct the full field at $t + 1$ by querying RS predictions everywhere

Key message: learn a forecast step in observation space, and still recover a state-like field.

Toy Dynamics: Advection–Diffusion with Source and Damping

We simulate a tracer/heating field $\phi(x, z, t)$ on a 2D domain:

$$(x, z) \in [0, L_x] \times [0, L_z].$$

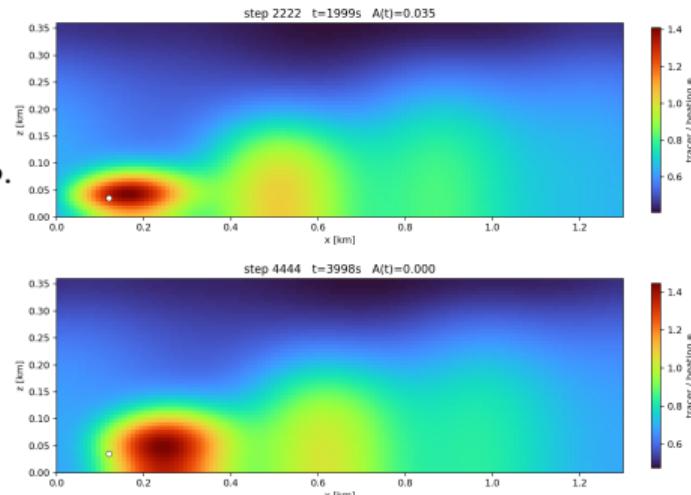
Dynamics (PDE)

$$\frac{\partial \phi}{\partial t} = -u \frac{\partial \phi}{\partial x} - w \frac{\partial \phi}{\partial z} + K \nabla^2 \phi + A(t) S(x, z) - \lambda \phi - \lambda_{\text{top}}(z) \phi.$$

Key design choices

- ▶ Mean wind: (u, w) right + upward \Rightarrow visible transport
- ▶ Source $S(x, z)$: localized bottom-left heating region
- ▶ Pulsed forcing $A(t)$ \Rightarrow visible trace stripes
- ▶ Damping λ + top sponge λ_{top} \Rightarrow equilibrium

Time stepping: RK4.



Boundary Conditions: Wrap-Around Transport (Periodic in x)

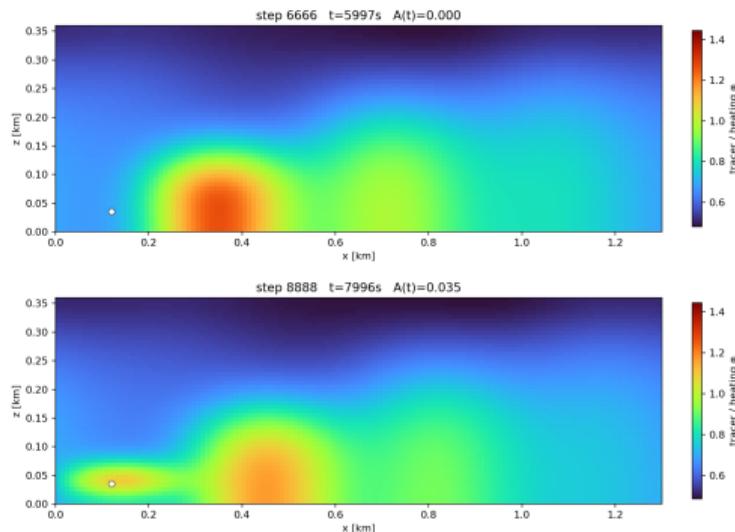
We want a flow where structures move out to the right and re-enter from the left.

Boundary conditions

- ▶ Periodic in x : outflow wraps around
- ▶ Reflective/top treatment in z (plus sponge)

Effect

- ▶ A persistent tracer train forms
- ▶ The system reaches a statistical steady state
- ▶ Ideal for learning time transitions



This creates a clean, cyclic “atmosphere” for ML demos.

RK4 Implementation (Core Loop)

Right-hand side

$$\text{RHS}(\phi, t) = -u \partial_x \phi - w \partial_z \phi + K \nabla^2 \phi + A(t)S(x, z) - (\lambda + \lambda_{\text{top}})\phi.$$

One RK4 step (schematic)

```
k1 = rhs(phi, t)
k2 = rhs(phi + 0.5*dt*k1, t + 0.5*dt)
k3 = rhs(phi + 0.5*dt*k2, t + 0.5*dt)
k4 = rhs(phi + dt*k3,      t + dt)
phi_next = phi + (dt/6)*(k1 + 2*k2 + 2*k3 + k4)
```

Snapshot logic

- ▶ Choose n_{vis} time indices between 0 and n_{steps}
- ▶ Save numbered PNGs: 1_dyn_XX.png

This gives clean training snapshots + nice lecture figures.

Energy / Mass Budget Monitoring (Sanity Checks)

We monitor injected heat content and losses:

$$C(t) = \int \phi \, dx \, dz, \quad I(t) = \int A(t)S \, dx \, dz, \quad L(t) = \int (\lambda + \lambda_{top})\phi \, dx \, dz.$$

Cumulative budget (discrete)

$$\text{acc_in} = \sum_n I(t_n)\Delta t, \quad \text{acc_loss} = \sum_n L(t_n)\Delta t.$$

- ▶ $\int \phi$ stabilizes near equilibrium
- ▶ injected vs lost energy becomes balanced
- ▶ confirms numerical stability and correct forcing/damping design

Always build monitoring into the toy model: it prevents silent nonsense.

Observations: Radiosondes (Sparse Vertical Profiles)

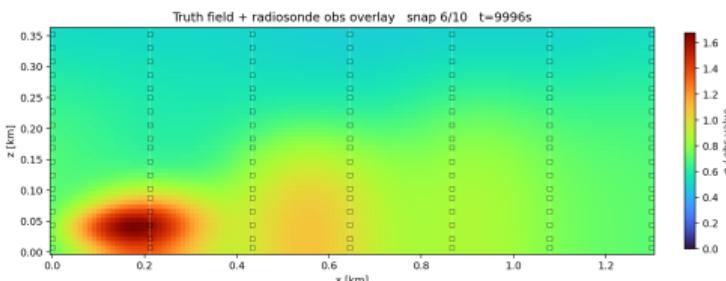
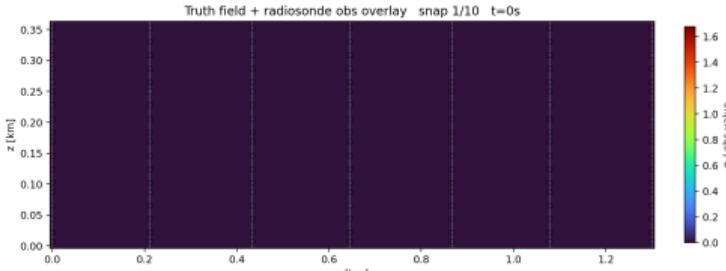
Radiosonde observation operator

At station column x_s and selected heights z_k :

$$y_t^{rs}(s, k) = \phi(x_s, z_k, t) + \epsilon.$$

Geometry is discrete

- ▶ n_{rs} station columns: `rs_ix`
- ▶ n_{vert} vertical levels: `rs_iz`
- ▶ output array: `yrs[time, station, vert]`



RS gives a sparse but physically intuitive reference view.

Observations: Satellite (Vertical Weighting Integrals)

Satellite observation operator

For each channel c (Gaussian vertical weights $w_c(z)$):

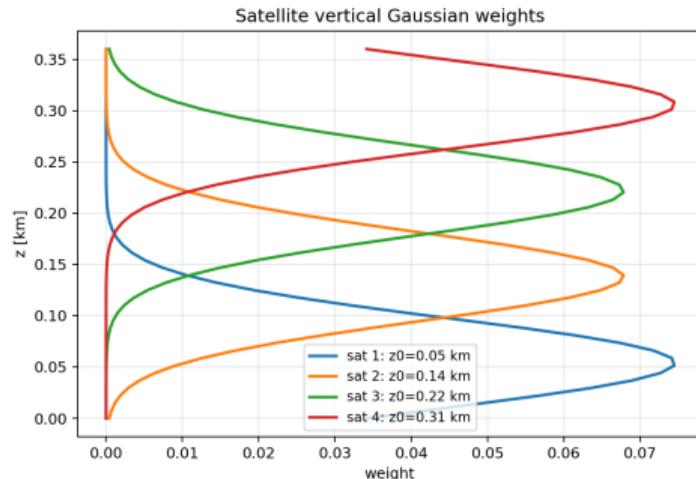
$$y_t^{\text{sat}}(c, x_i) = \sum_{j=1}^{n_z} w_c(z_j) \phi(x_i, z_j, t) \Delta z.$$

Properties

- ▶ available for all columns x_i
- ▶ integrated information \Rightarrow ill-posed inverse problem
- ▶ multiple channels \Rightarrow multi-layer sensitivity

Array

- ▶ `ysat[time, channel, x]`



Gaussian vertical weighting functions used for the SAT channels.

Saved Dataset: Truth + Two Observation Types

We store everything into `dyn_truth_obs.npz`:

- ▶ **Truth fields**

$$x_{true} \in \mathbb{R}^{T_{snap} \times n_z \times n_x}$$

- ▶ **Radiosondes**

$$yrs \in \mathbb{R}^{T_{snap} \times n_{rs} \times n_{vert}}$$

- ▶ **Satellite**

$$ysat \in \mathbb{R}^{T_{snap} \times n_{sat} \times n_x}$$

- ▶ Snapshot times: `t_snap`
- ▶ Geometry: `rs_ix`, `rs_iz`
- ▶ SAT vertical weights: `sat_w`

This allows a clean separation: Notebook 1 = generate dataset, Notebook 2 = learn transitions.

Learning Task: Next-Step Obs Forecasting

We train an ML model that predicts the next observation state:

$$(y_t^{sat}, y_t^{rs}) \mapsto (\hat{y}_{t+1}^{sat}, \hat{y}_{t+1}^{rs}(\cdot)).$$

Inputs at time t

- ▶ Satellite curtain: $y_t^{sat}(c, x)$ for all columns
- ▶ RS set of points: $\{(x_m, z_m, y_m)\}_{m=1}^{N_{in}}$

Outputs at time $t + 1$

- ▶ Predicted satellite curtain $\hat{y}_{t+1}^{sat}(c, x)$
- ▶ Predicted RS values at query points $\hat{y}_{t+1}^{rs}(x_q, z_q)$

Important: query points can be anywhere \Rightarrow generalization to new RS placements .

Normalization: The Key Practical Ingredient

We normalize observations before training:

SAT: per-channel statistics

$$y_n^{sat} = \frac{y^{sat} - \mu_{sat}}{\sigma_{sat}}, \quad \mu_{sat} = \mathbb{E}_{t,x}[y^{sat}], \quad \sigma_{sat} = \text{std}_{t,x}[y^{sat}].$$

RS / truth: global

$$\phi_n = \frac{\phi - \mu_\phi}{\sigma_\phi}, \quad y_n^{rs} = \frac{y^{rs} - \mu_\phi}{\sigma_\phi}.$$

Without normalization the training becomes unstable and “learns the wrong scale”.

Flexible RS Input: Two Dataset Modes (A/B Switch)

We implement a dataset switch (important for interpretability):

Mode A (synthetic RS from truth)

- ▶ sample RS input points randomly from $xtrue[t]$
- ▶ add realistic noise σ_{rs}
- ▶ excellent generalization for arbitrary RS geometry

Mode B (use only stored RS observations)

- ▶ RS input points are exactly $yrs[t,:,:]$ at fixed rs_ix , rs_iz
- ▶ strict statement: “**input uses only observations at time t** ”
- ▶ weaker geometric diversity, but realistic RS network

Reality: We have only fixed radiosondes, but we have airplanes !

Saved:

data_dyn_obs/dyn_truth_obs.npz

Content:

`xtrue` shape=(10,50,130)
`yrs` shape=(10,7,18)
`ysat` shape=(10,4,130)

`t_snap` shape=(10,)
`rs_ix` shape=(7,)
`rs_iz` shape=(18,)
`sat_w` shape=(4,50)

`x` shape=(130,)
`z` shape=(50,)

Architecture: CNN on SAT + Set Encoder for RS + Query Head

Why a CNN?

SAT is a **curtain** $y^{sat}(c, x)$ along x .

Spatial patterns advect \Rightarrow translation-like structure.

Components

- ▶ **SAT encoder:** 1D-CNN
 - ▶ extracts local features along x
- ▶ **RS encoder:** DeepSets / pooling
 - ▶ handles variable-size RS point sets
- ▶ **Query decoder:** predicts $\hat{y}_{t+1}^{rs}(x_q, z_q)$

This enforces coupling: RS inference is driven by SAT structure .

Query-style output

$$\hat{y}_{t+1}^{rs}(x_q, z_q) = g_\theta(\text{CNN}(y_t^{sat}), \text{Set}(y_t^{rs}), x_q, z_q).$$

Practical win

- ▶ RS can be placed anywhere
- ▶ same model predicts:
 - ▶ sparse profiles
 - ▶ full fields (query everywhere)

DeepSets / pooling: why order does not matter

RS inputs form a set $\mathcal{S} = \{(x_m, z_m, y_m)\}_{m=1}^N$. A set has no order, so the encoding must satisfy $F(\mathcal{S}) = F(\pi(\mathcal{S}))$ for any permutation π . We enforce this with a permutation-invariant encoder:

$$e_m = \psi_\theta(x_m, z_m, y_m), \quad E = \text{pool}(e_1, \dots, e_N)$$

where pooling is sum/mean/max (commutative \Rightarrow order-invariant). Locations matter via (x_m, z_m) inside ψ_θ .

Training Loss: Joint SAT + RS Query Targets

We train on snapshot transitions $t \rightarrow t + 1$ using a joint loss:

SAT loss (next curtain)

$$\mathcal{L}_{sat} = \|\hat{y}_{t+1}^{sat} - y_{t+1}^{sat}\|_2^2.$$

RS query loss (next profile at query points)

$$\mathcal{L}_{rs} = \frac{1}{N_q} \sum_{q=1}^{N_q} (\hat{y}_{t+1}^{rs}(x_q, z_q) - y_{t+1}^{rs}(x_q, z_q))^2.$$

Total

$$\mathcal{L} = \mathcal{L}_{sat} + \alpha \mathcal{L}_{rs}.$$

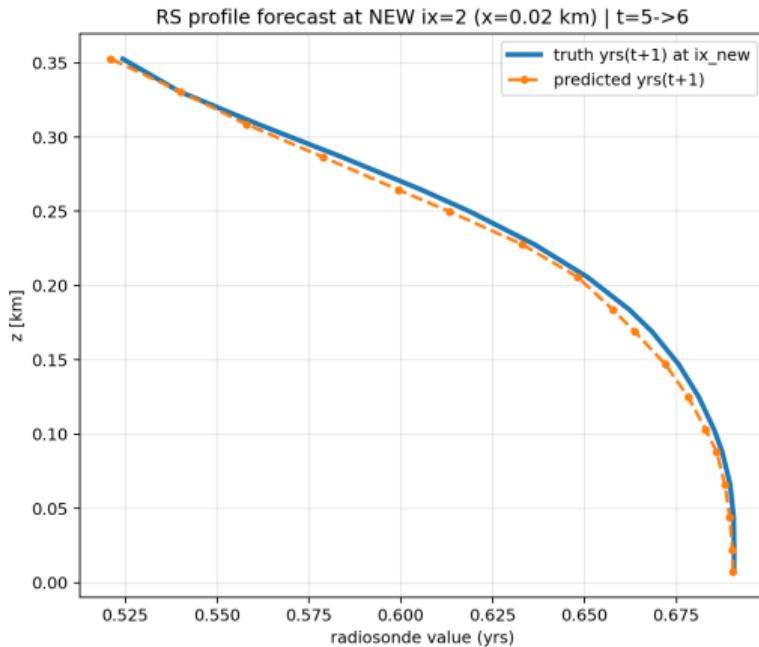
Note: RS targets for arbitrary (x_q, z_q) are taken from truth $xtrue[t + 1]$.

Evaluation 1: Predict a New RS Profile at an Unseen Location

We test generalization :

- ▶ Input RS: taken from stored network $yrs[t, :, :]$
- ▶ Choose a new x column not used by RS
- ▶ Query all heights $\{z_q\}$ at that x
- ▶ Compare predicted profile \hat{y}_{t+1}^{rs} vs truth $xtrue[t + 1]$

This checks whether the model learned physics-like transport information from SAT.



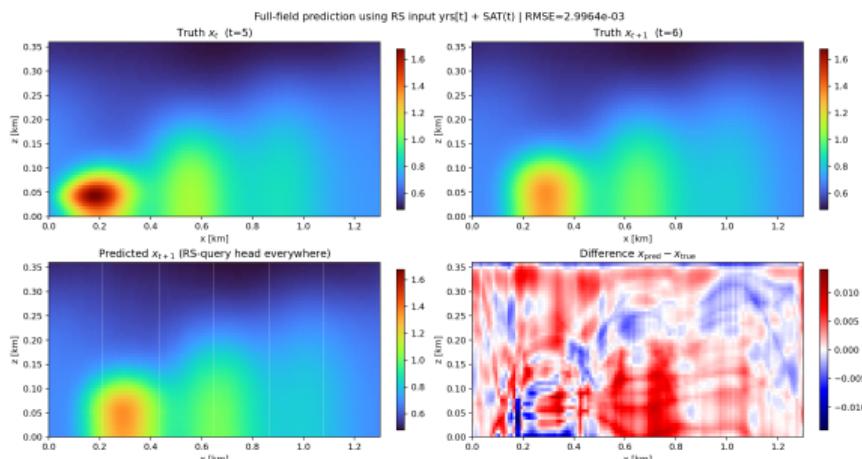
Evaluation 2: Full-Field Reconstruction at $t + 1$ (Query Everywhere)

The RS query head can be evaluated on the whole grid:

$$x_{\text{pred}}(x_i, z_j, t + 1) := \hat{y}_{t+1}^{rs}(x_i, z_j).$$

We visualize

- ▶ $x_{\text{true}}[t]$ (reference)
- ▶ $x_{\text{true}}[t + 1]$ (truth)
- ▶ $x_{\text{pred}}[t + 1]$ (prediction)
- ▶ difference $x_{\text{pred}} - x_{\text{true}}$



This turns obs-to-obs learning into a state-like field prediction.

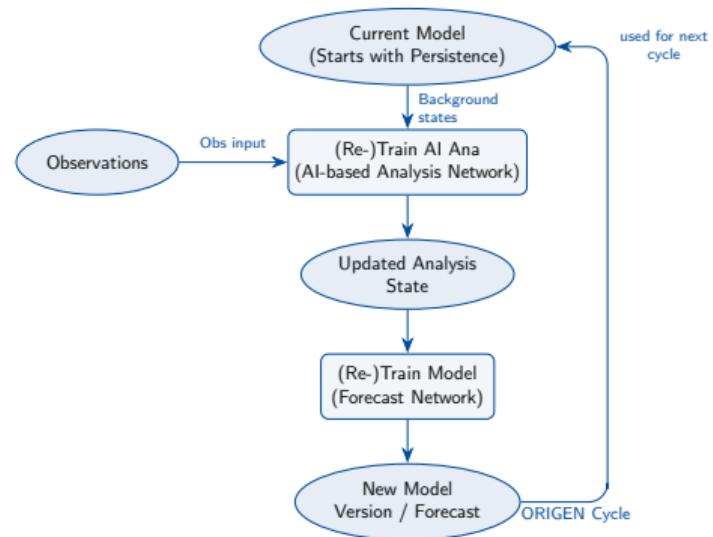
Observation-based Reconstruction & Inversion for Generative Emulation of Nonlinear Systems: ORIGEN

ORIGEN cycle schematic

- ▶ A closed learning cycle linking observations , analysis reconstruction , and forecast models .

Conceptual loop

- ▶ Start from a current model (here: persistence / baseline)
- ▶ Use observations to reconstruct state:
AI analysis / inversion
- ▶ Update the model/forecast emulator from reconstructed states
- ▶ Iterate: next cycle uses the new model



ORIGEN on a Minimal Example: Simple Oscillator

Minimal testbed (2D circle)

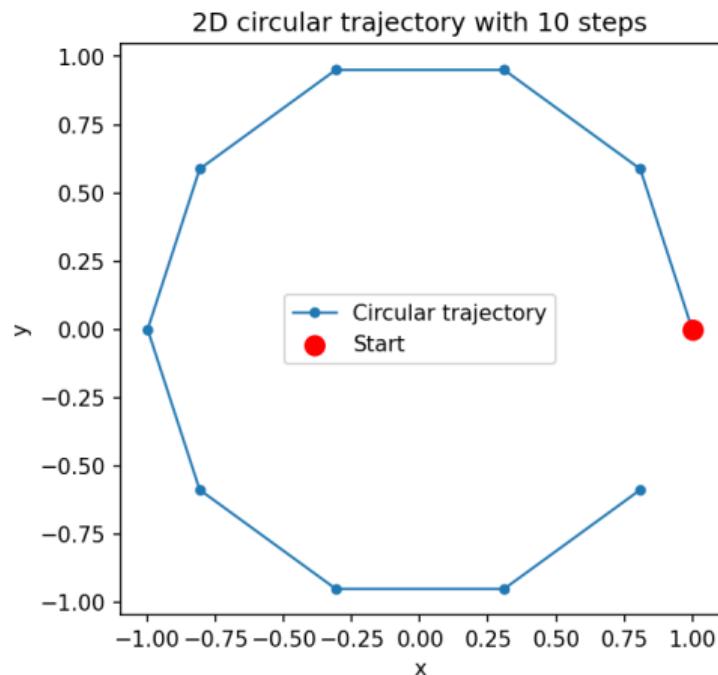
- ▶ Truth trajectory $x_k = (x_{1,k}, x_{2,k})$ on a circle

$$x_k = \begin{bmatrix} \cos \theta_k \\ \sin \theta_k \end{bmatrix}, \quad \theta_k = \frac{2\pi k}{n}$$

- ▶ Observations are scalar and partial :

$$y_k = H_k x_k + \epsilon_k, \quad H_k \in \{[1, 0], [0, 1]\}$$

Aim: Explain ORIGEN mechanics without complex dynamics.



Observations: Partial Measurements of x_1 or x_2

Time-dependent observation operator

- At each step k we observe only one component

Selector:

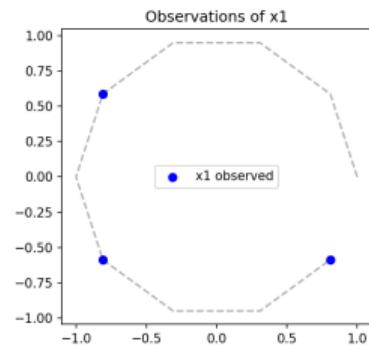
$$s_k \in \{1, 2\}$$

- Observation operator:

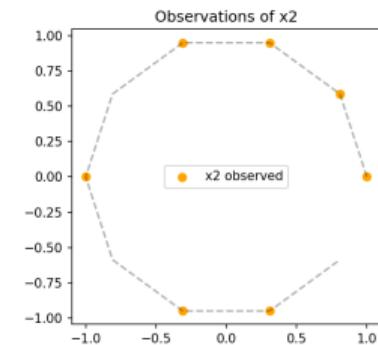
$$H_k = \begin{cases} [1, 0] & s_k = 1 \text{ (observe } x_1) \\ [0, 1] & s_k = 2 \text{ (observe } x_2) \end{cases}$$

- Observation equation:

$$y_k = H_k x_k + \epsilon_k, \quad \epsilon_k \sim \mathcal{N}(0, R)$$



Blue: x_1 observed



Orange: x_2 observed

Interpretation: This mimics heterogeneous sensors and missing data .

3D-Var Step: Background → Analysis (Single Observation)

We combine background + one scalar obs

- ▶ Background state: $x_k^b \in \mathbb{R}^2$
- ▶ Obs: $y_k \in \mathbb{R}$, operator $H_k \in \mathbb{R}^{1 \times 2}$

Innovation

$$d_k = y_k - H_k x_k^b$$

Gain (scalar obs)

$$K_k = B H_k^\top \left(H_k B H_k^\top + R \right)^{-1}$$

Analysis update

$$x_k^a = x_k^b + K_k d_k$$

Effect: update only in observed direction, but shaped by B .

Initial choice of covariances

$$B = \begin{bmatrix} \sigma_b^2(x_1) & 0 \\ 0 & \sigma_b^2(x_2) \end{bmatrix}, \quad R = \sigma_o^2$$

Cycling (persistence model)

$$x_{k+1}^b = M(x_k^a), \quad M(x) = x$$

- ▶ This produces a full sequence:

$$x_0^b \rightarrow x_0^a \rightarrow x_1^b \rightarrow \dots$$

- ▶ We visualize results next: truth vs background vs analysis

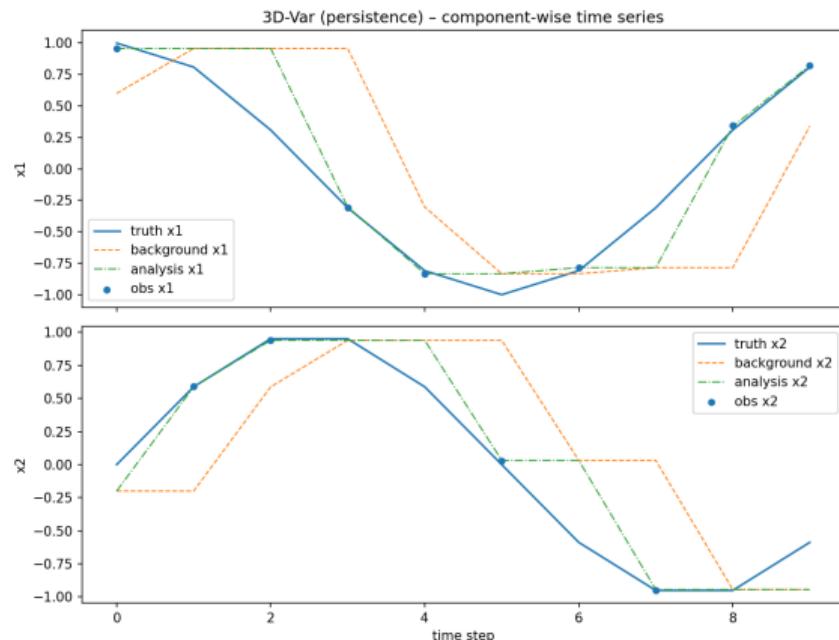
3D-Var Cycle Result: Time Series of x_1 and x_2

What we compare

- ▶ Truth x_k
- ▶ Background x_k^b
- ▶ Analysis x_k^a
- ▶ and the scalar observations y_k

Key behavior

- ▶ analyses are pulled towards the obs
- ▶ cycling propagates improvements forward
- ▶ unobserved component still benefits indirectly (via B and cycling)



Idea: Information is collected iteratively.

Sequential Rounds: With All Observations We Converge in Two Cycles

Idea: assimilate in two rounds

- ▶ Round 1: observe x_1 only

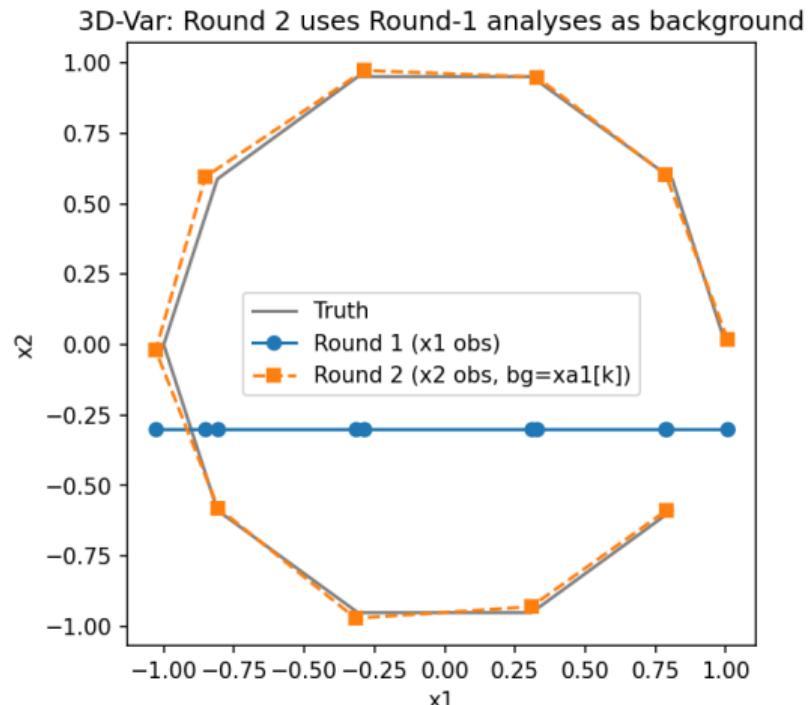
$$y_k^{(1)} = x_{1,k} + \epsilon_k$$
- ▶ Round 2: observe x_2 only, using

$$x_k^{b,(2)} := x_k^{a,(1)}$$

If obs error is small

- ▶ after Round 1: x_1 aligns with truth
- ▶ after Round 2: x_2 aligns with truth
- ▶ both components observed ⇒ fast convergence

Message: With complete information,
ORIGEN-style cycling converges quickly.



Iterative 3D-Var: Reconstruction Improves over Cycles (Demo)

What is iterated here?

- We reconstruct an entire trajectory

$$\{x_k^a\}_{k=0}^{n-1}$$

from noisy, partial observations

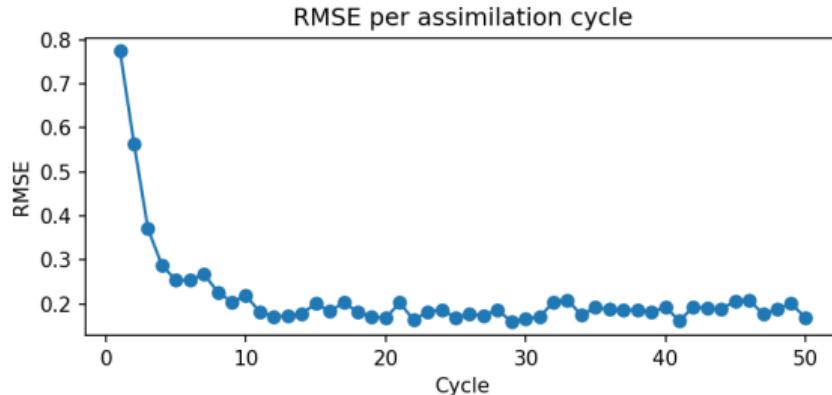
$$y_k = H_k x_k + \epsilon_k$$

- Each cycle produces a refined estimate of the full sequence:

$$\{x_k^{a,(c)}\} \Rightarrow \{x_k^{a,(c+1)}\}$$

Crucial point

- The “dynamics” is implicit in the reconstructed series
- We are not propagating with a separate model; we repeatedly improve the sequence itself



RMSE per assimilation cycle: iterative 3D-Var improves the reconstructed trajectory

ORIGEN: from the reconstructed $\{x_k^a\}$ we learn a forecast map

$$x^a(t) \mapsto x^a(t + \Delta t),$$

in the notebook we focus on trajectory reconstruction.

Reconstruction: 2D Trajectory over Iterative 3D-Var Cycles

Iterative 3D-Var reconstruction

- ▶ Each cycle reconstructs the full sequence:

$$\{x_k^a\}_{k=0}^{n-1}$$

- ▶ Observations remain **noisy and partial**:

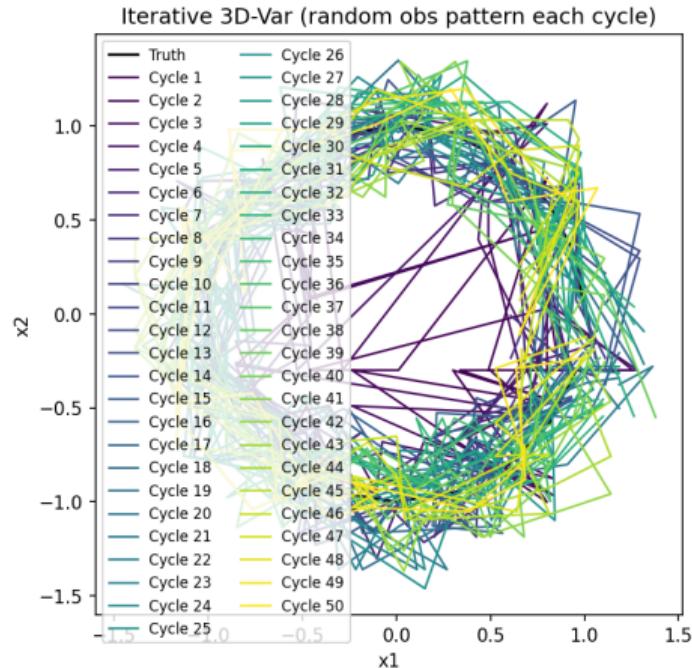
$$y_k = H_k x_k + \epsilon_k$$

- ▶ Repeated cycles reduce reconstruction error

What you see on the right

- ▶ Truth trajectory (black)
- ▶ Selected analysis cycles (colored)
- ▶ Later cycles are closer to truth

Message: 3D-Var keeps error coming in from the observation error



Reconstruction: Cycle Time Series with Noisy Observations

Final-cycle diagnostics

- ▶ Compare truth vs final reconstructed trajectory
- ▶ Show observations for the last random pattern:

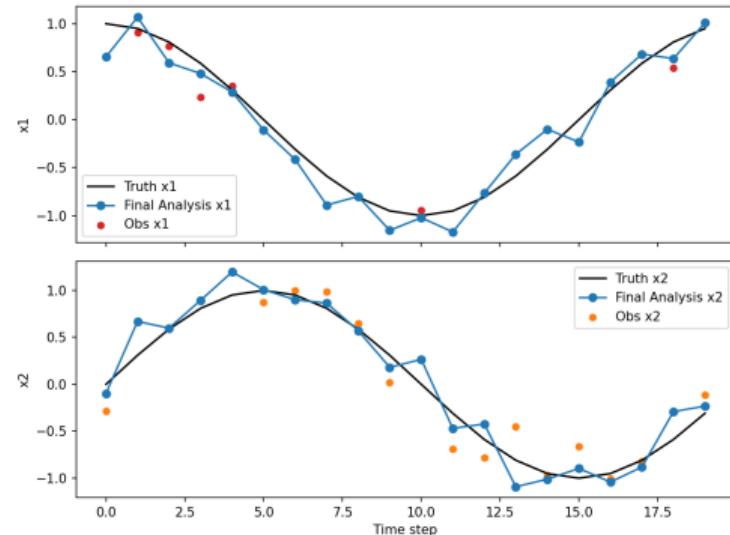
$$H_k \in \{[1, 0], [0, 1]\}$$

- ▶ Observations are noisy \Rightarrow analysis does not overfit

3D-Var property

- ▶ analysis is a weighted compromise
- ▶ controlled by B and R

Message: 3D-VAR is not enough



Final cycle: time series for x_1 and x_2 with noisy partial observations

Adding a Covariance Update: Kalman Filter–Type Uncertainty Reduction

Upgrade: update B during cycling

- ▶ Until now: fixed background covariance B
- ▶ Now: after each analysis step we also update uncertainty

Kalman filter covariance update

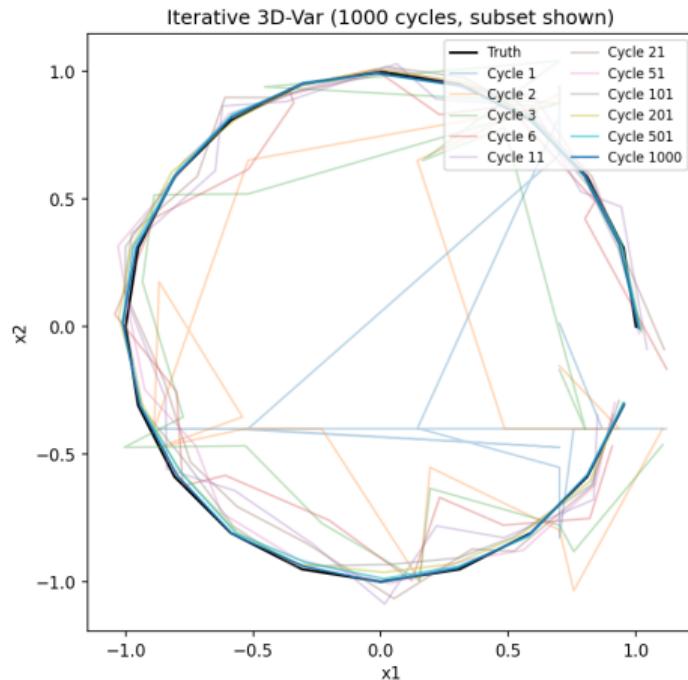
$$A_k = (I - K_k H_k) B_k \quad \Rightarrow \quad B_k \leftarrow A_k$$

Effect:

- ▶ uncertainty shrinks in observed directions
- ▶ gain adapts across cycles
- ▶ analysis trajectory converges strongly to truth

Message:

With adaptive B , iterative DA converges.



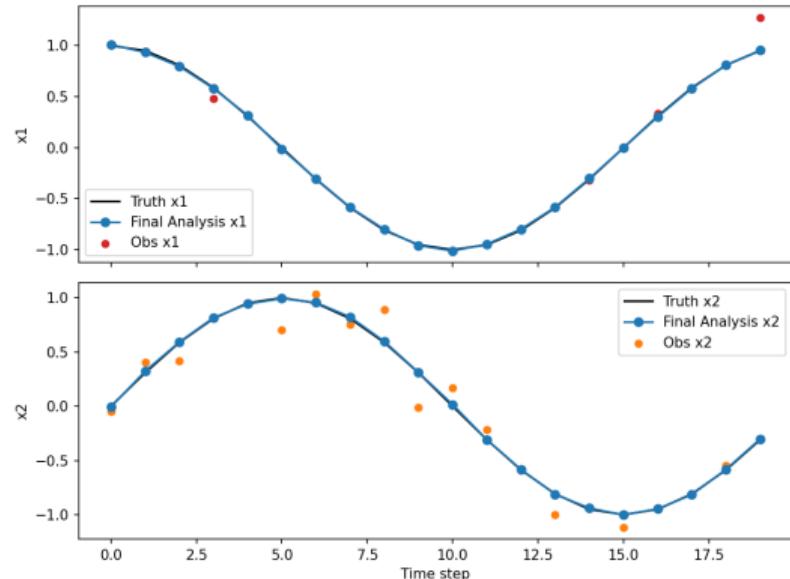
Full Convergence: Final Trajectory with Covariance Update

Final-cycle reconstruction

- ▶ With B update, repeated cycling yields near-perfect reconstruction
- ▶ Even with noisy and partial observations

Why it converges

- ▶ analysis reduces both error and uncertainty
- ▶ smaller uncertainty \Rightarrow more consistent updates
- ▶ loop stabilizes around the truth trajectory



Message: ORIGEN principle: reconstruction + uncertainty adaptation \Rightarrow convergence.

Final cycle: time series x_1, x_2 converged

Lorenz-63: ORIGEN Reconstruction & Learned Forecast Model

ORIGEN loop (concept)

1. Assimilate random observed subsets (x, y, z, xy, xz, yz)
2. Update background $\mathbf{x}^b \leftarrow \mathbf{x}^a$
3. Update covariance $B \leftarrow P^a$ (Kalman-style)

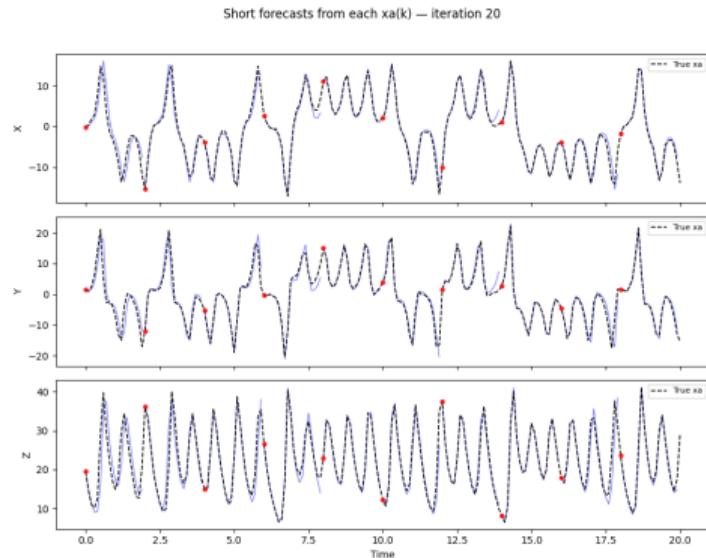
Model learning

- ▶ Train NN to learn one-step map

$$\mathbf{x}^a(k) \mapsto \mathbf{x}^a(k+1)$$

- ▶ **Fine-Tuning with Rollout**

Rollout: many subsequent short forecasts



Forecast rollouts (blue) starting from analysis states (red dots), compared to the reconstructed reference trajectory (black dashed).